

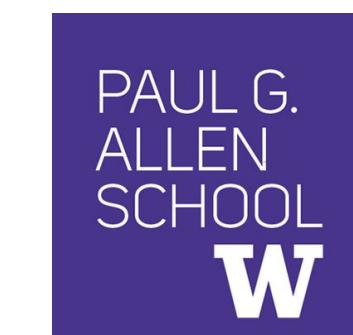
Bridging Textual and Tabular Data for Cross-Domain Text-to-SQL Semantic Parsing

Paul G. Allen School of Computer Science and Engineering



Ph.D. Qualification Exam - Winter'21

Advisors: Luke Zettlemoyer and Yejin Choi



Victoria Lin



Work done at Salesforce AI Research

Personal

Intelligent

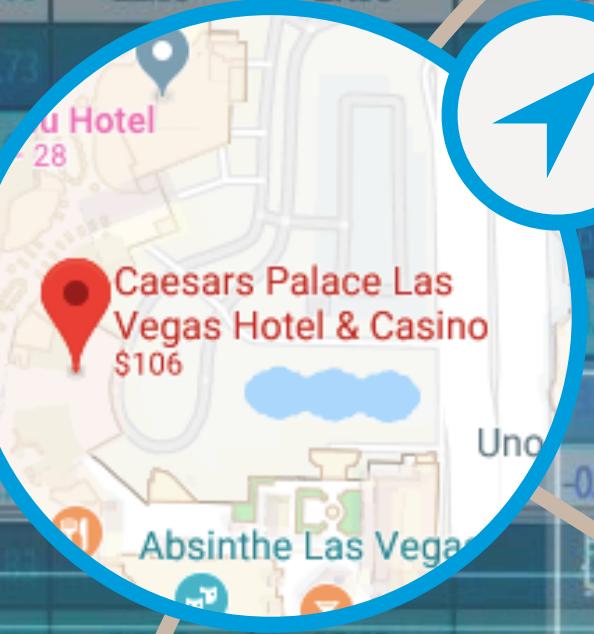
Conversational

Tables and databases are commonly used data structures that power a variety of downstream applications.

Personal



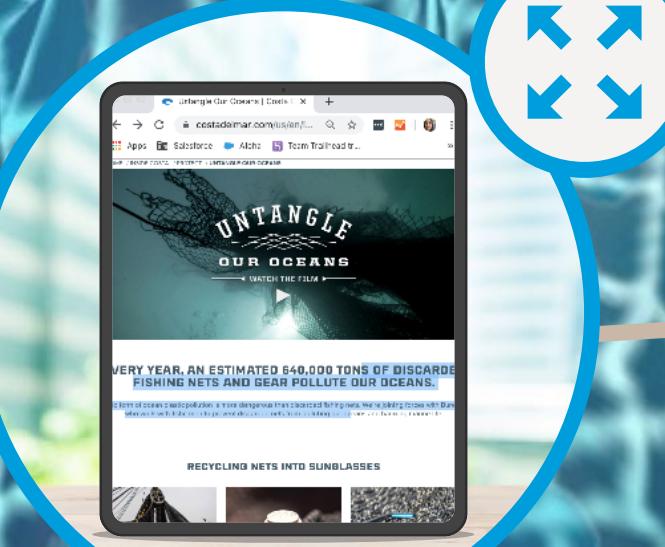
Intelligent



Conversational



The Internet contains billions of public tables. And more tabular data can be found on corporate intranets and personal devices.



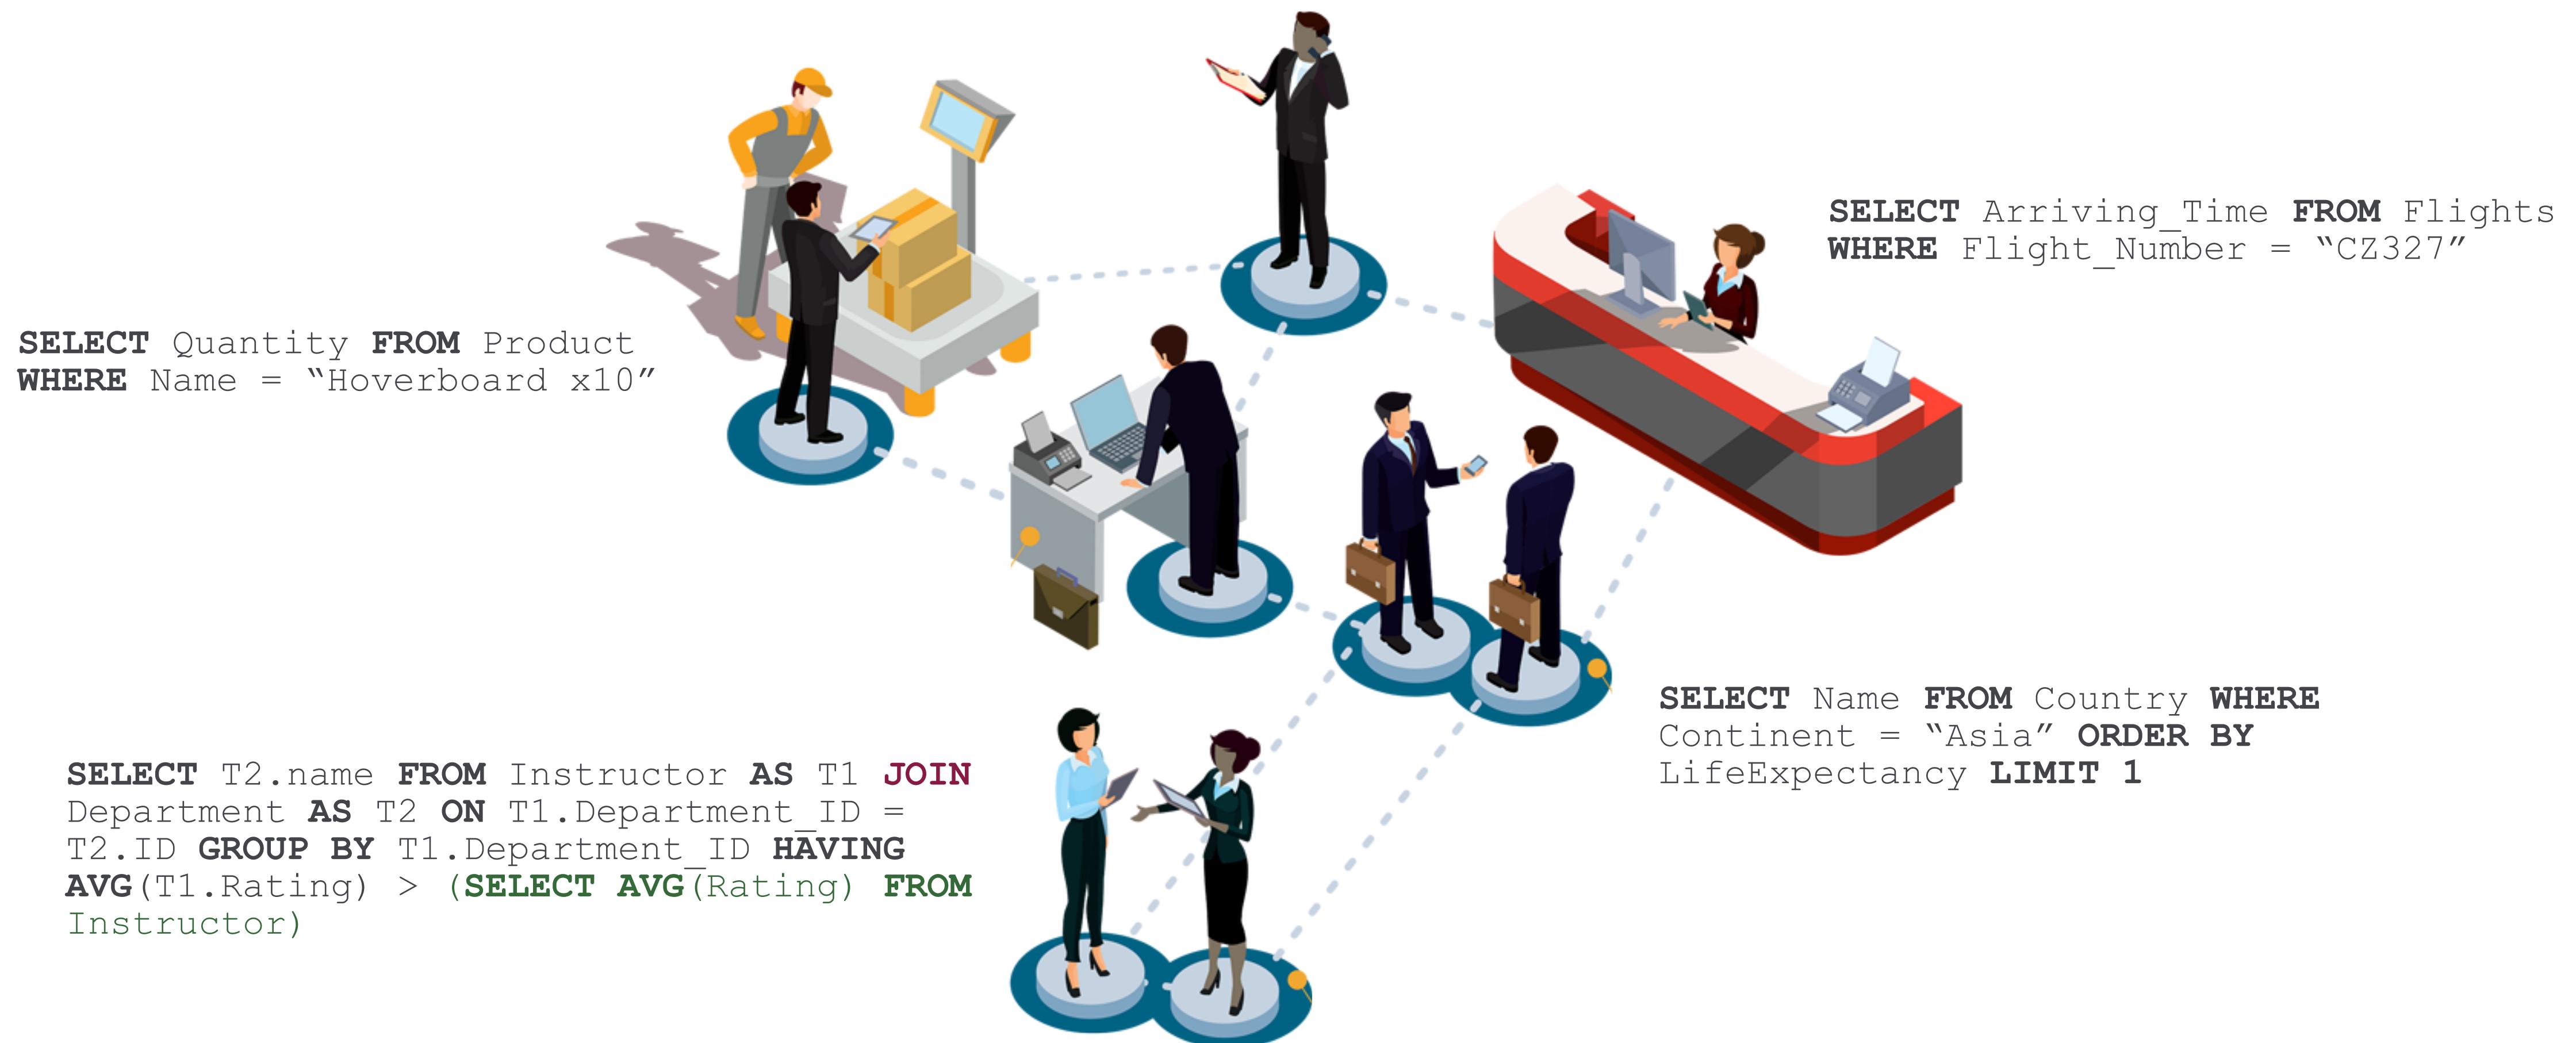
Natural Language Interface to Databases

Traditionally, users access databases using structured query language (SQL).



Natural Language Interface to Databases

Traditionally, users access databases using structured query language (SQL).



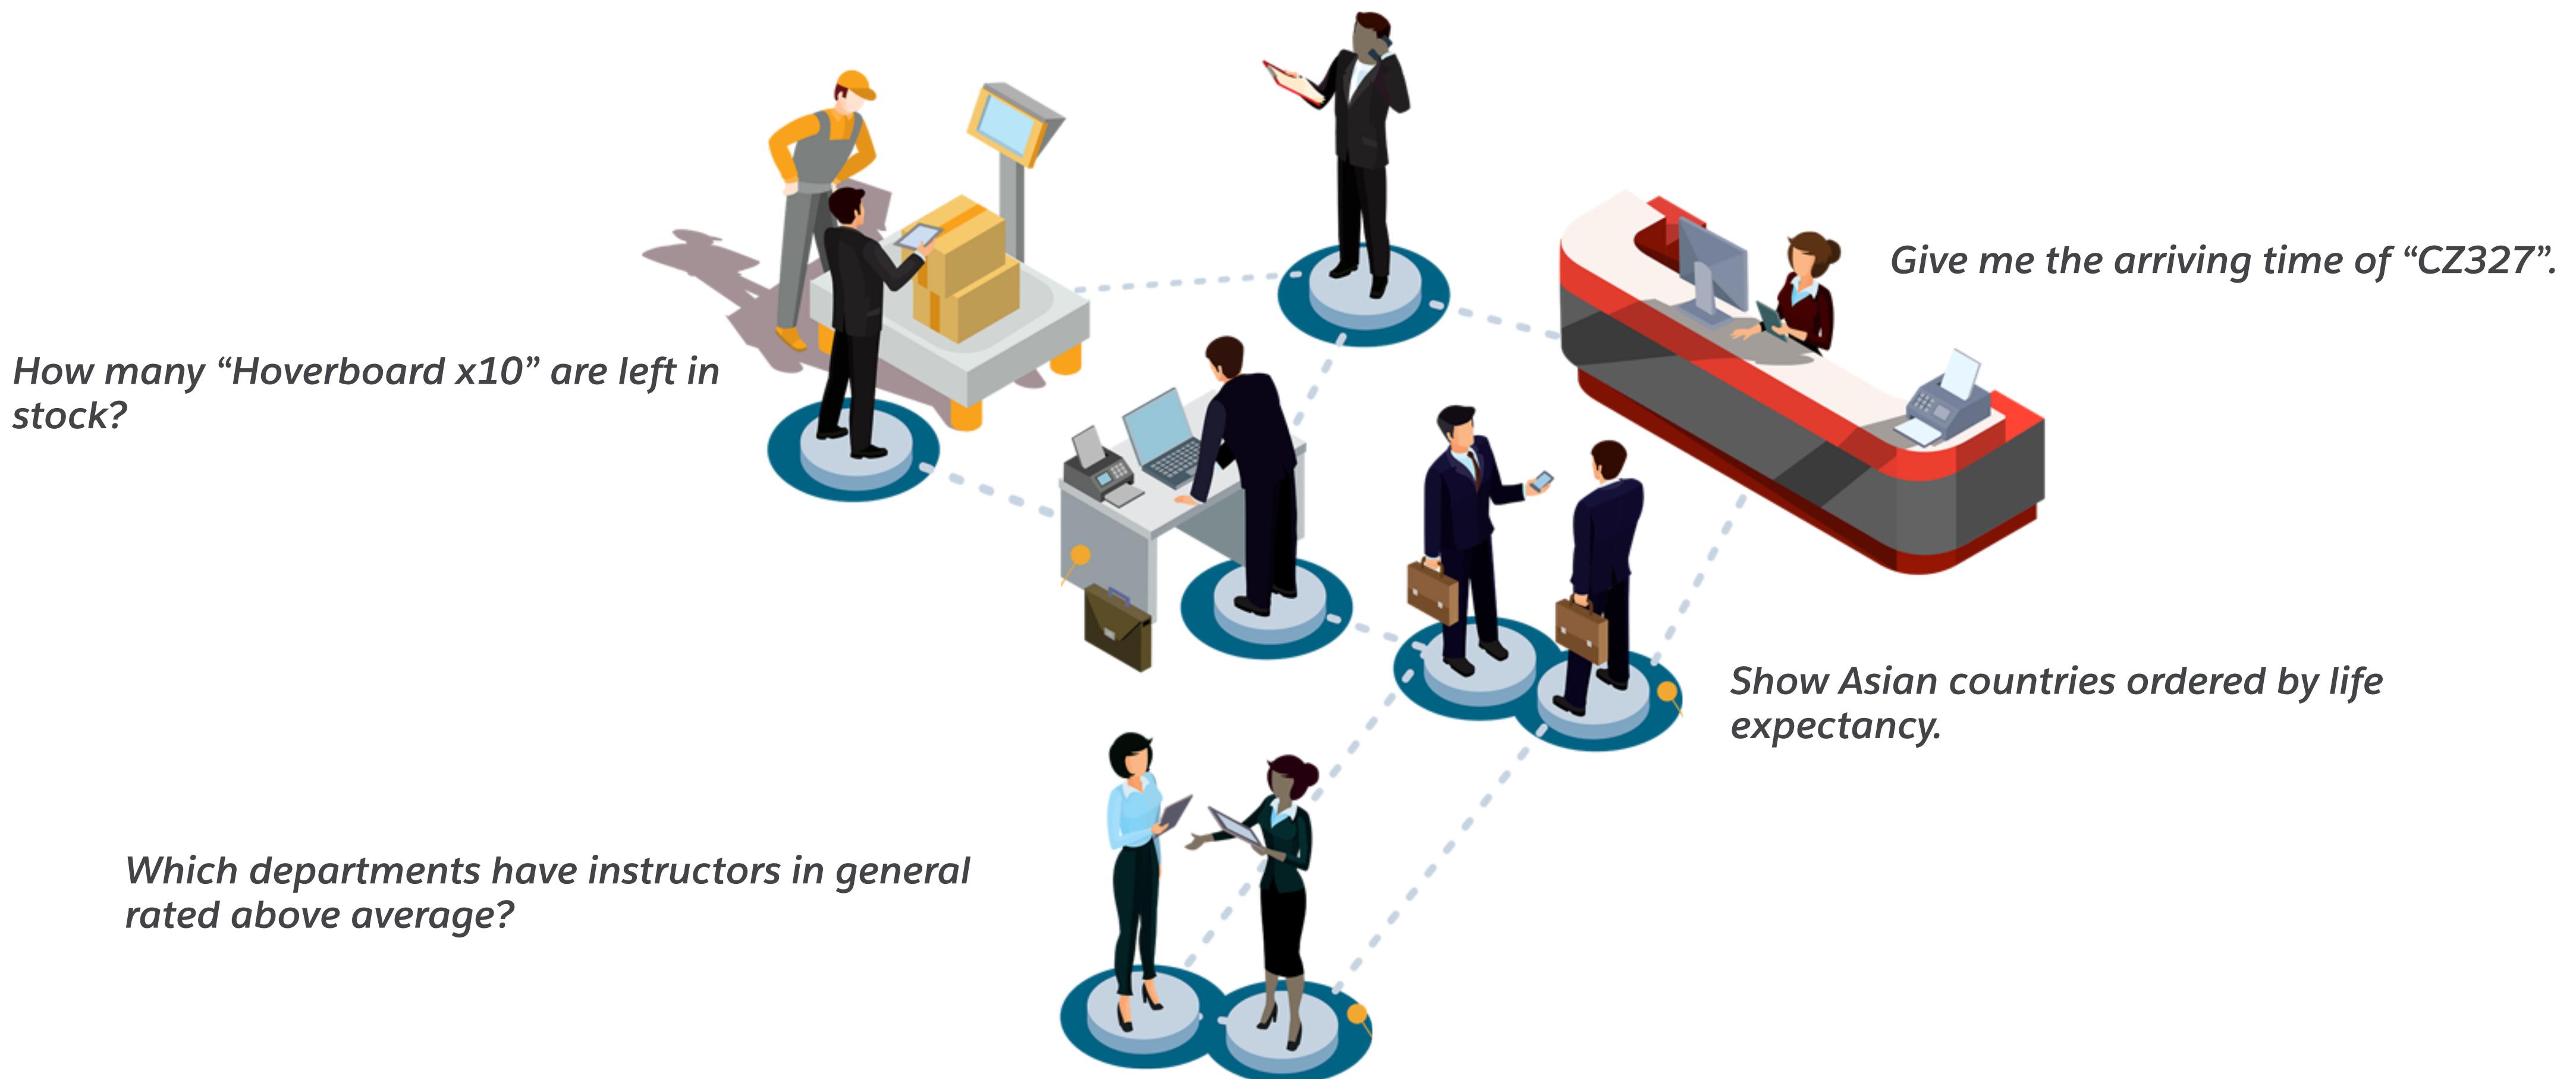
Natural Language Interface to Databases

Our goal is to learn semantic parsers that map natural language utterances to **executable** SQL queries for **any database**.



Natural Language Interface to Databases

Our goal is to learn semantic parsers that map natural language utterances to **executable SQL queries for any database**.





research

<https://naturalsql.com>

Photon About

real_estate_property

Tables

- Other_Available_Features**

feature_id	feature_type_code	feature_name	feature_description
2	Amenity	AirCon	Air Conditioning.
3	Amenity	Pool	Swimming Pool.
4	Security	BurglarAlarm	Burglar Alarm

5 << < 1 > >>

- Ref_Feature_Types**

feature_type_code	feature_type_name
Amenity	Amenity, eg Pool.
Security	Securiyt, eg Burglar Alarm.

5 << < 1 > >>

- Other_Property_Features**

property_id	feature_id	property_feature_description
15	3	dolorem
12	4	earum
6	2	illo
7	2	hic
6	3	et

5 << < 1 2 3 > >>

- Properties**

property_id	property_type_code	date_on_market	date_sold	property_name	property_address
1	House	1991-06-21 23:52:10	1979-05-13 16:58:06	park	4745 Emerson Stravenue Suite 829 South
2	House	1990-05-25 23:01:51	1990-11-14 19:16:38	the cole	098 Tremaine Highway Suite 569 South W
3	Other	1986-11-26 04:12:18	1981-06-26 21:28:28	prism	062 Micaela Court Apt. 707 Margretville, V
4	Field	2017-09-14 15:49:23	2003-02-27 18:17:11	riverside	49578 Ayden Mountains New Russellhave
5	Apartment	2016-05-06 16:53:39	2012-08-19 07:36:57	parc east	2765 Schulist Stream Lindmouth, UT 0339

5 << < 1 2 3 > >>

salesforce research

Chat started by Photon • 3:43:58 AM

list name of all technicians and the machines they repair

```
SELECT technician.Name FROM repair_assignment JOIN machine ON repair_assignment.Machine_ID = machine.Machine_ID JOIN technician ON repair_assignment.technician_id = technician.technician_id
```

Did I get it right?

no

Sorry...

Please rephrase your previous question or ask a new question

List names of technician and the machines series they repaired

```
SELECT technician.Name, machine.Machine_series FROM repair_assignment JOIN machine ON repair_assignment.Machine_ID = machine.Machine_ID JOIN technician ON repair_assignment.technician_id = technician.technician_id
```

Did I get it right?

yes

Great!

Please input your question in NL or SQL to query the DB

Type Here

Text-to-SQL translation

Query arbitrary subsets of tables in any database

Support table join and other complex SQL operators



research

<https://naturalsql.com>

Photon About

real_estate_property

Tables

- Other_Available_Features**

feature_id	feature_type_code	feature_name	feature_description
2	Amenity	AirCon	Air Conditioning.
3	Amenity	Pool	Swimming Pool.
4	Security	BurglarAlarm	Burglar Alarm

5 << < 1 > >>

- Ref_Feature_Types**

feature_type_code	feature_type_name
Amenity	Amenity, eg Pool.
Security	Securiyt, eg Burglar Alarm.

5 << < 1 > >>

- Other_Property_Features**

property_id	feature_id	property_feature_description
15	3	dolorem
12	4	earum
6	2	illo
7	2	hic
6	3	et

5 << < 1 2 3 > >>

- Properties**

property_id	property_type_code	date_on_market	date_sold	property_name	property_address
1	House	1991-06-21 23:52:10	1979-05-13 16:58:06	park	4745 Emerson Stravenue Suite 829 South
2	House	1990-05-25 23:01:51	1990-11-14 19:16:38	the cole	098 Tremaine Highway Suite 569 South W
3	Other	1986-11-26 04:12:18	1981-06-26 21:28:28	prism	062 Micaela Court Apt. 707 Margretville, V
4	Field	2017-09-14 15:49:23	2003-02-27 18:17:11	riverside	49578 Ayden Mountains New Russellhave
5	Apartment	2016-05-06 16:53:39	2012-08-19 07:36:57	parc east	2765 Schulist Stream Lindmouth, UT 0339

5 << < 1 2 3 > >>

salesforce research

Chat started by Photon • 3:43:58 AM

list name of all technicians and the machines they repair

```
SELECT technician.Name FROM repair_assignment JOIN machine ON repair_assignment.Machine_ID = machine.Machine_ID JOIN technician ON repair_assignment.technician_id = technician.technician_id
```

Did I get it right?

no

Sorry...

Please rephrase your previous question or ask a new question

List names of technician and the machines series they repaired

```
SELECT technician.Name, machine.Machine_series FROM repair_assignment JOIN machine ON repair_assignment.Machine_ID = machine.Machine_ID JOIN technician ON repair_assignment.technician_id = technician.technician_id
```

Did I get it right?

yes

Great!

Please input your question in NL or SQL to query the DB

Type Here

Text-to-SQL translation

Query arbitrary subsets of tables in any database

Support table join and other complex SQL operators

Relational Database Terms

User Profile

UserID	Name	Nationality	PartitionID	# Followers	...
103041	Smith, John	Canada	P10	3	...
103042	Hanks, Tom	United States	P11	16.6M	...
103043	Obama, Barack	United States	P11	127M	...
...

 A table represents a entity type (or event type).

Relational Database Terms

User Profile ← Table Name

UserID	Name	Nationality	PartitionID	# Followers	...
103041	Smith, John	Canada	P10	3	...
103042	Hanks, Tom	United States	P11	16.6M	...
103043	Obama, Barack	United States	P11	127M	...
...

Relational Database Terms

User Profile

UserID	Name	Nationality	PartitionID	# Followers	...
103041	Smith, John	Canada	P10	3	...
103042	Hanks, Tom	United States	P11	16.6M	...
103043	Obama, Barack	United States	P11	127M	...
...

Table Header

Relational Database Terms

User Profile

UserID	Name	Nationality	PartitionID	# Followers	...
103041	Smith, John	Canada	P10	3	...
103042	Hanks, Tom	United States	P11	16.6M	...
103043	Obama, Barack	United States	P11	127M	...
...

Column Name

Relational Database Terms

User Profile

UserID	Name	Nationality	PartitionID	# Followers	...
103041	Smith, John	Canada	P10	3	...
103042	Hanks, Tom	United States	P11	16.6M	...
103043	Obama, Barack	United States	P11	127M	...
...

Row



A row is an instantiation of the entity/event.

Relational Database Terms

User Profile

UserID	Name	Nationality	PartitionID	# Followers	...
103041	Smith, John	Canada	P10	3	...
103042	Hanks, Tom	United States	P11	16.6M	...
103043	Obama, Barack	United States	P11	127M	...
...

column/field



Relational Database Terms

User Profile

UserID	Name	Nationality	PartitionID	# Followers	...
103041	Smith, John	Canada	P10	3	...
103042	Hanks, Tom	United States	P11	16.6M	...
103043	Obama, Barack	United States	P11	127M	...
...

Integer String Integer

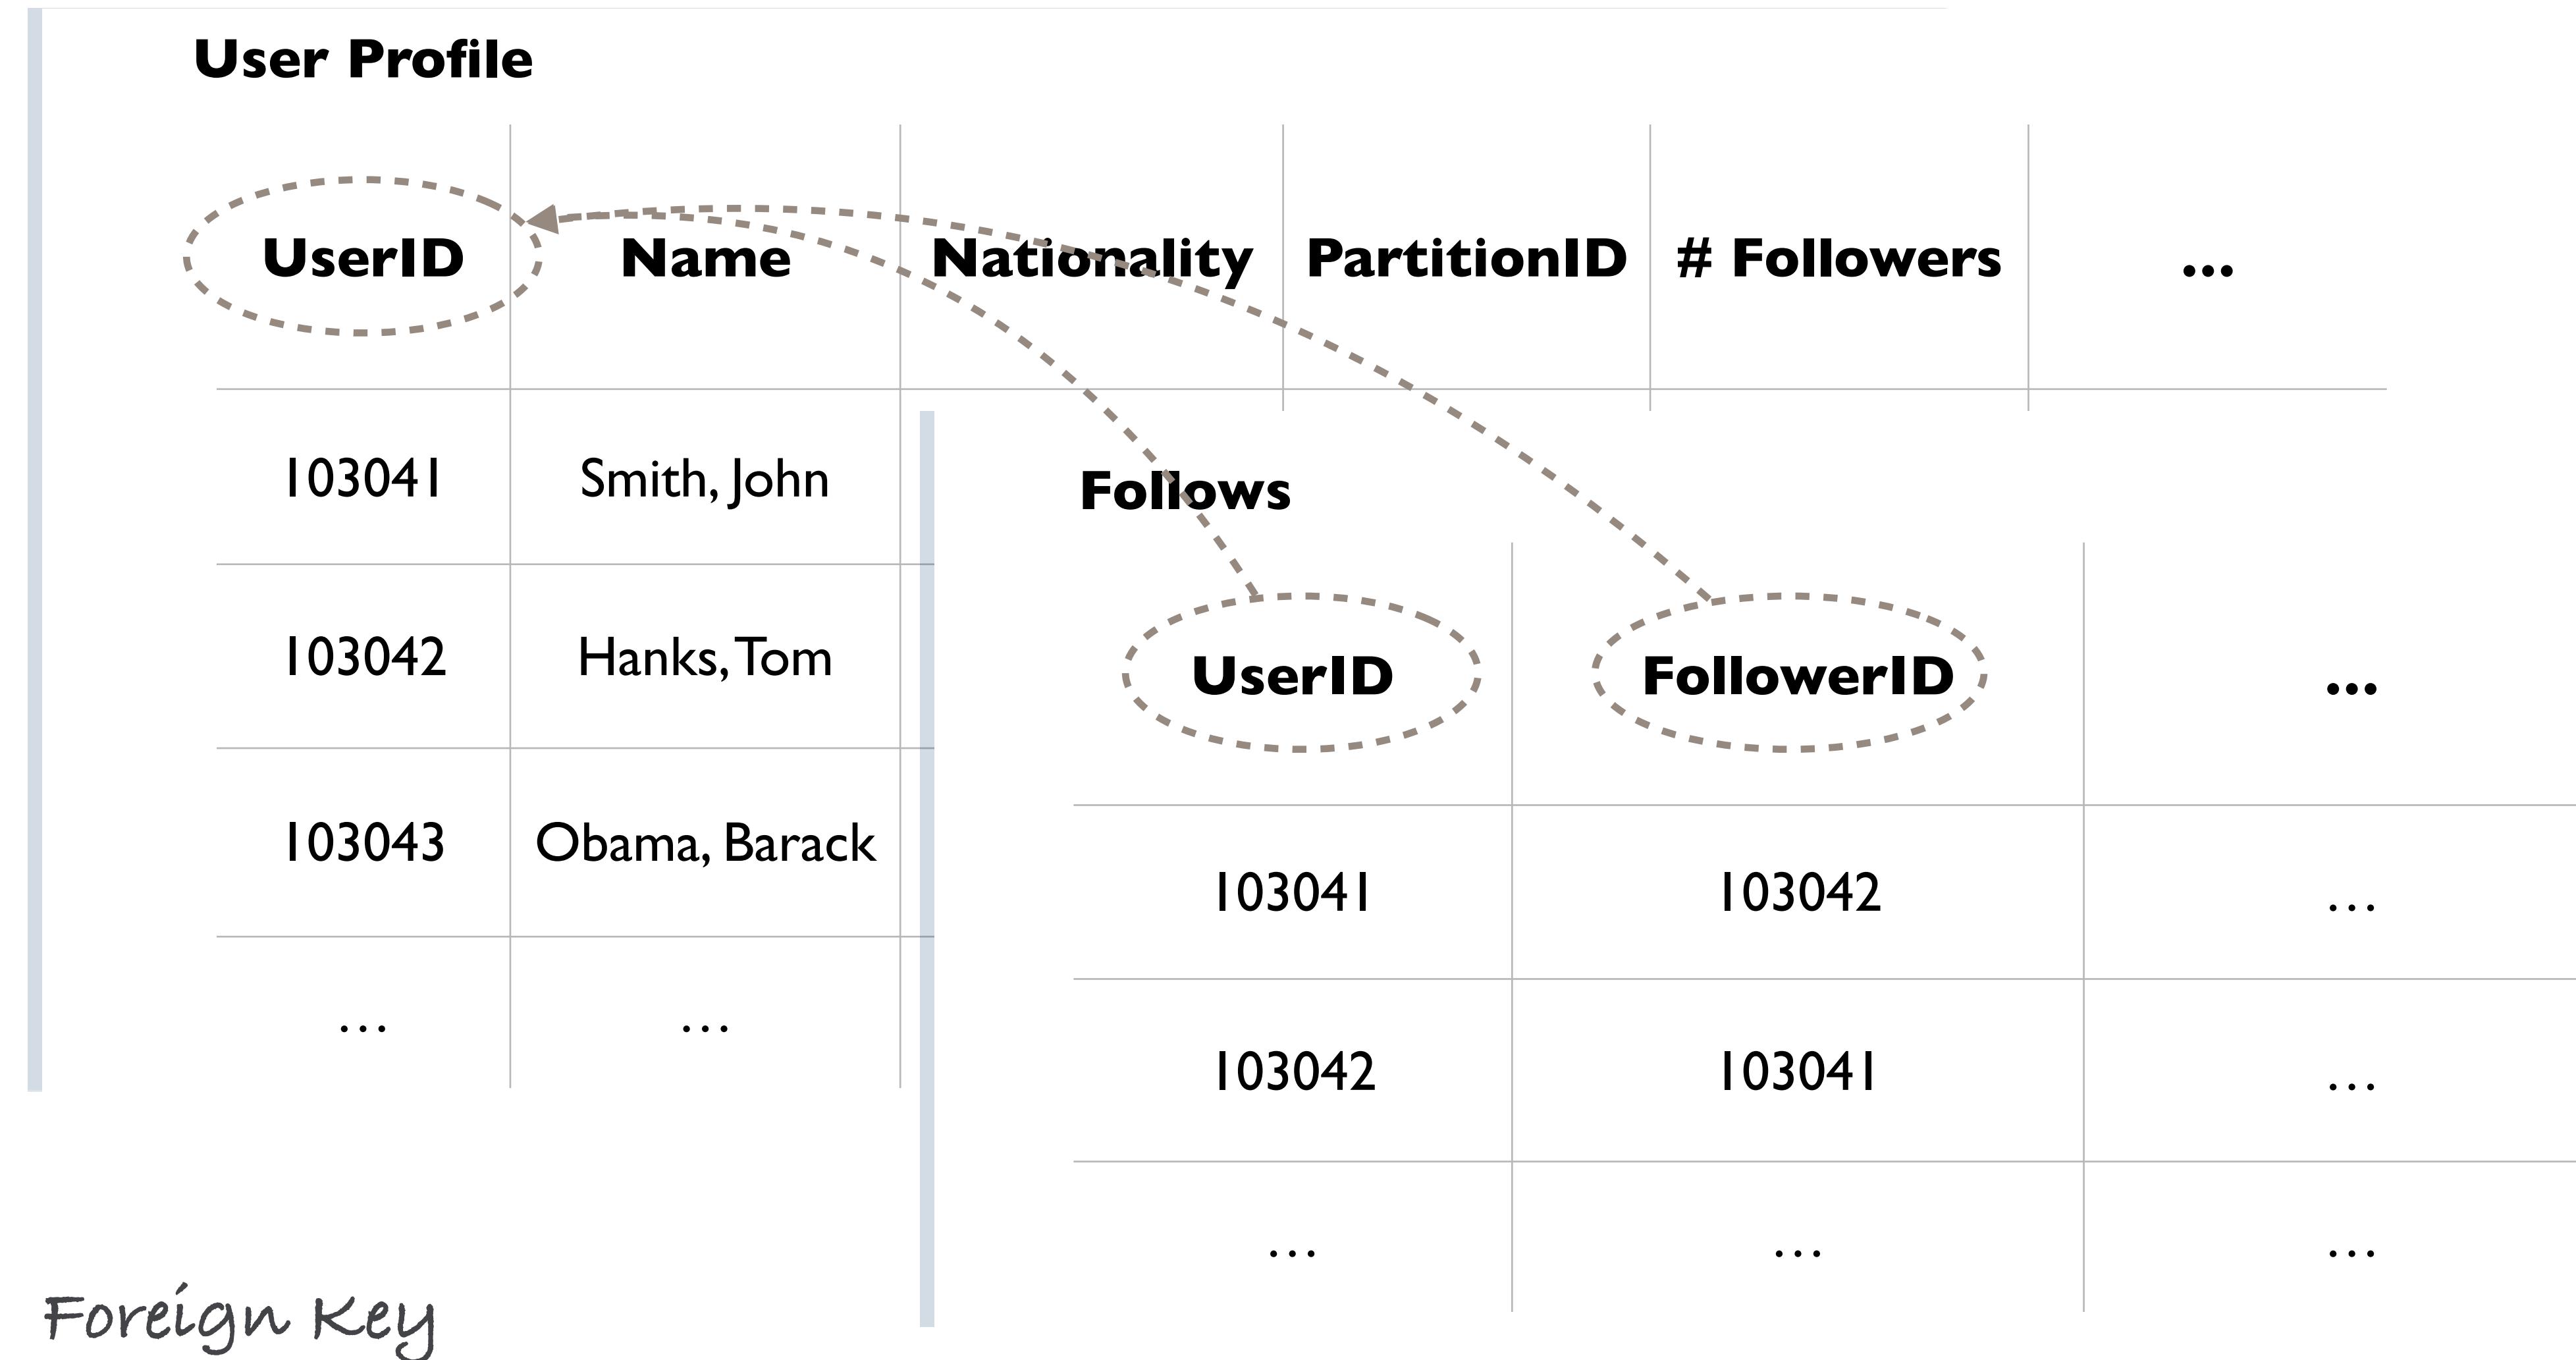
Relational Database Terms

User Profile

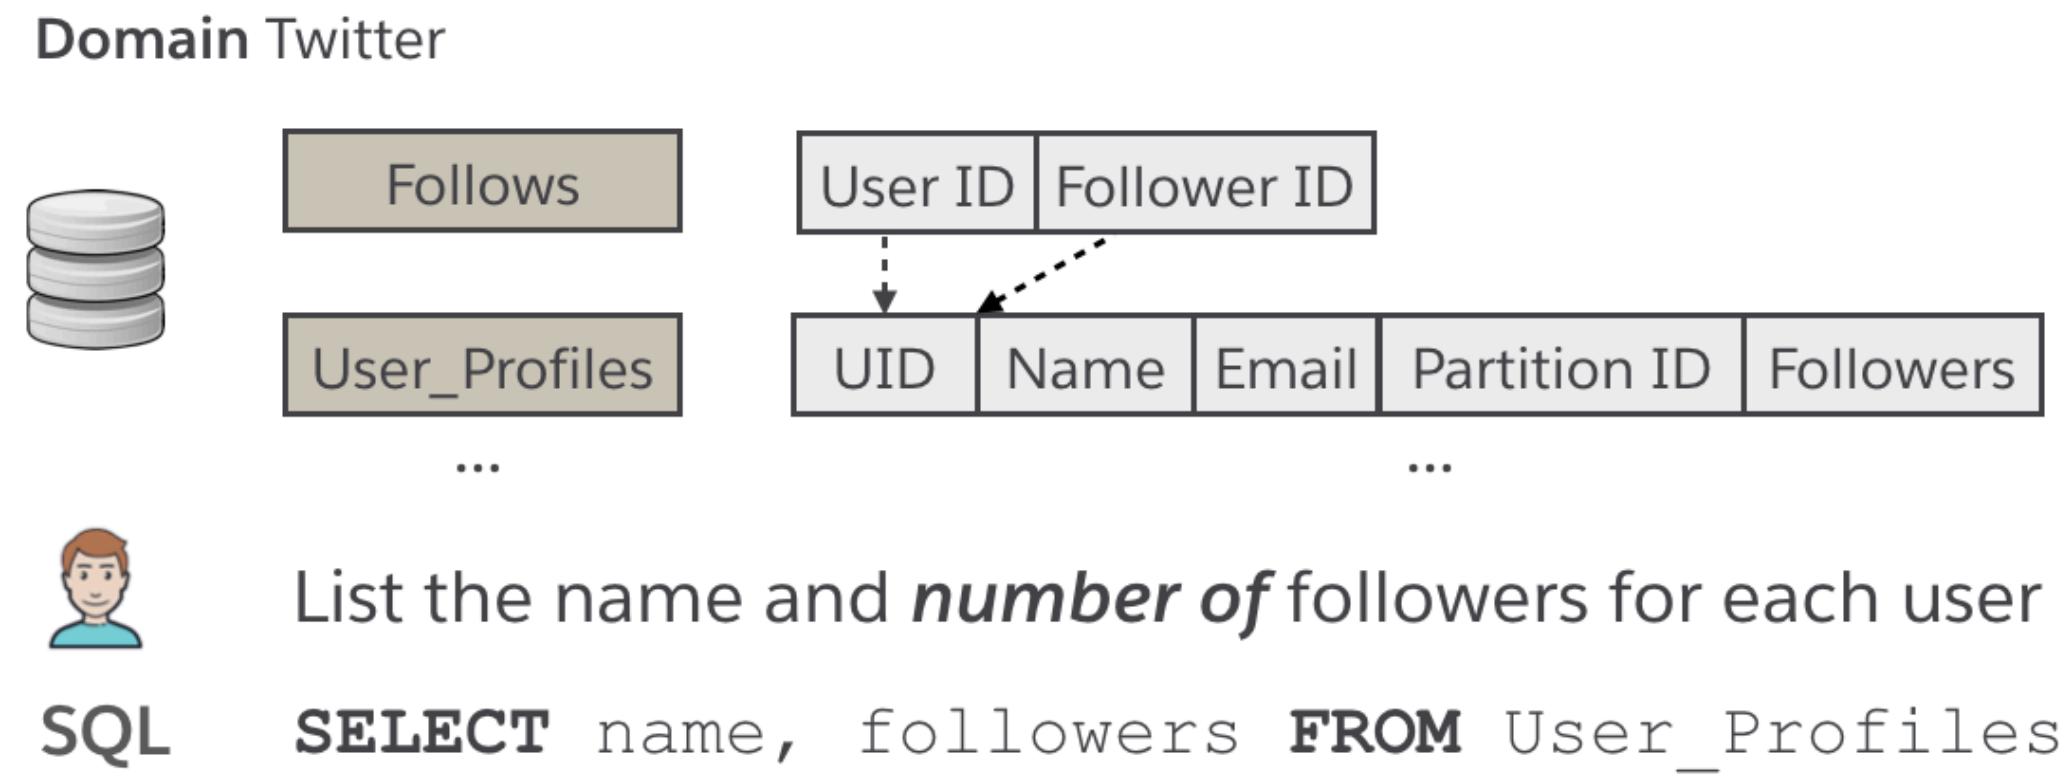
UserID	Name	Nationality	PartitionID	# Followers	...
103041	Smith, John	Canada	P10	3	...
103042	Hanks, Tom	United States	P11	16.6M	...
103043	Obama, Barack	United States	P11	127M	...
...

Primary Key
↑

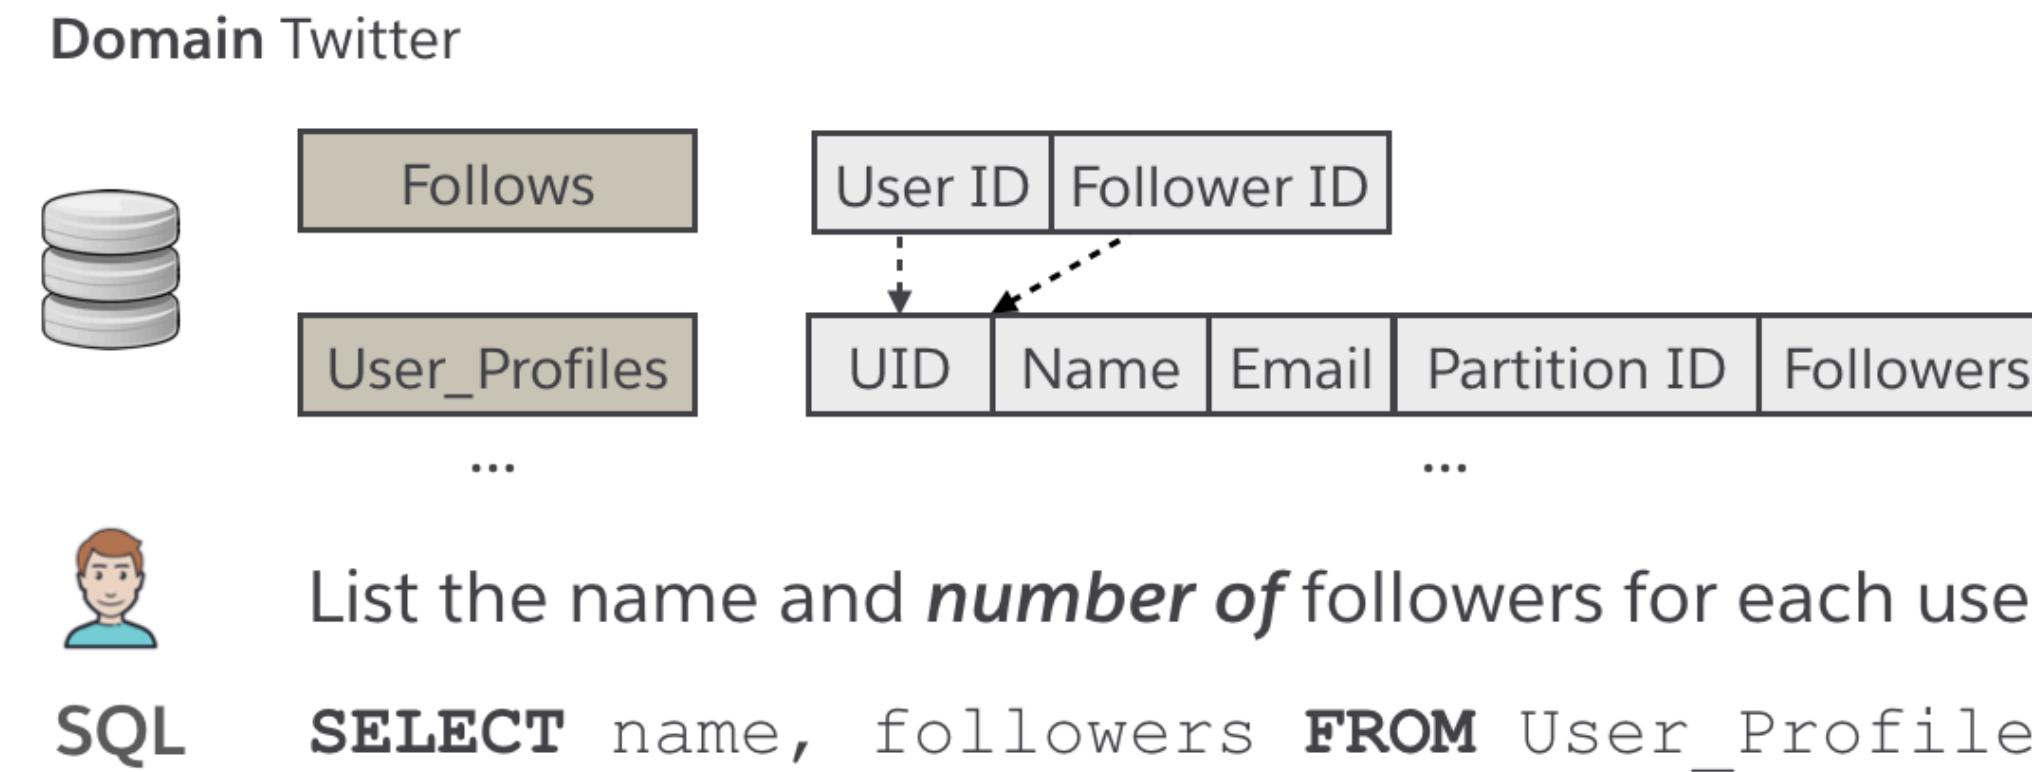
Relational Database Terms



Text-to-SQL Semantic Parsing

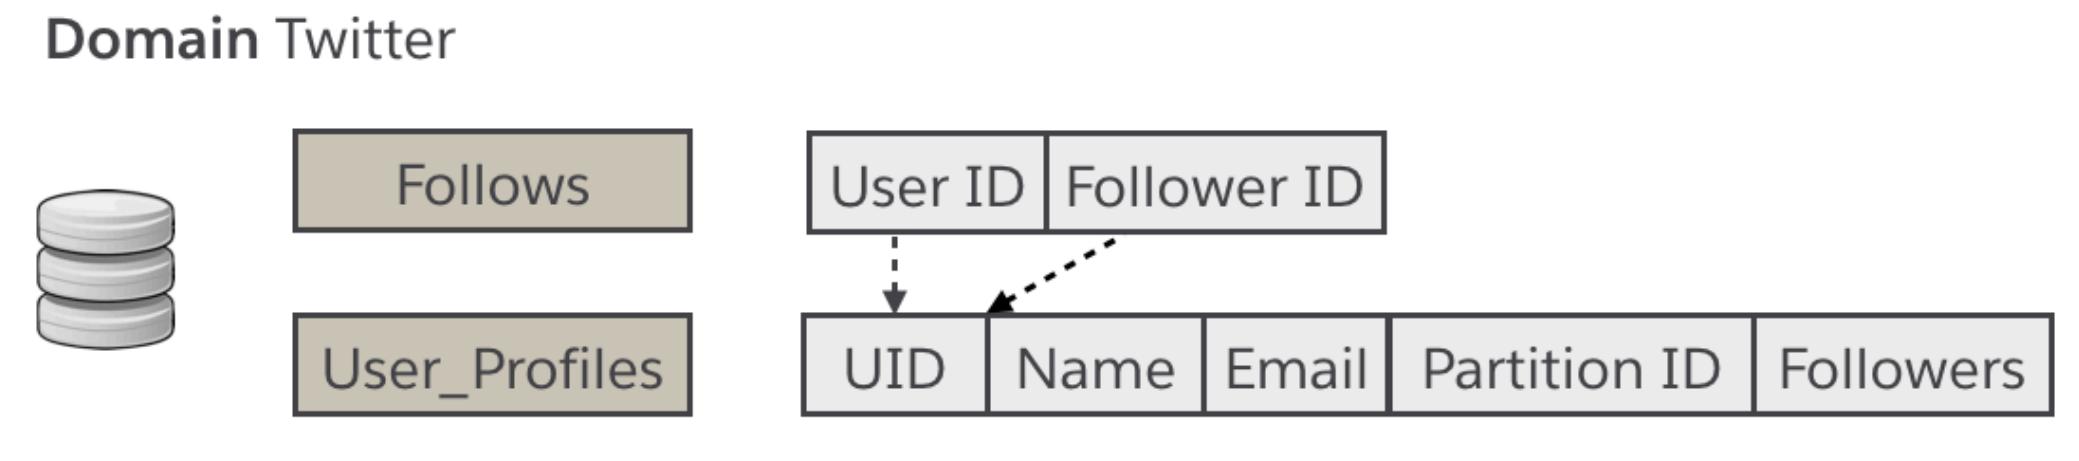


Text-to-SQL Semantic Parsing



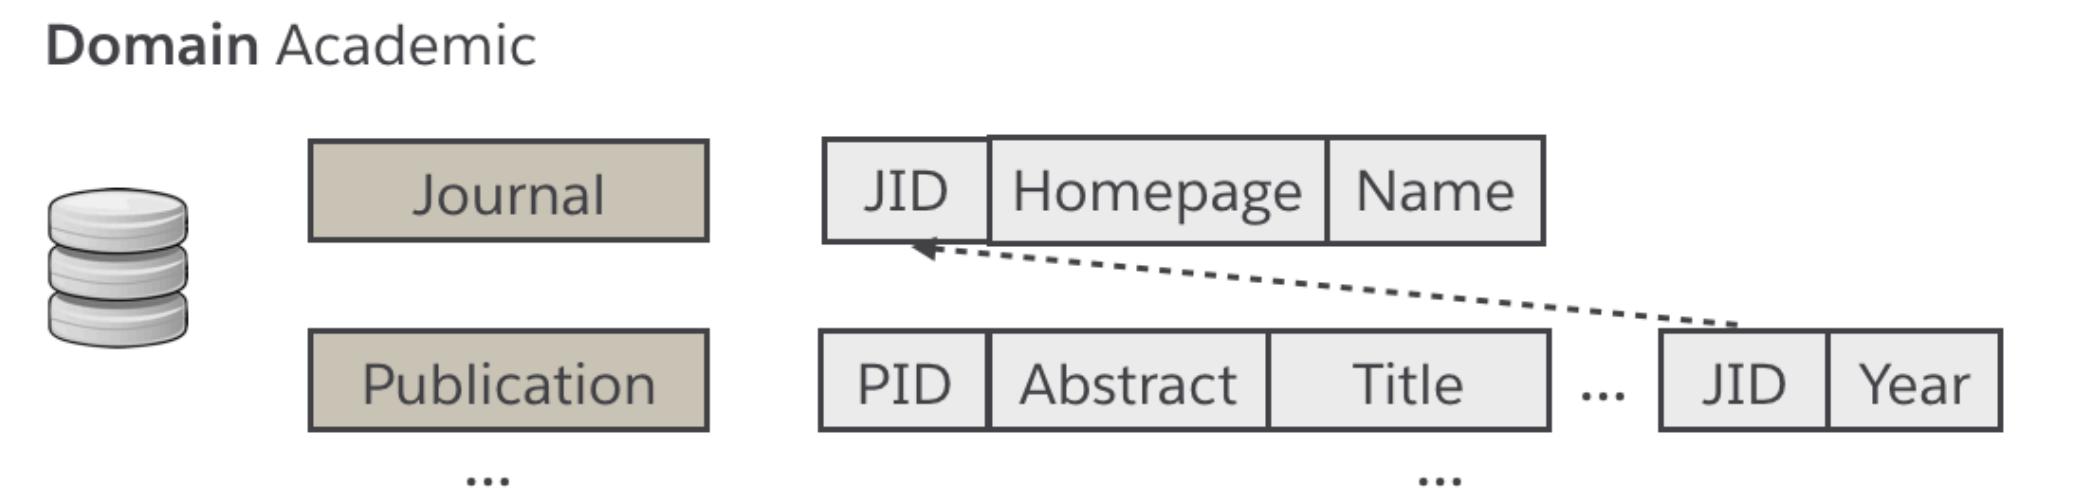
Strong/Full
supervision

Cross Domain Text-to-SQL Semantic Parsing



>List the name and *number of* followers for each user

SQL `SELECT name, followers FROM User_Profiles`



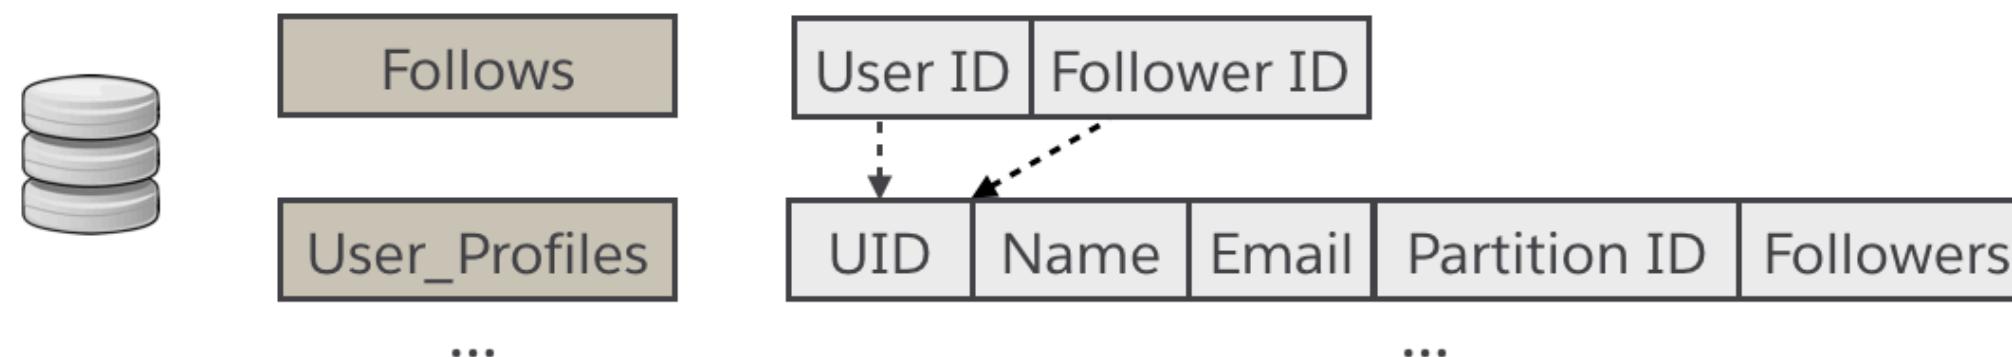
Return me the *number of* papers on PVLDB

SQL `SELECT COUNT(DISTINCT t2.title)
FROM Publication AS T2 JOIN Journal AS T1
ON T2.JID = T1.JID WHERE T1.name = "PVLDB"`

Challenges

Challenge 1: Questions with similar intent may map to very different SQL logical forms when issued to different DBs.

Domain Twitter



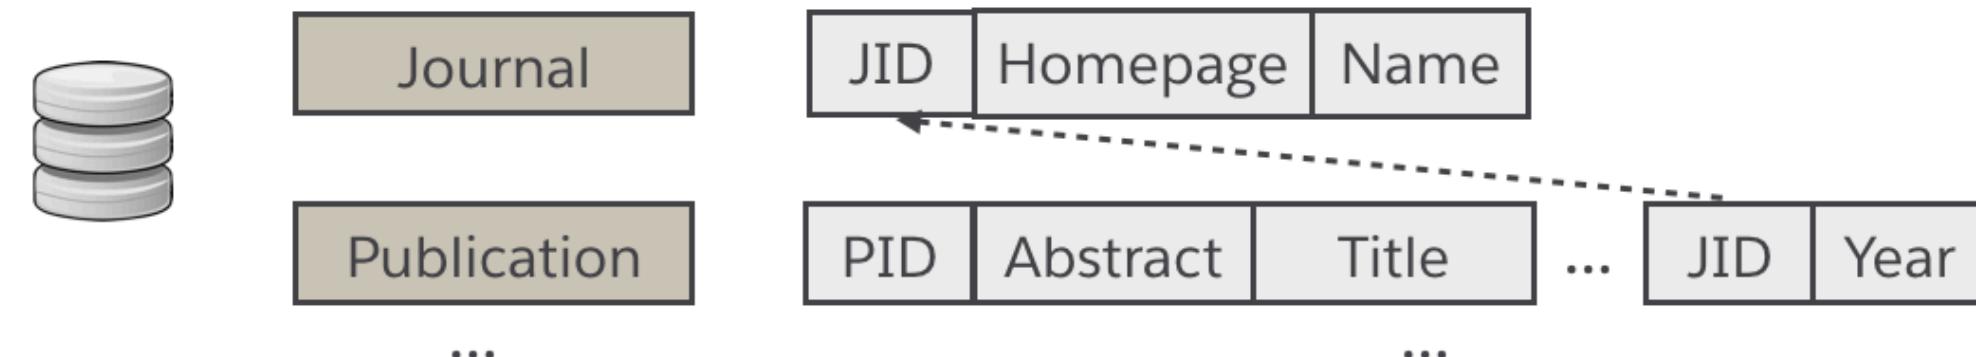
List the name and *number of* followers for each user

SQL

```
SELECT name, followers FROM User_Profiles
```

Cross-
Database

Domain Academic



Return me the *number of* papers on PVLDB

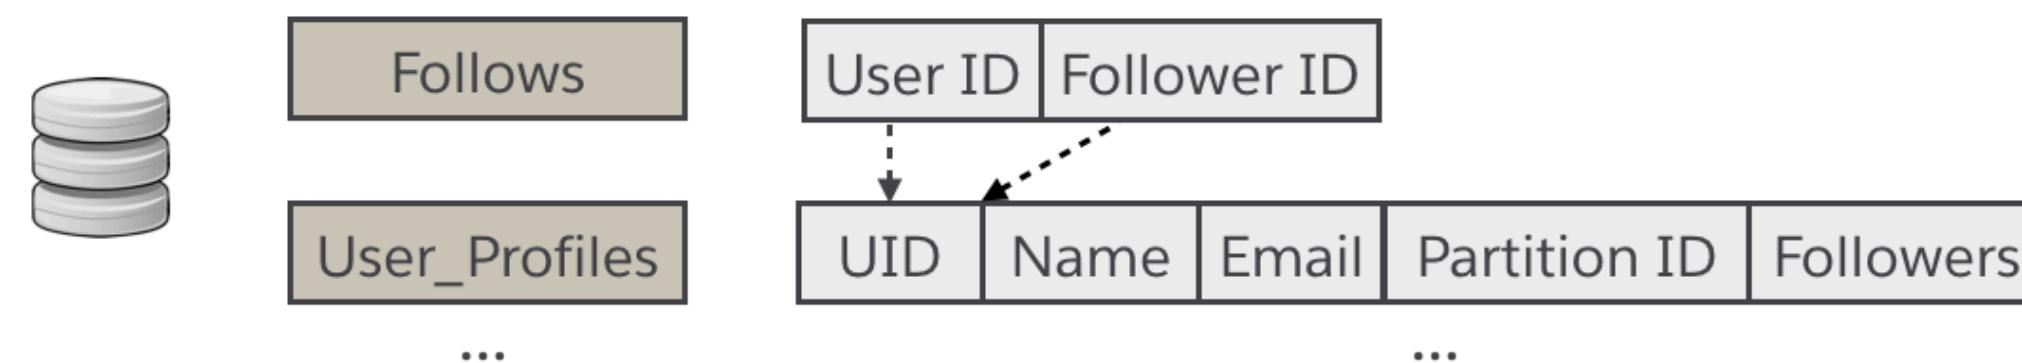
SQL

```
SELECT COUNT(DISTINCT t2.title)
FROM Publication AS T2 JOIN Journal AS T1
ON T2.JID = T1.JID WHERE T1.name = "PVLDB"
```

Challenges

Challenge 2: Domain-specific entities are frequently encountered.

Domain Twitter

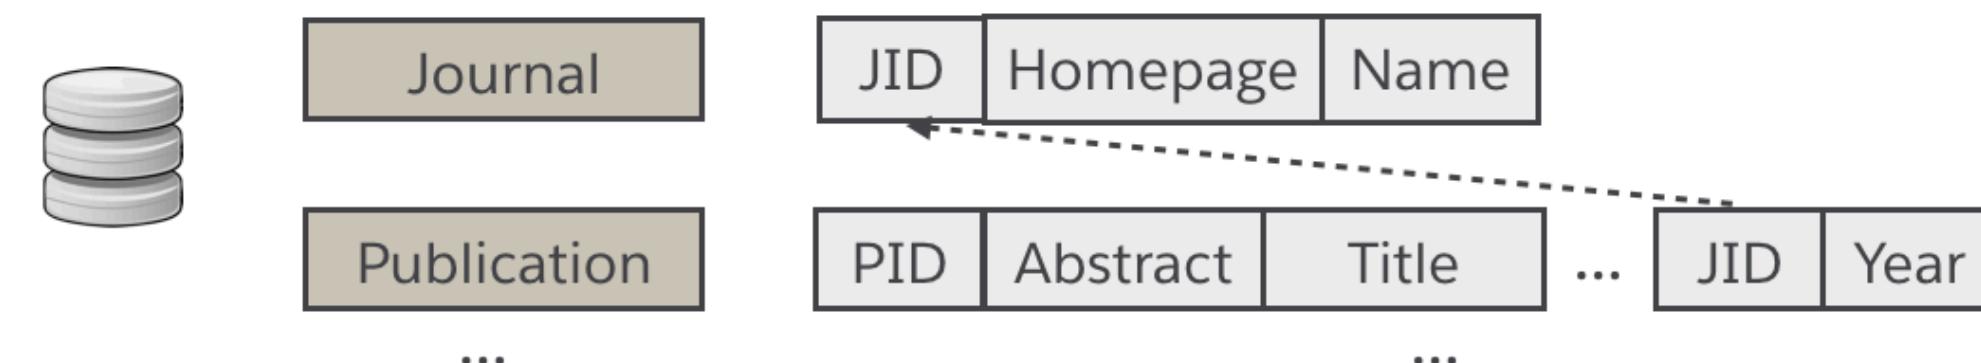


List the name and *number of* followers for each user

SQL

```
SELECT name, followers FROM User_Profiles
```

Domain Academic



Return me the *number of* papers on PVLDB

SQL

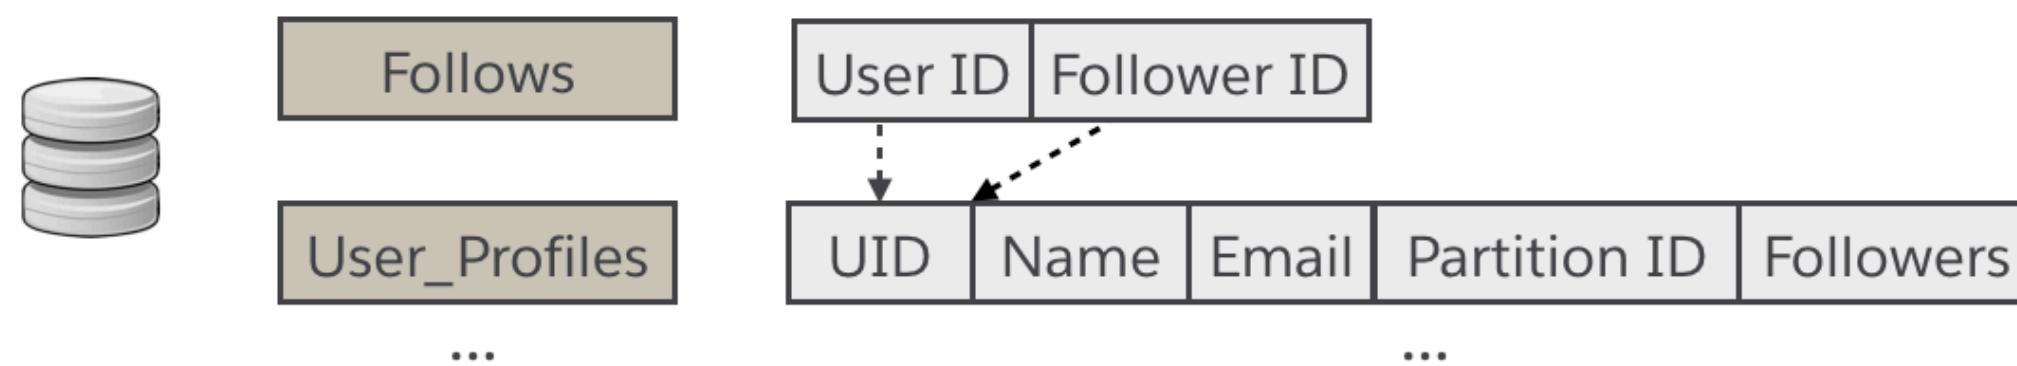
```
SELECT COUNT(DISTINCT t2.title)
FROM Publication AS T2 JOIN Journal AS T1
ON T2.JID = T1.JID WHERE T1.name = "PVLDB"
```

Domain Diversity

Observations

 Observation 1: It is important to contextualize the question and the database (DB), similar to the setup in machine reading comprehension

Domain Twitter

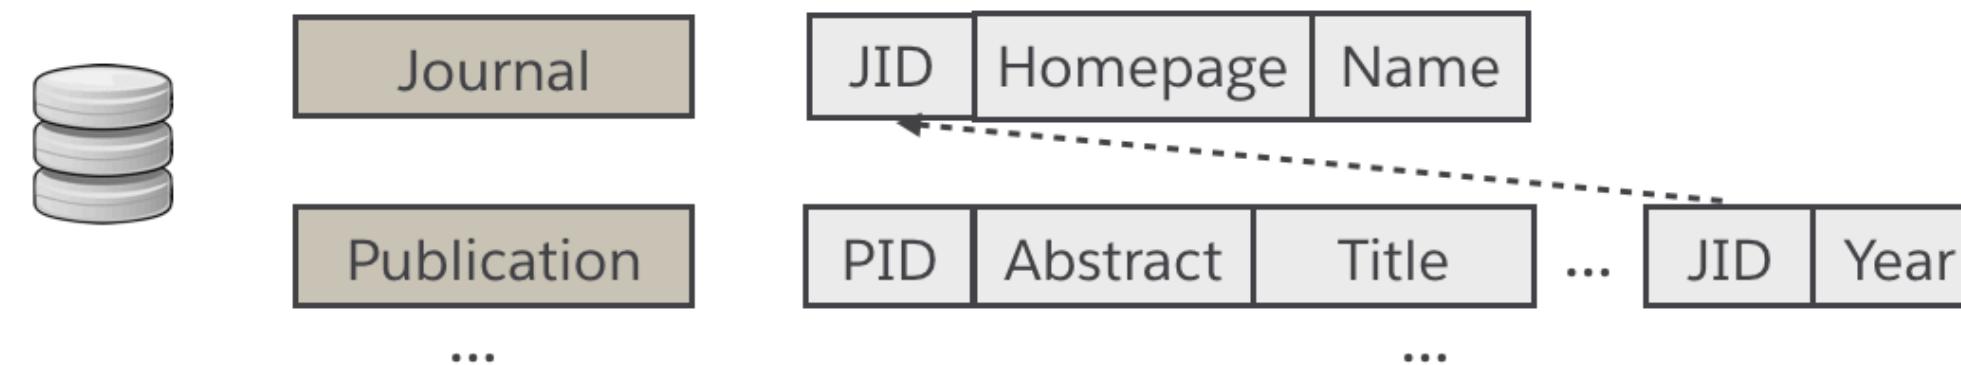


List the name and *number of* followers for each user

SQL

```
SELECT name, followers FROM User_Profiles
```

Domain Academic



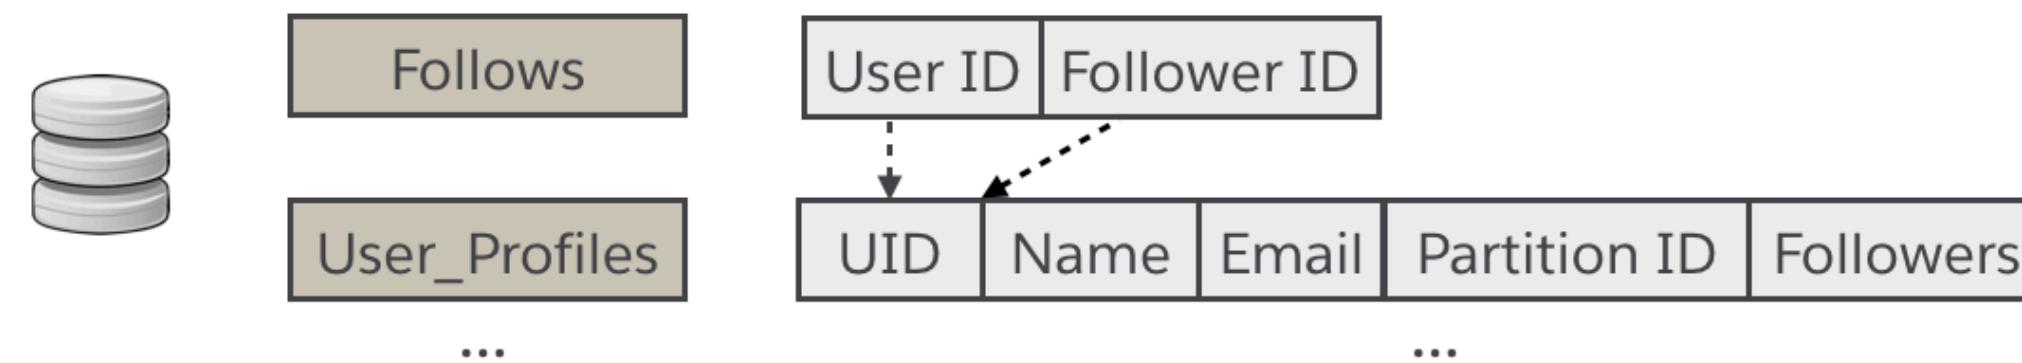
Return me the *number of* papers on PVLDB

SQL

```
SELECT COUNT(DISTINCT t2.title)
FROM Publication AS T2 JOIN Journal AS T1
ON T2.JID = T1.JID WHERE T1.name = "PVLDB"
```

Observations

Domain Twitter

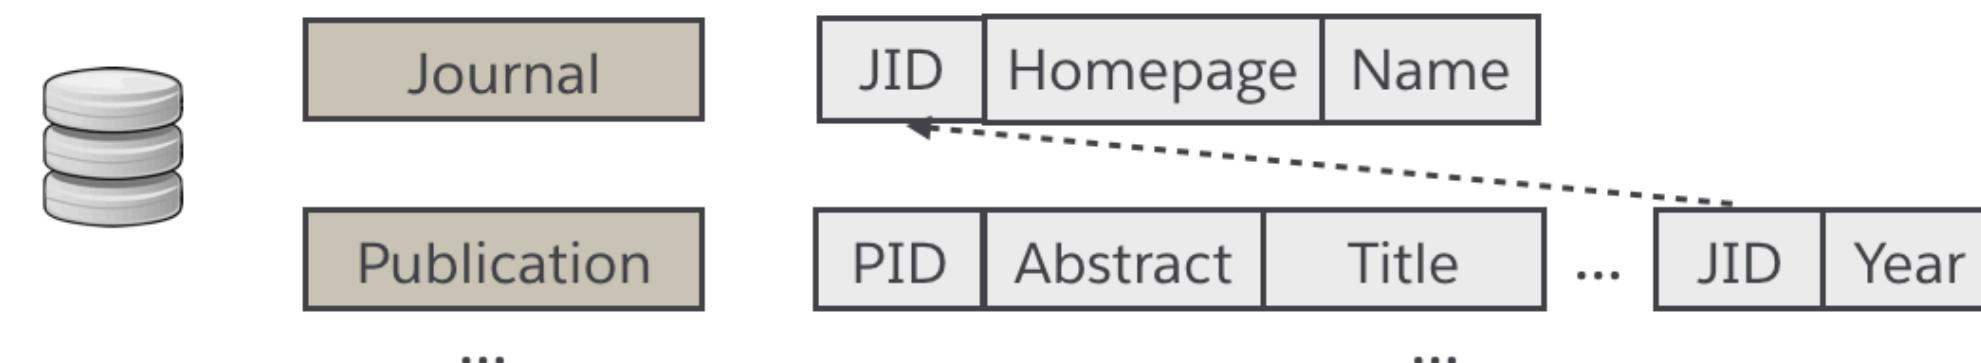


List the name and *number of* followers for each user

SQL

```
SELECT name, followers FROM User_Profiles
```

Domain Academic



Return me the *number of* papers on PVLDB

SQL

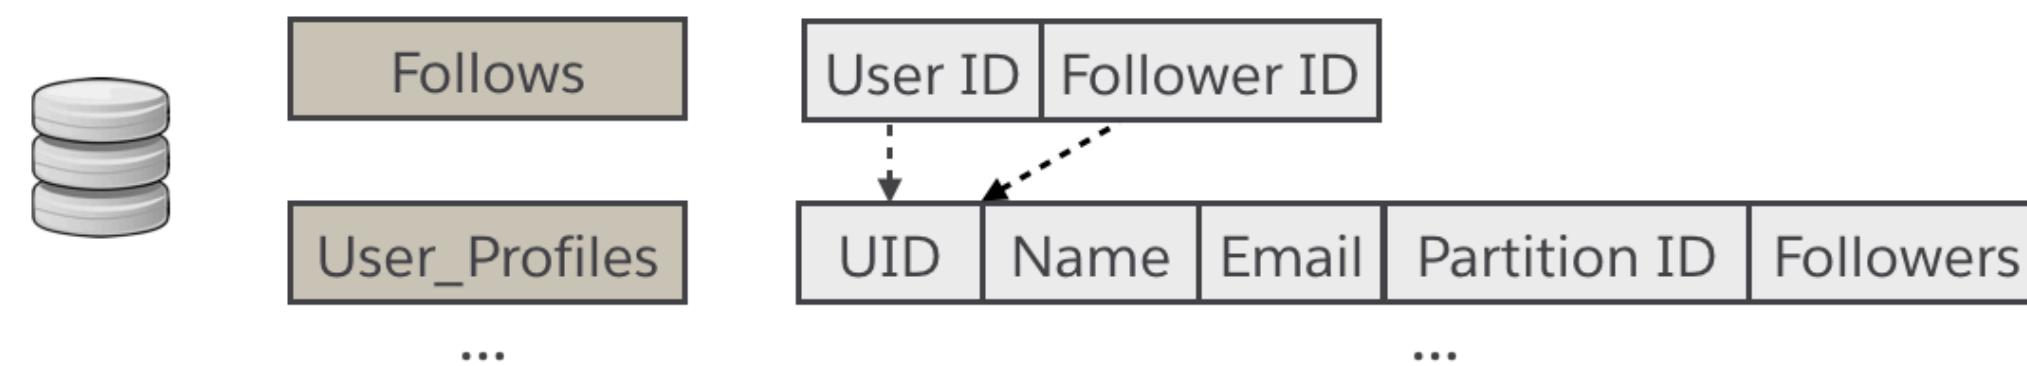
```
SELECT COUNT(DISTINCT t2.title)
FROM Publication AS T2 JOIN Journal AS T1
ON T2.JID = T1.JID WHERE T1.name = "PVLDB"
```



Observation 2: Database understanding should take into account both the DB schema and the DB content

Observations

Domain Twitter

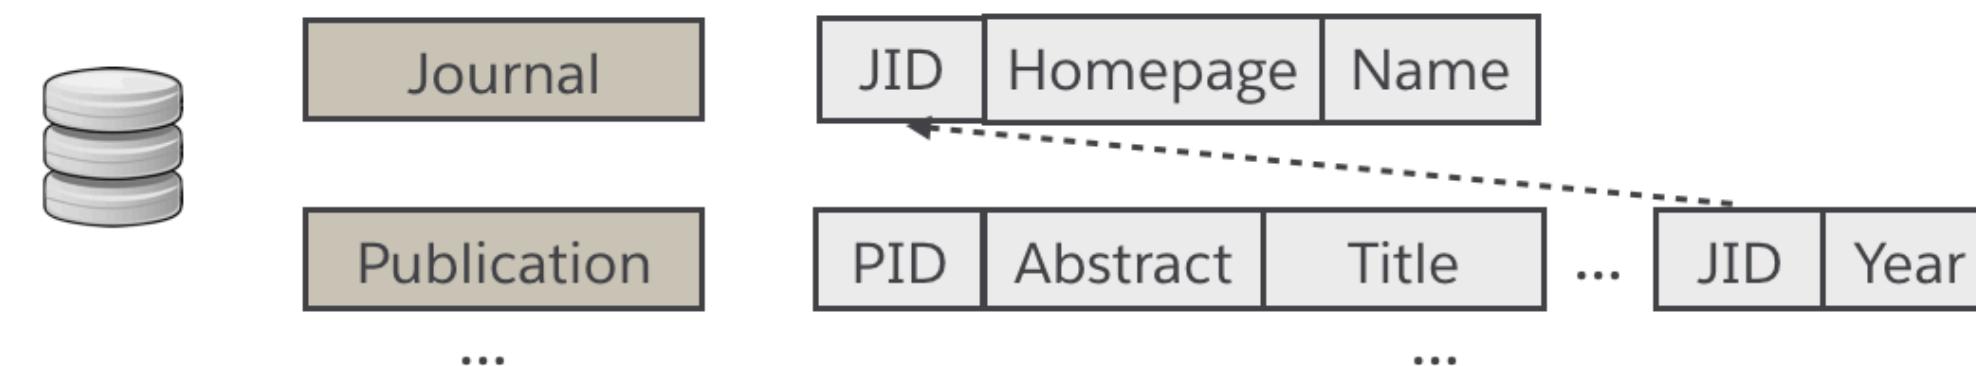


List the name and *number of* followers for each user

SQL

```
SELECT name, followers FROM User_Profiles
```

Domain Academic



Return me the *number of* papers on PVLDB

SQL

```
SELECT COUNT(DISTINCT t2.title)
FROM Publication AS T2 JOIN Journal AS T1
ON T2.JID = T1.JID WHERE T1.name = "PVLDB"
```



Observation 3: Most “rare entities” mentioned in the question correspond to tables, fields, or DB cells

Problem Setup

Question

Database Schema



Show names of properties that are either houses or apartments

Properties

Property id	Property type code	Property name	Date on market	Date sold	...
Apartment
Field
House
Shop
Other

Reference Property Types

Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

Picklists

Apartment
Field
House
Shop
Other

Problem Setup

Research Question:

How can we learn a representation that effectively captures the language grounding of an input question, the DB schema and the DB content?

Question

Database Schema



Show names of properties that are either houses or apartments

Properties					
Property id	Property type code	Property name	Date on market	Date sold	...
Apartment
Field
House
Shop
Other

Reference Property Types

Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

Picklists

Apartment	...
Field	...
House	...
Shop	...
Other	...

Joint Textual-Tabular Data Encoding

Serialize Table Header/DB Schema

Show names... Properties ... Property Type Code ... Reference Property Types ... Property Type Code ...

 Show names of properties that are either houses or apartments

Properties					
Property id	Property type code	Property name	Date on market	Date sold	...
Apartment
Field
House
Shop
Other

Reference Property Types	
Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

Joint Textual-Tabular Data Encoding

Serialize Table Header/DB Schema



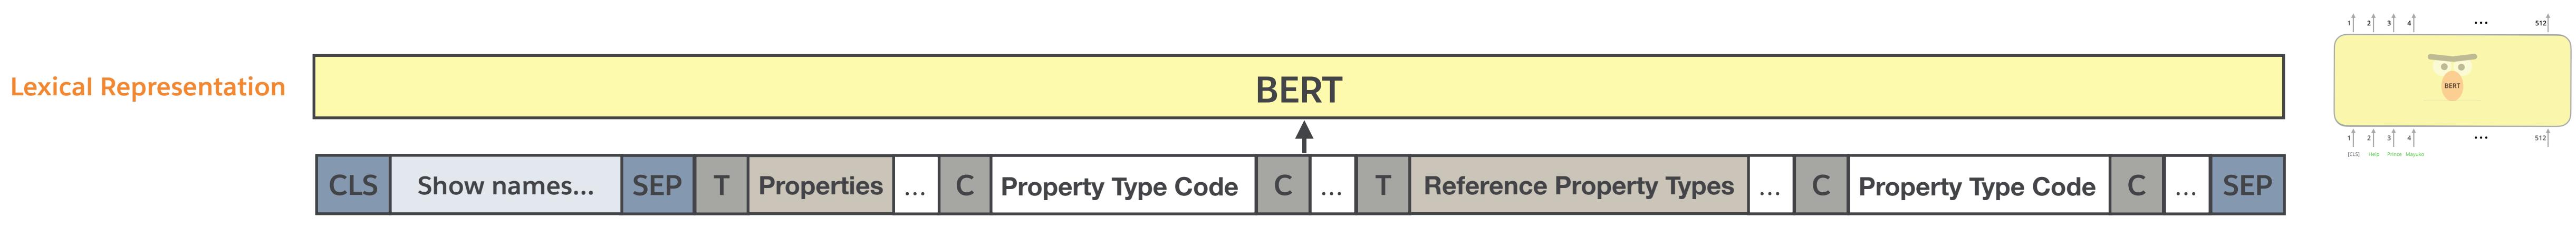
>Show names of properties that are either houses or apartments

Properties	Property id	Property type code	Property name	Date on market	Date sold	...
Apartment
Field
House
Shop
Other

Reference Property Types	Property type code	Property type description
Apartment
Field
House
Shop
Other

Joint Textual-Tabular Data Encoding

Serialize Table Header/DB Schema



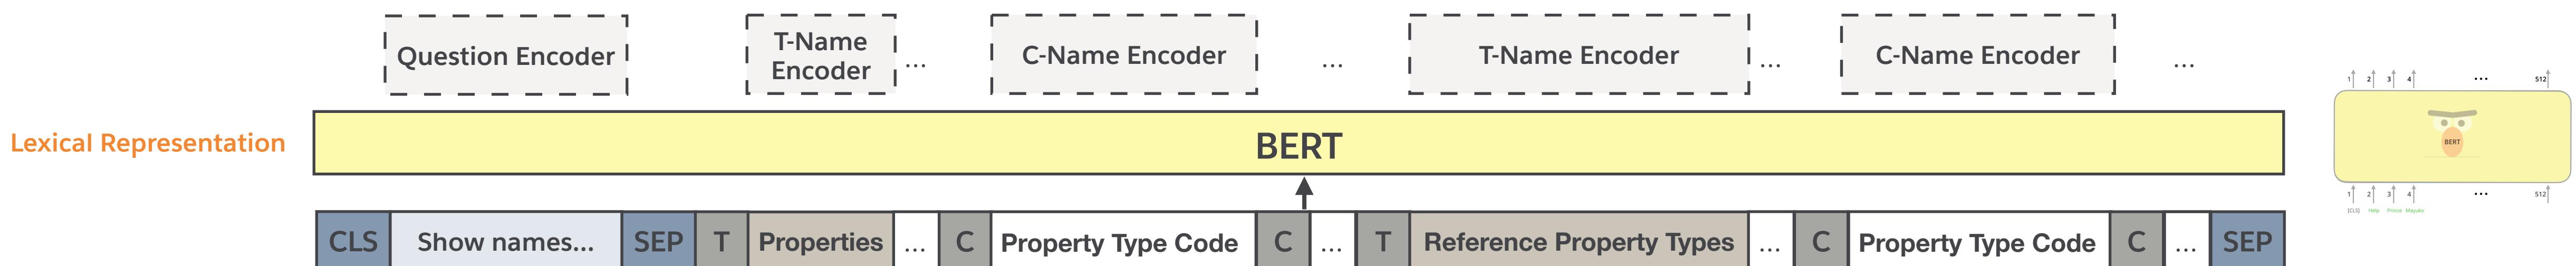
>Show names of properties that are either houses or apartments

Properties				
Property id	Property type code	Property name	Date on market	Date sold
Apartment
Field
House
Shop
Other

Reference Property Types	
Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

Joint Textual-Tabular Data Encoding

Component Encoding Layers



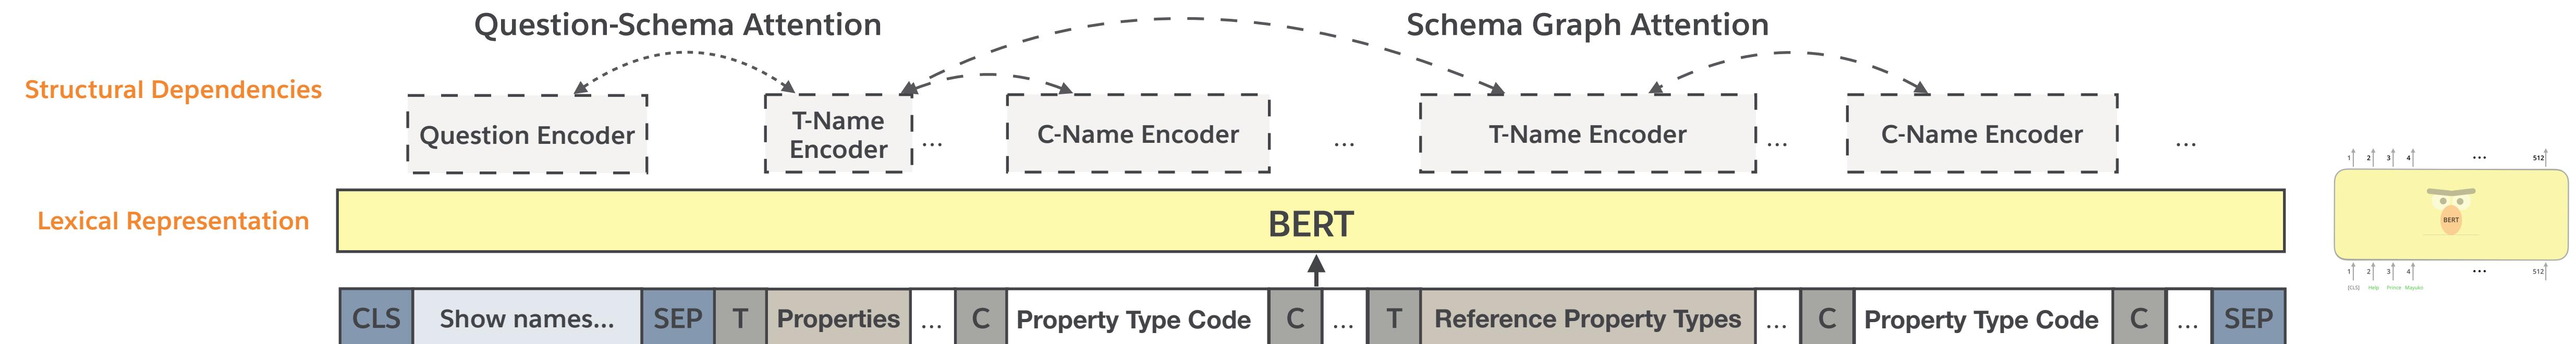
>Show names of properties that are either houses or apartments

Properties					
Property id	Property type code	Property name	Date on market	Date sold	...
Apartment
Field
House
Shop
Other

Reference Property Types	
Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

Joint Textual-Tabular Data Encoding

Attention Layers



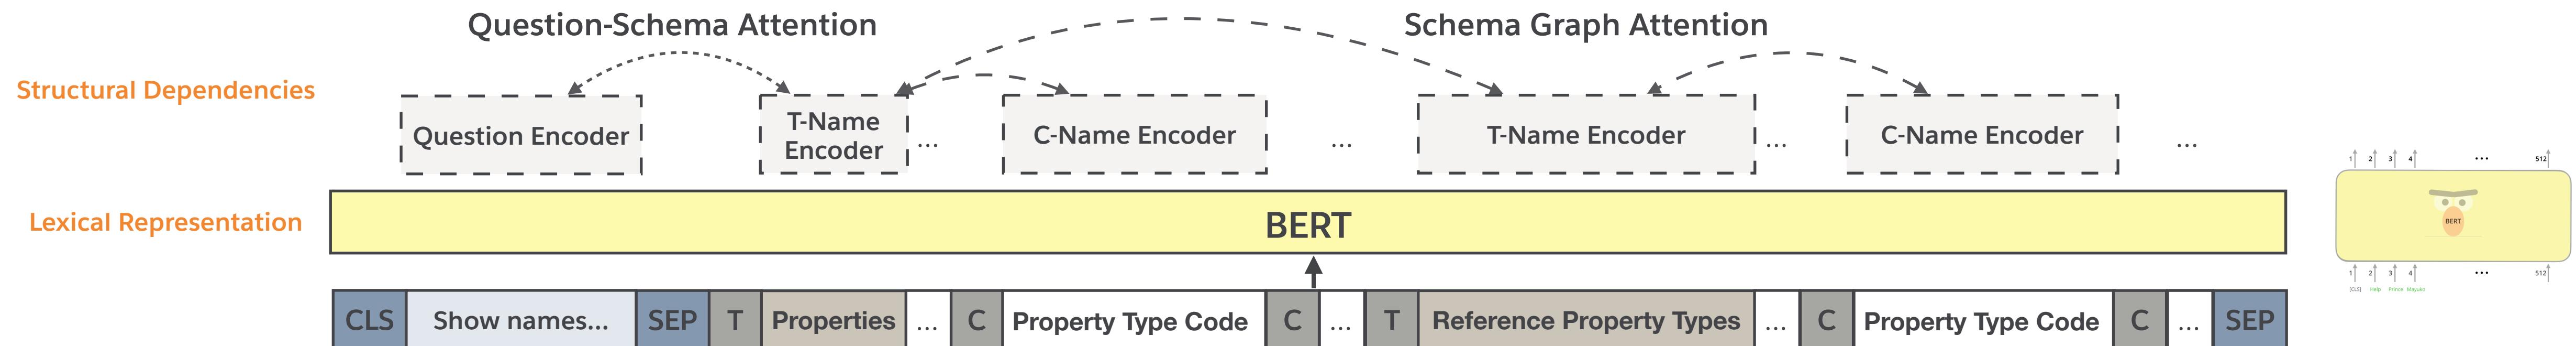
>Show names of properties that are either houses or apartments

Properties		Reference Property Types	
Property id	Property type code	Property name	Date on market
Apartment
Field
House
Shop
Other

Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

Joint Textual-Tabular Data Encoding

Attention Layers



>Show names of properties that are either houses or apartments

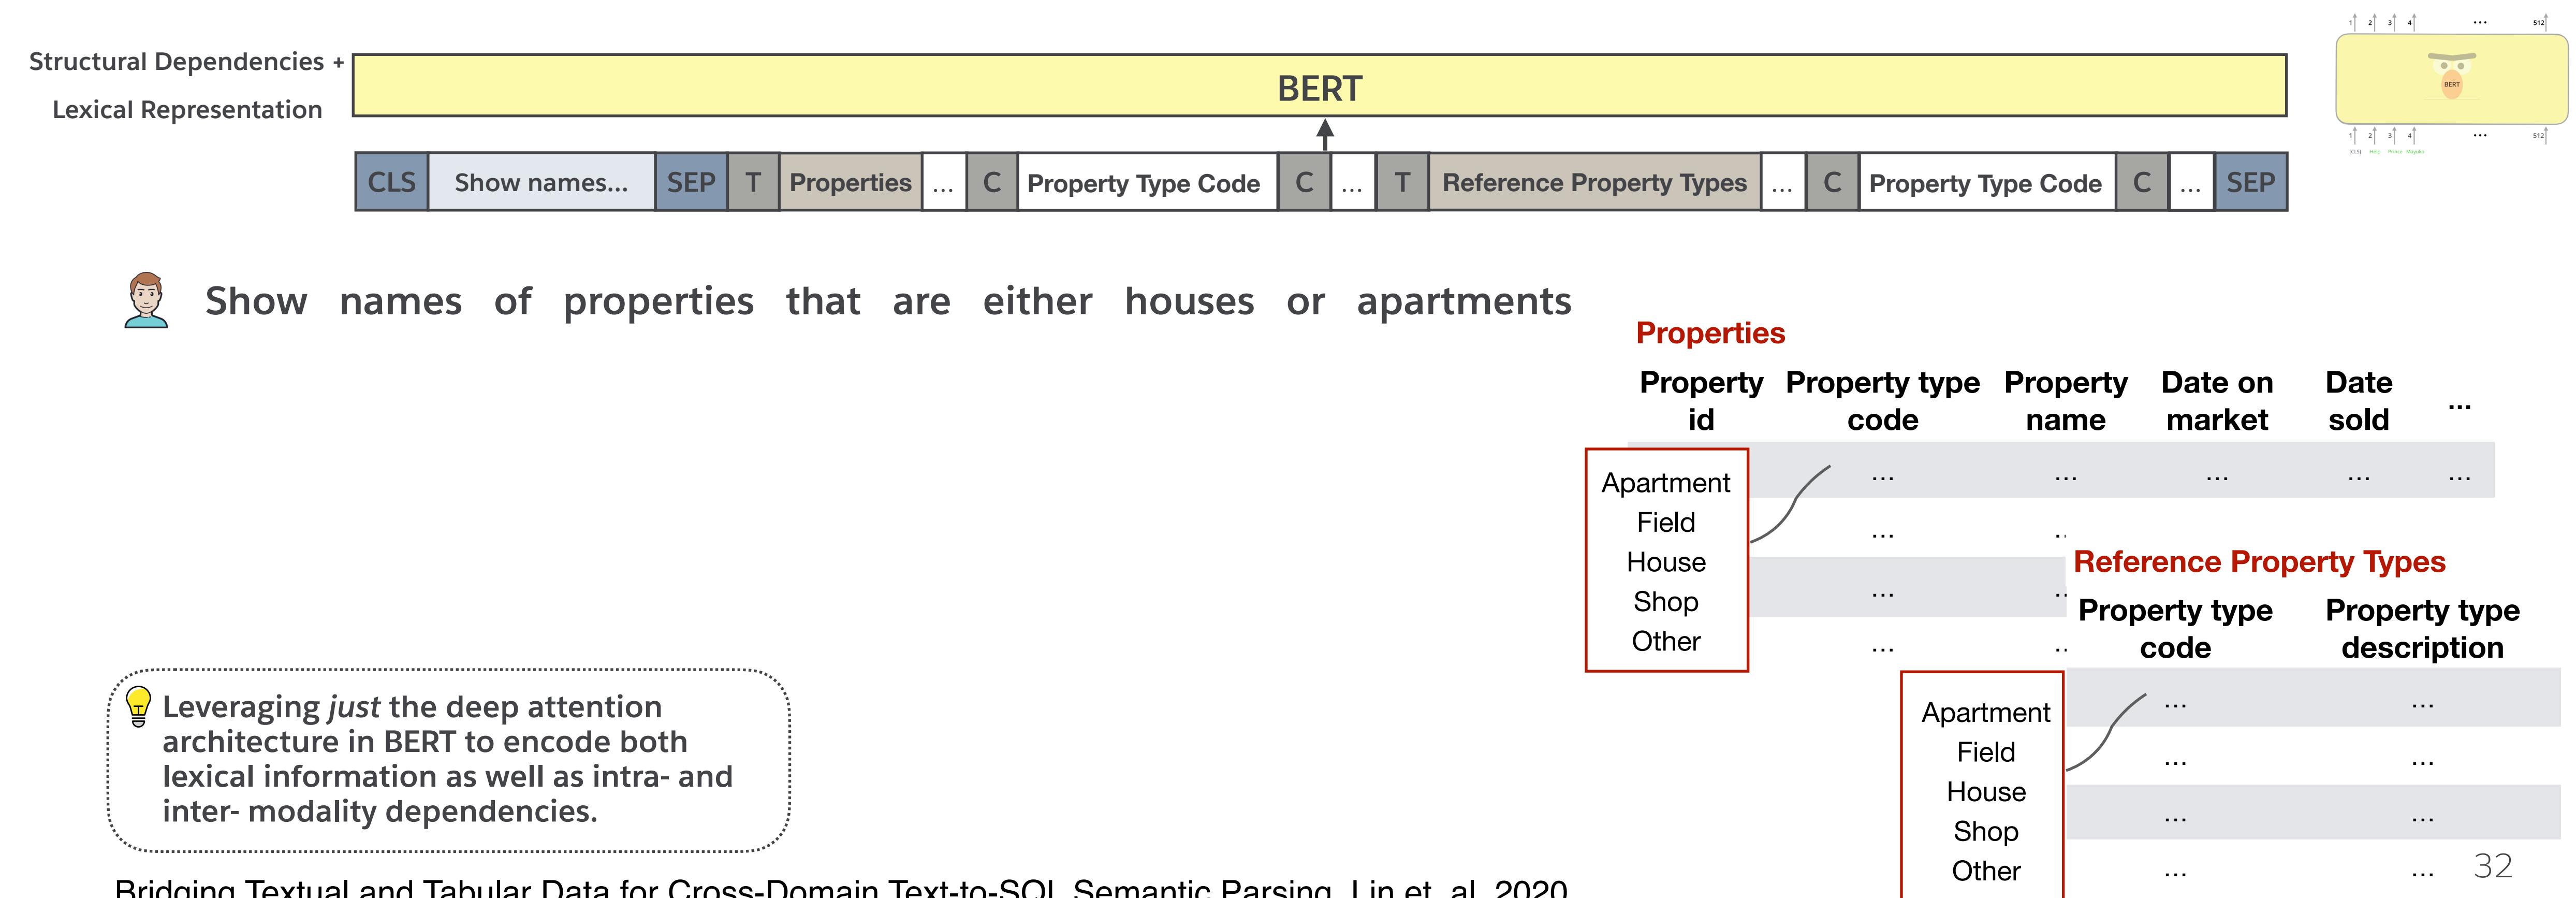
Properties					
Property id	Property type code	Property name	Date on market	Date sold	...
Apartment
Field
House
Shop
Other

Reference Property Types	
Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

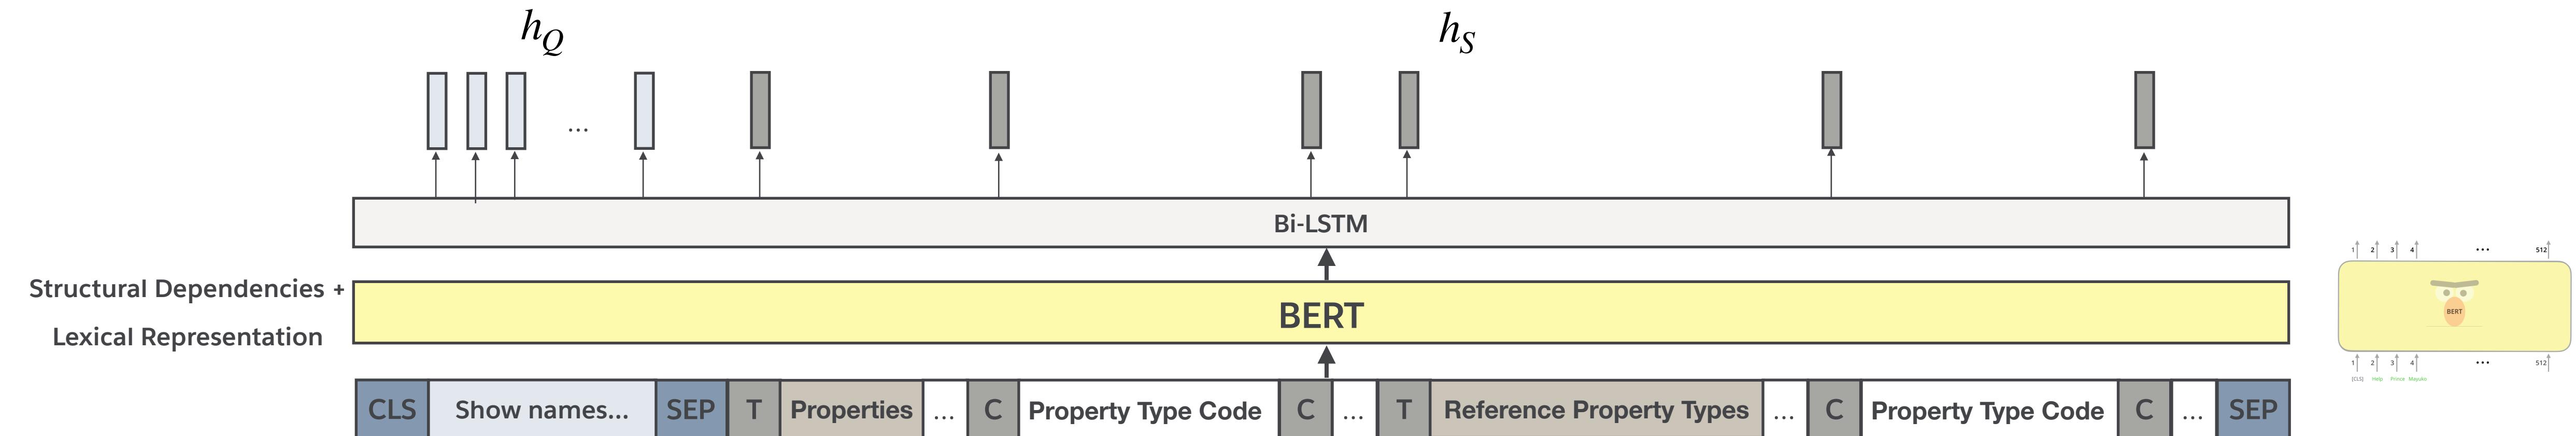
+6 Relational Self-Attention
Layers on top of BERT-large.
Model complexity quickly
increases.

Architecture redundancy

Joint Textual-Tabular Data Encoding



Joint Textual-Tabular Data Encoding



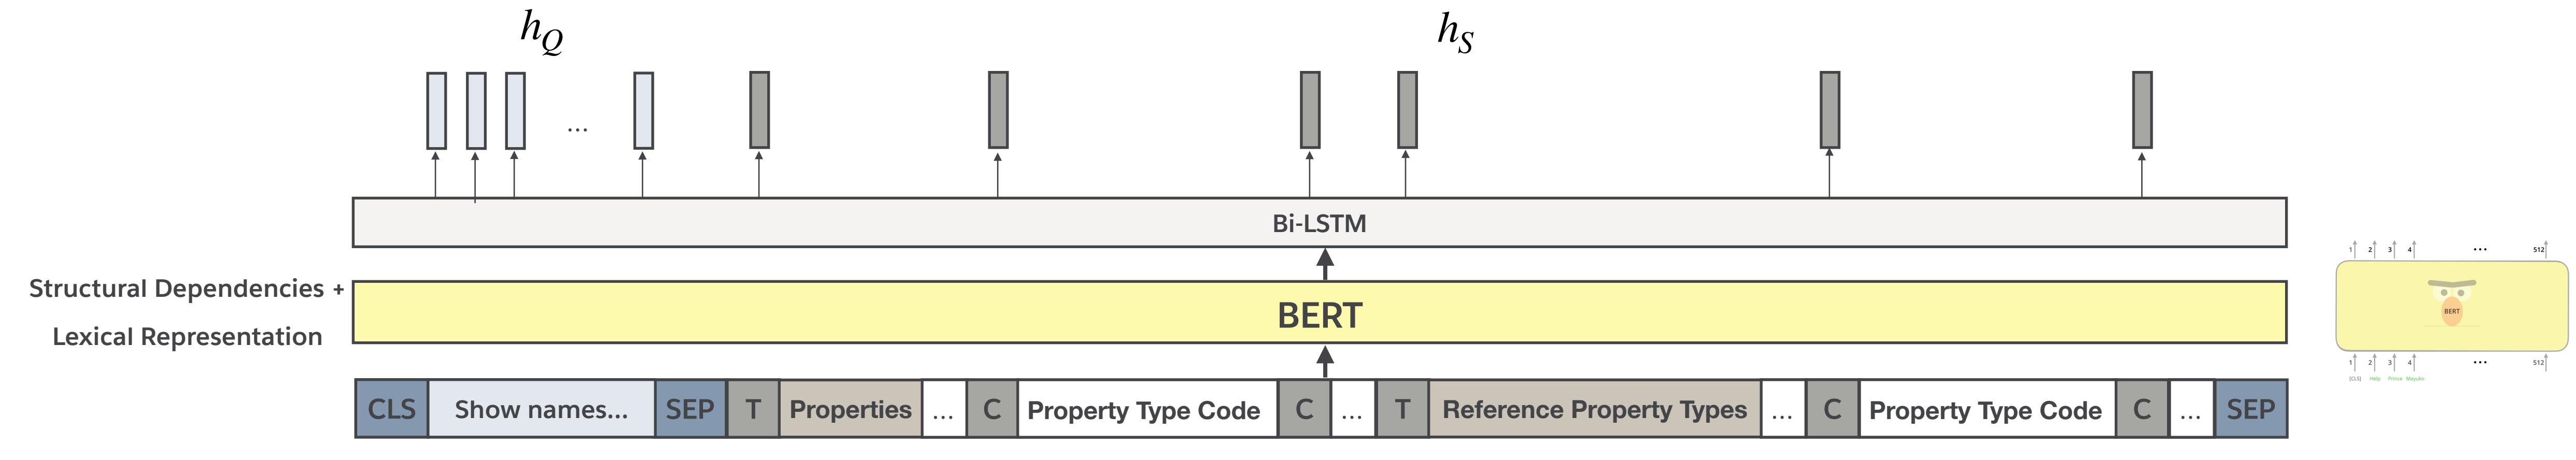
Show names of properties that are either houses or apartments

Properties					
Property id	Property type code	Property name	Date on market	Date sold	...
Apartment
Field
House
Shop
Other

Reference Property Types	
Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

Leveraging just the deep attention architecture in BERT to encode both lexical information as well as intra- and inter-modality dependencies.

Joint Textual-Tabular Data Encoding



Show names of properties that are either houses or apartments

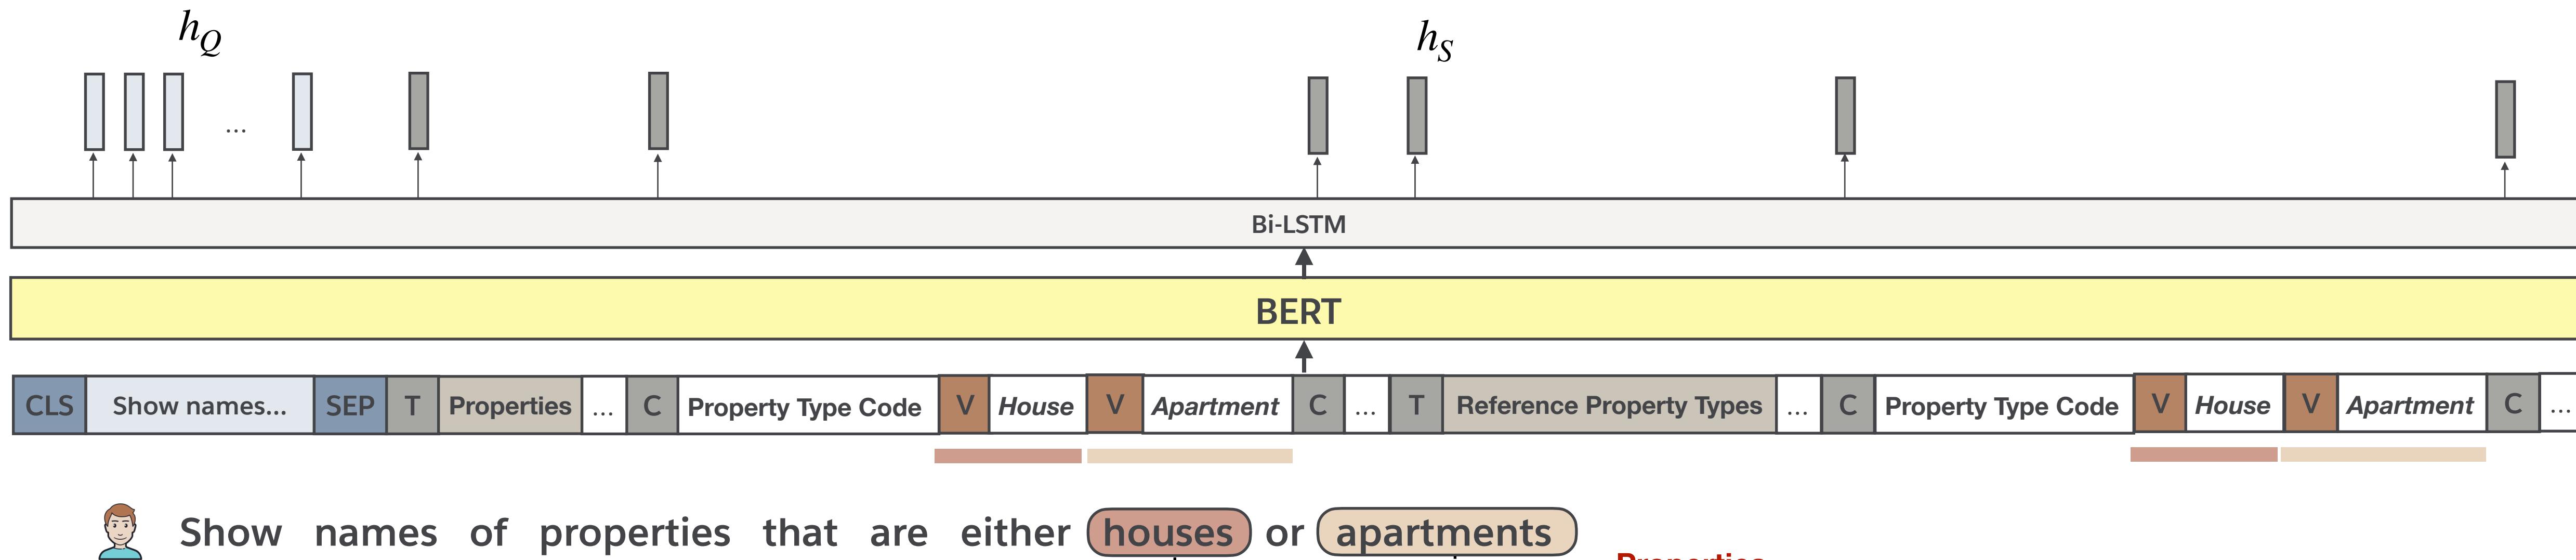
Properties

Property id	Property type code	Property name	Date on market	Date sold
Apartment
Field
House
Shop
Other

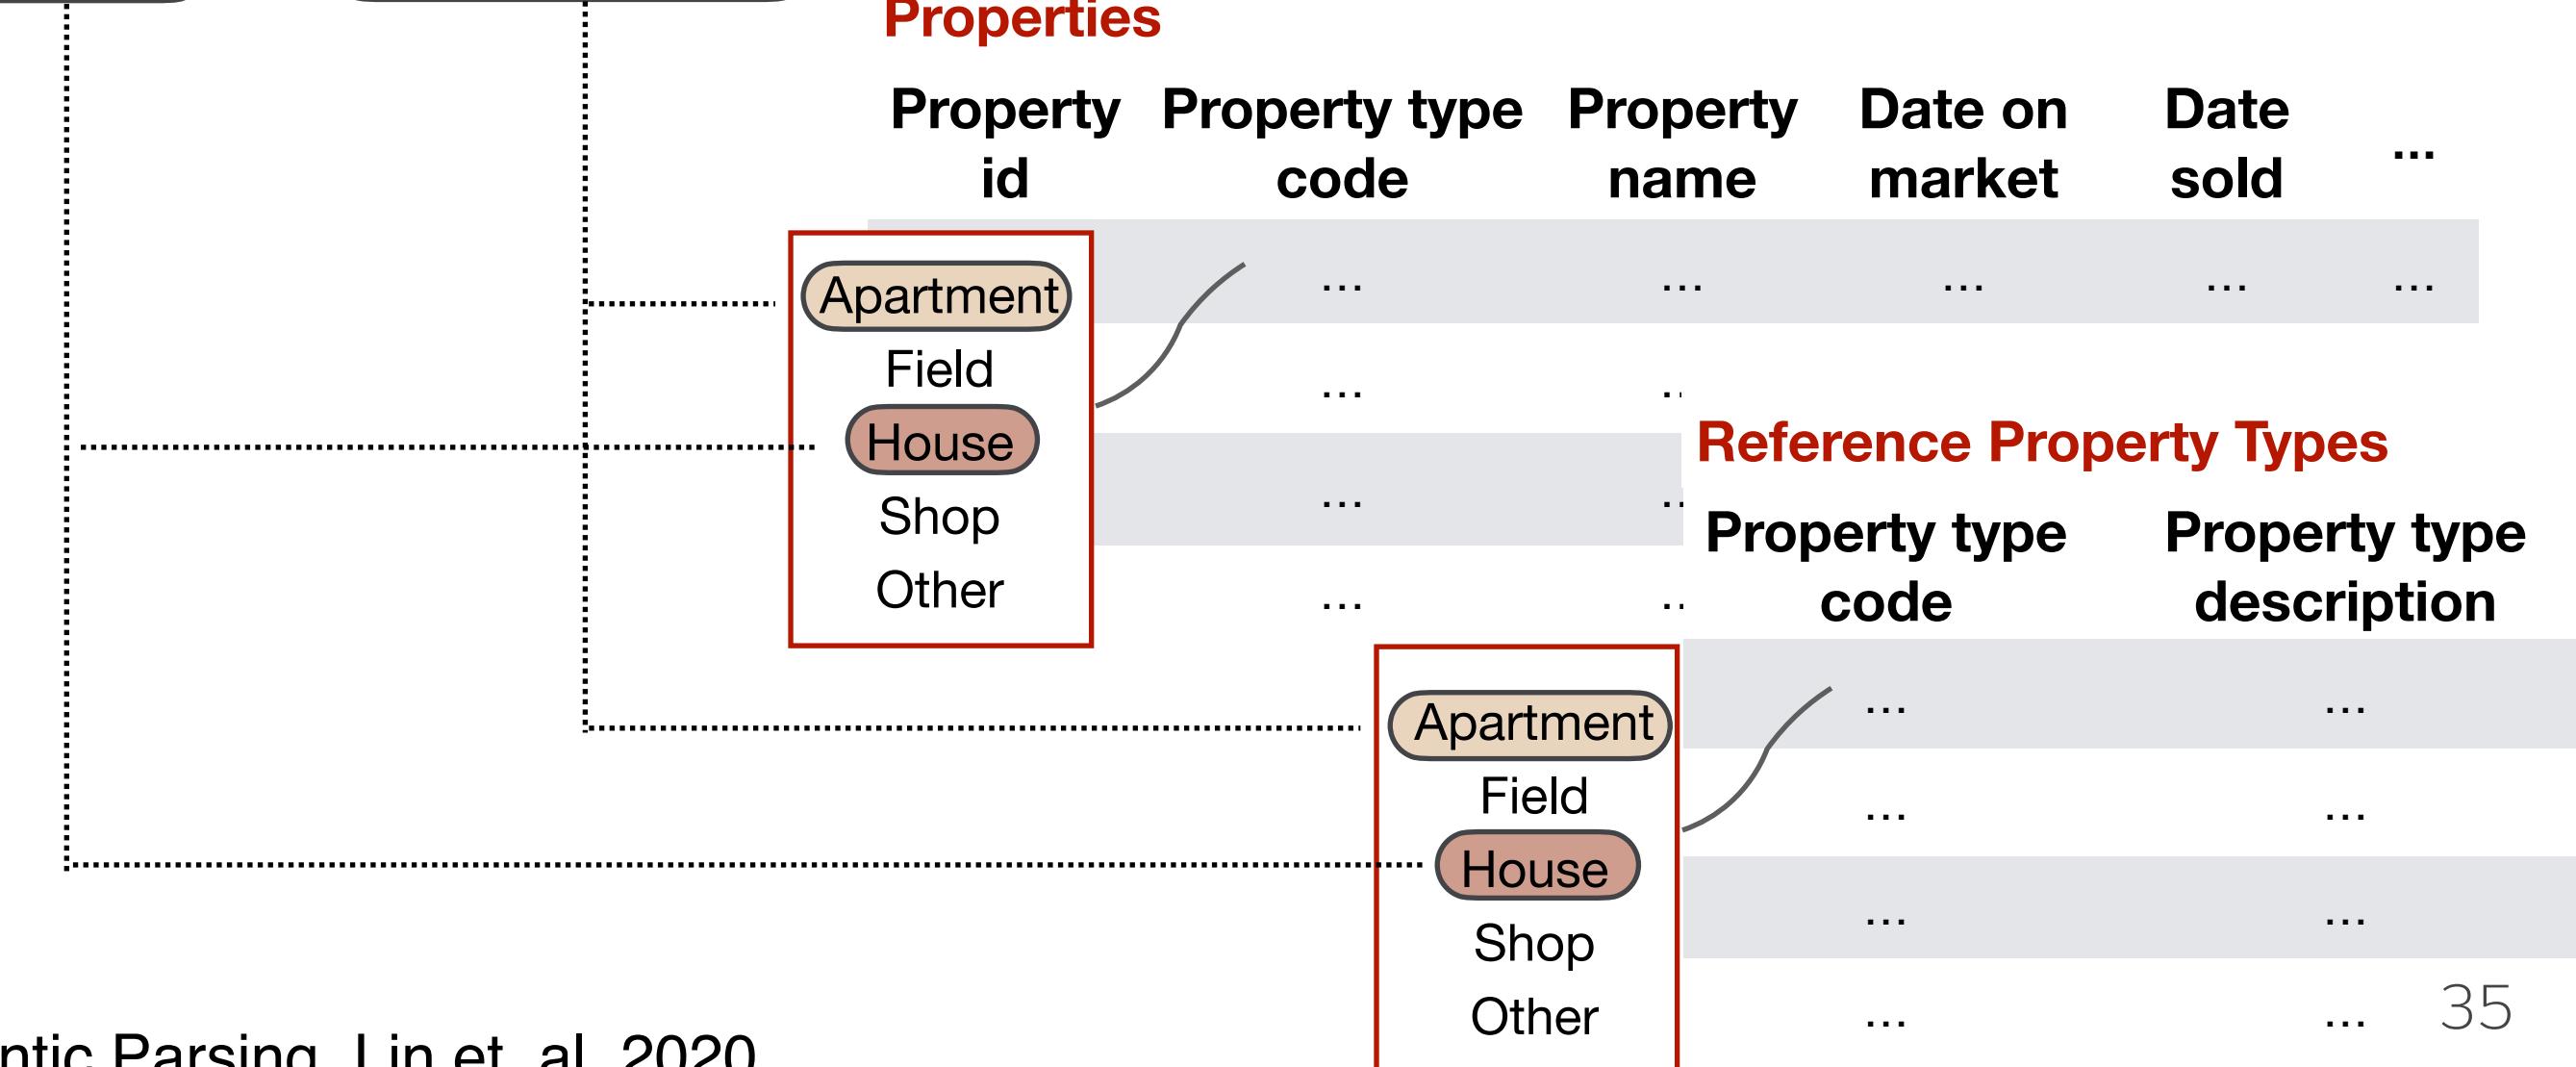
Reference Property Types

Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

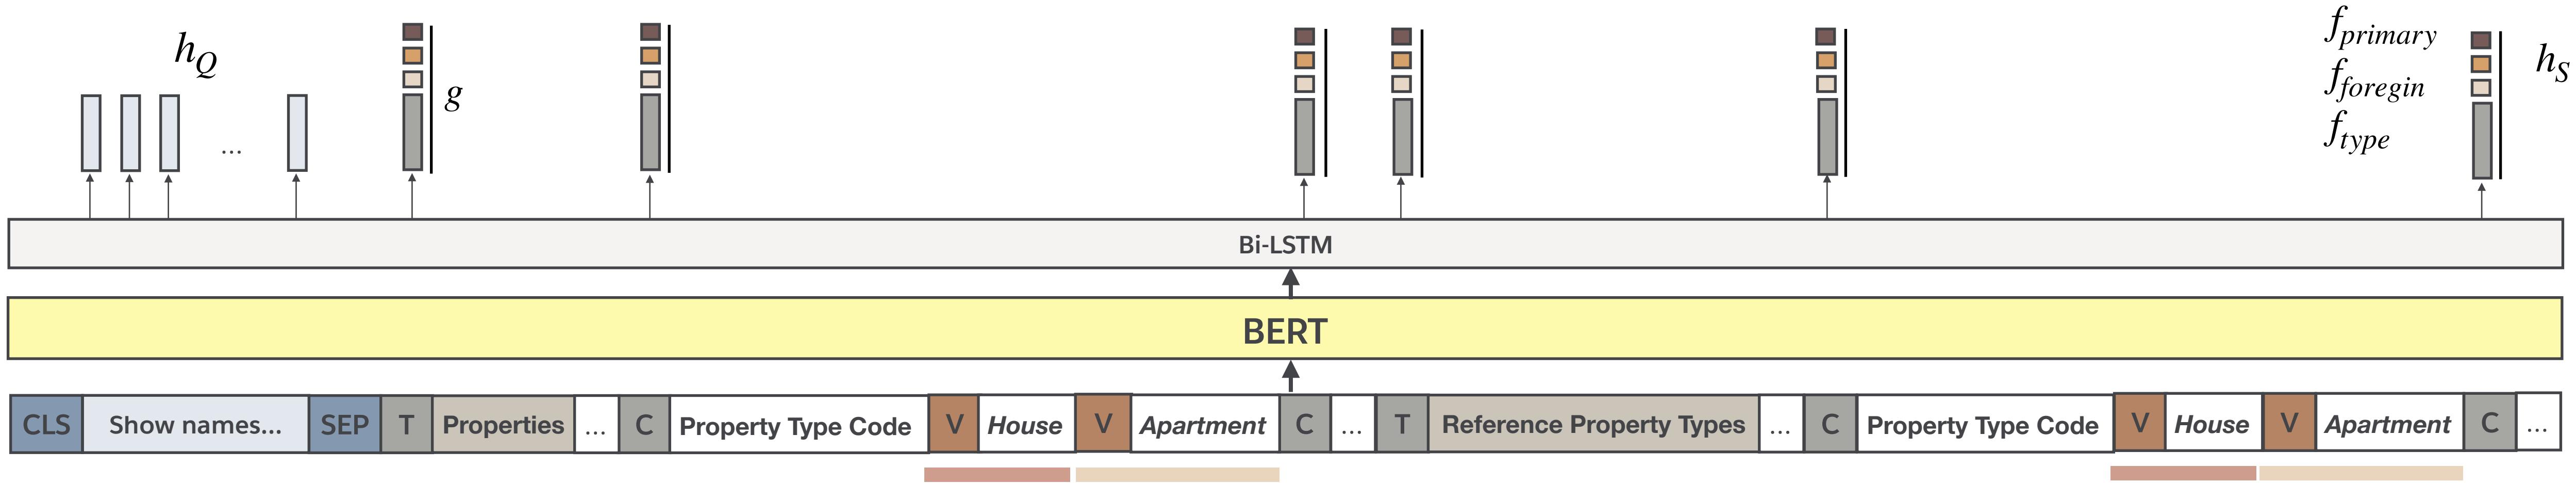
Bridging



Show names of properties that are either houses or apartments



Meta-Data Encoding



Show names of properties that are either houses or apartments

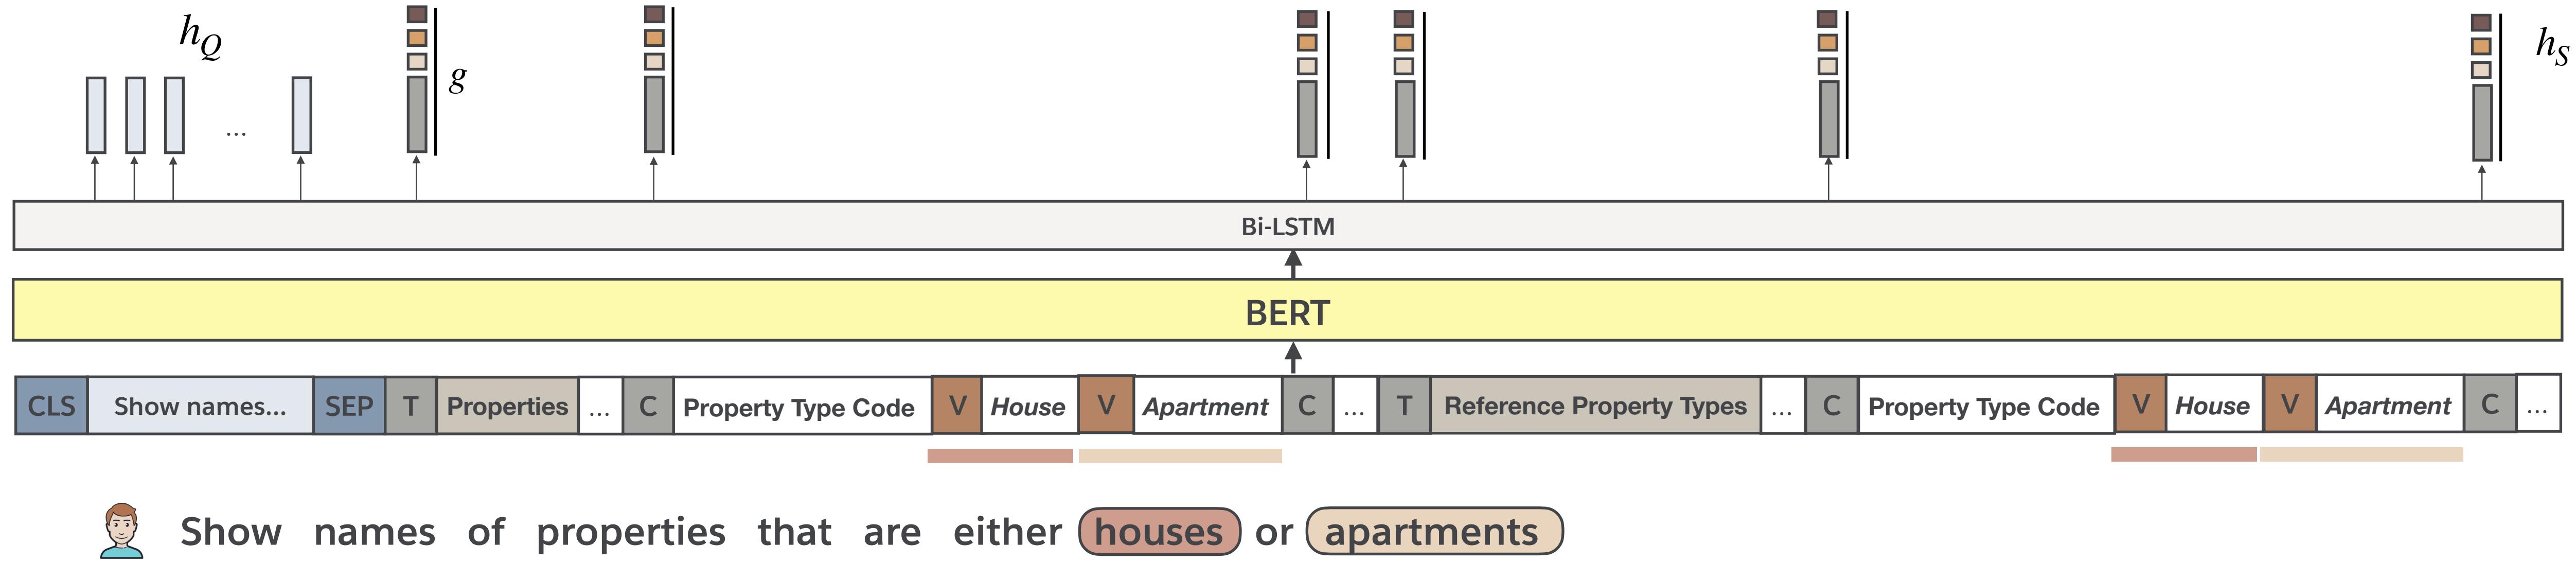
Properties

Property id	Property type code	Property name	Date on market	Date sold	...
Apartment
Field
House
Shop
Other

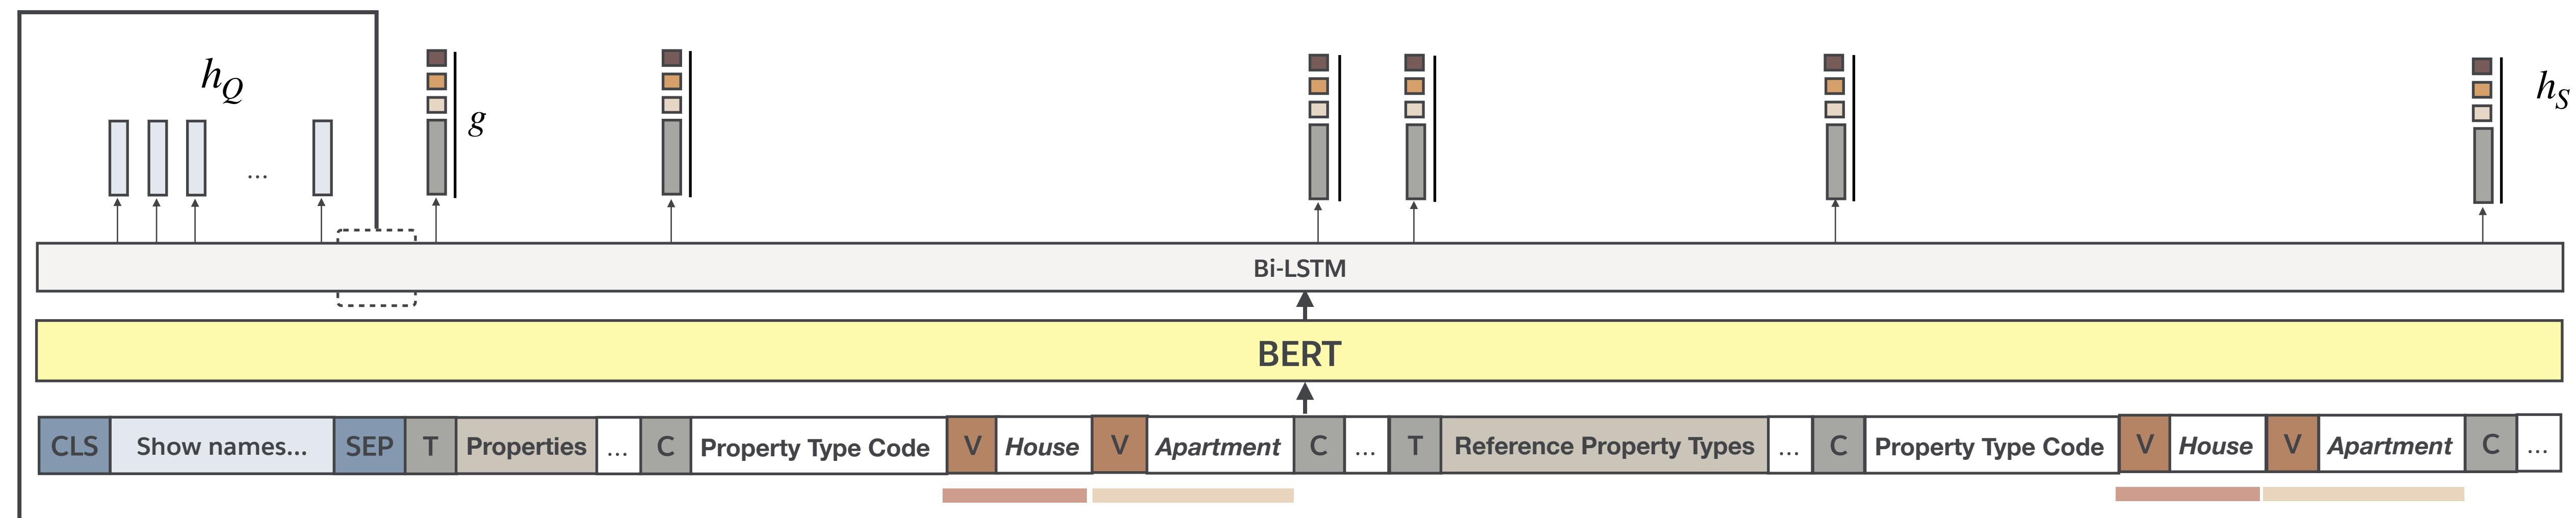
Reference Property Types

Property type code	Property type description
Apartment	...
Field	...
House	...
Shop	...
Other	...

Decoder

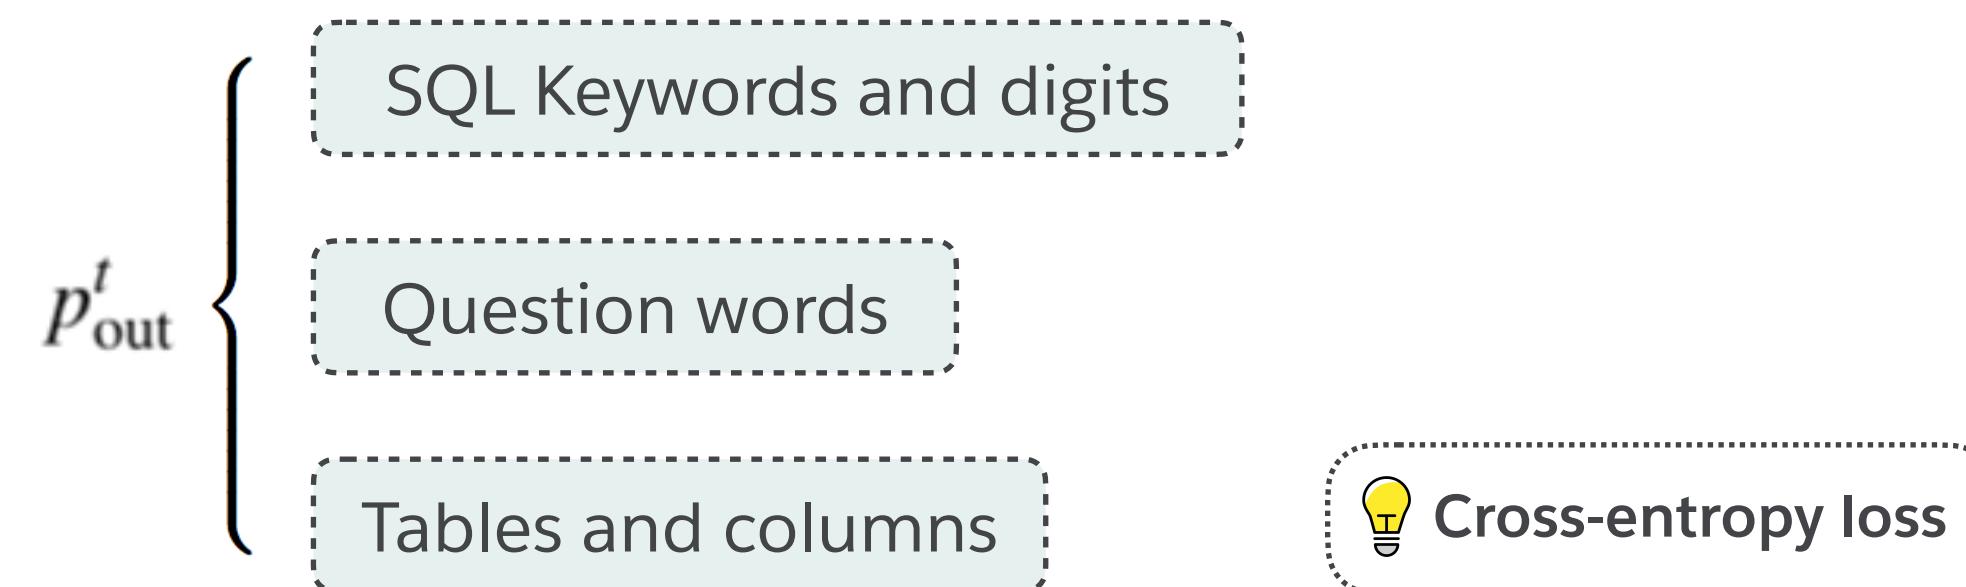


Decoder



>Show names of properties that are either houses or apartments

→LSTM-based pointer-generator (See et al. 2017)



Schema-Consistency Guided Decoding

Pruning the search space of a sequential pointer-generator decoder

- **Observation:** The FROM clauses set the scope of a SQL query and the table fields appeared in the rest of the clauses can only belong to the tables in FROM

```
SELECT T2.name FROM Instructor AS T1 JOIN Department AS T2 ON T1.Department_ID = T2.ID  
GROUP BY T1.Department_ID HAVING AVG(T1.Rating) > (SELECT AVG(Rating) FROM Instructor)
```

Schema-Consistency Guided Decoding

Pruning the search space of a sequential pointer-generator decoder

- **Observation:** The FROM clauses set the scope of a SQL query and the table fields appeared in the rest of the clauses can only belong to the tables in FROM

```
SELECT T2.name FROM Instructor AS T1 JOIN Department AS T2 ON T1.Department_ID = T2.ID  
GROUP BY T1.Department_ID HAVING AVG(T1.Rating) > (SELECT AVG(Rating) FROM Instructor)
```

 Rewrite a SQL query in execution order, with FROM clause at the beginning of each sub-query

```
FROM Instructor AS T1 JOIN Department AS T2 ON T1.Department_ID = T2.ID SELECT T2.name  
GROUP BY T1.Department_ID HAVING AVG(T1.Rating) > (FROM Instructor SELECT AVG(Rating))
```

Schema-Consistency Guided Decoding

Pruning the search space of a sequential pointer-generator decoder

- **Observation:** The FROM clauses set the scope of a SQL query and the table fields appeared in the rest of the clauses can only belong to the tables in FROM

```
SELECT T2.name FROM Instructor AS T1 JOIN Department AS T2 ON T1.Department_ID = T2.ID  
GROUP BY T1.Department_ID HAVING AVG(T1.Rating) > (SELECT AVG(Rating) FROM Instructor)
```

 Rewrite a SQL query in execution order, with FROM clause at the beginning of each sub-query

```
FROM Instructor AS T1 JOIN Department AS T2 ON T1.Department_ID = T2.ID SELECT T2.name  
GROUP BY T1.Department_ID HAVING AVG(T1.Rating) > (FROM Instructor SELECT AVG(Rating))
```

Lemma: Let Y_{exec} be a SQL query with clauses arranged in execution order, then any table field in Y_{exec} will appear after the corresponding table token.

Schema-Consistency Guided Decoding

- Generate SQL queries in execution order and unmask DB fields dynamically



Schema-Consistency Guided Decoding

- Generate SQL queries in execution order and unmask DB fields dynamically



FROM

Schema-Consistency Guided Decoding

- Generate SQL queries in execution order and unmask DB fields dynamically



FROM Instructor

Schema-Consistency Guided Decoding

- Generate SQL queries in execution order and unmask DB fields dynamically



FROM Instructor JOIN

Schema-Consistency Guided Decoding

- Generate SQL queries in execution order and unmask DB fields dynamically



`FROM Instructor JOIN Department`

Schema-Consistency Guided Decoding

- Generate SQL queries in execution order and unmask DB fields dynamically



FROM Instructor JOIN Department ON

Schema-Consistency Guided Decoding

- Generate SQL queries in execution order and unmask DB fields dynamically



```
FROM Instructor JOIN Department ON Instructor.Department_ID = Department.ID SELECT
Department.name GROUP BY Instructor.Department_ID HAVING_AVG(Instruction.Rating) >
(FROM Instructor SELECT AVG(Instruction.Rating))
```

Schema-Consistency Guided Decoding

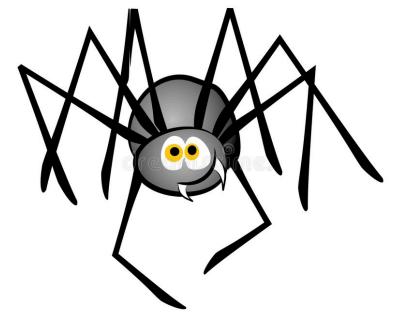
- Generate SQL queries in execution order and unmask DB fields dynamically



```
FROM Instructor JOIN Department ON Instructor.Department_ID = Department.ID SELECT  
Department.name GROUP BY Instructor.Department_ID HAVING_AVG(Instruction.Rating) >  
(FROM Instructor SELECT AVG(Instruction.Rating))
```

- ✓ Implemented via vector space computation
- ✓ Applied during inference only
- ✓ Cannot guarantee schema consistency, used in combination with static SQL correctness check
- ✓ Can be applied to other types of decoders

Dataset - Spider



(Yu et al. 2018)

Expert-annotated, cross-domain, complex
text-to-SQL dataset

No overlap between Train/Dev/Test
databases, enabling the development of text-
to-SQL models which generalize to unseen DBs

	Train	Dev	Test
Hidden			
# DBs	146	20	40
# Examples	8,659	1,034	2,147

Database

Instructor

Primary key

ID	Name	Department_ID	Salary	...
----	------	---------------	--------	-----

Department

Foreign key

ID	Name	Building	Budget	...
----	------	----------	--------	-----

...

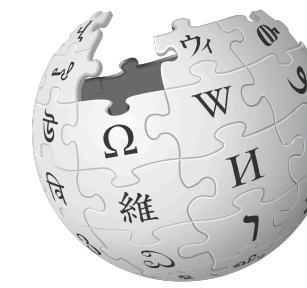
...

Question What are the name and budget of the departments
with average instructor salary above the overall average?

SQL

```
SELECT T2.name, T2.budget
FROM Instructor AS T1 JOIN Department AS T2 ON
T1.Department_ID = T2.ID
GROUP BY T1.Department_ID
HAVING AVG(T1.salary) >
(SELECT AVG(Salary) FROM Instructor)
```

Dataset - WikiSQL



(Zhong et al. 2017)

Generated over **Wikipedia tables** using the **semantic-parsing-overnight** approach (Wang et al. 2015)

SQL Template: `SELECT $AGG $COLUMN
WHERE $COLUMN $OP $VALUE
(AND $COLUMN $OP $VALUE) *`

Train/Dev/Test tables overlap, but **49.6%** of dev tables are not in the train set and **45.1%** of test tables are not in the train set.

WikiTable

Player	No.	Nationality	Position	Years in Toronto	School/Club
--------	-----	-------------	----------	------------------	-------------

Question Who is the player that wears number 45?

SQL `SELECT Player WHERE No. = 42`

	Train	Dev	Test
# Tables	17,984	1,621	2,787
# Examples	56,355	8,421	15,878

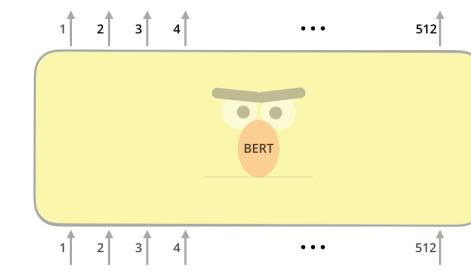
Experiment Setup

Pre-processing

- Compute fuzzy string match between the input question and the picklists of each DB field to obtain value mentions
- For each DB field, use the top- K matches in the DB schema representation ($K = 2$)

Decoding

- Beam search with length penalty
 - beam size = 16 for ablation study; beam size = 64 for leaderboard results



BERT-large-uncased, 24 layers
(Devlin et al. 2019)

Evaluation

- Exact set match
 - Logical form match ignoring values and SQL component order invariance
- Execution accuracy
 - Check if the execution results of the predicted SQL query matches the executions results of the ground-truth SQL query

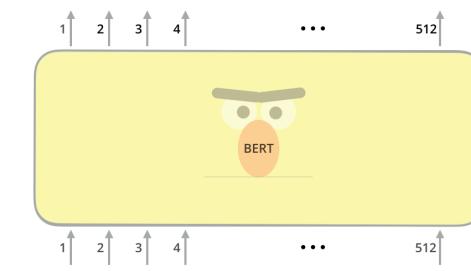
Experiment Setup

Pre-processing

- Compute fuzzy string match between the input question and the picklists of each DB field to obtain value mentions
- For each DB field, use the top- K matches in the DB schema representation ($K = 2$)

Decoding

- Beam search with length penalty
 - beam size = 16 for ablation study; beam size = 64 for leaderboard results



BERT-large-uncased, 24 layers
(Devlin et al. 2019)

Evaluation

- Exact set match
 - Logical form match ignoring values and SQL component order invariance
- Execution accuracy
 - Check if the execution results of the predicted SQL query matches the executions results of the ground-truth SQL query

Better evaluation for text-to-SQL is still an open research problem

Ablation Study

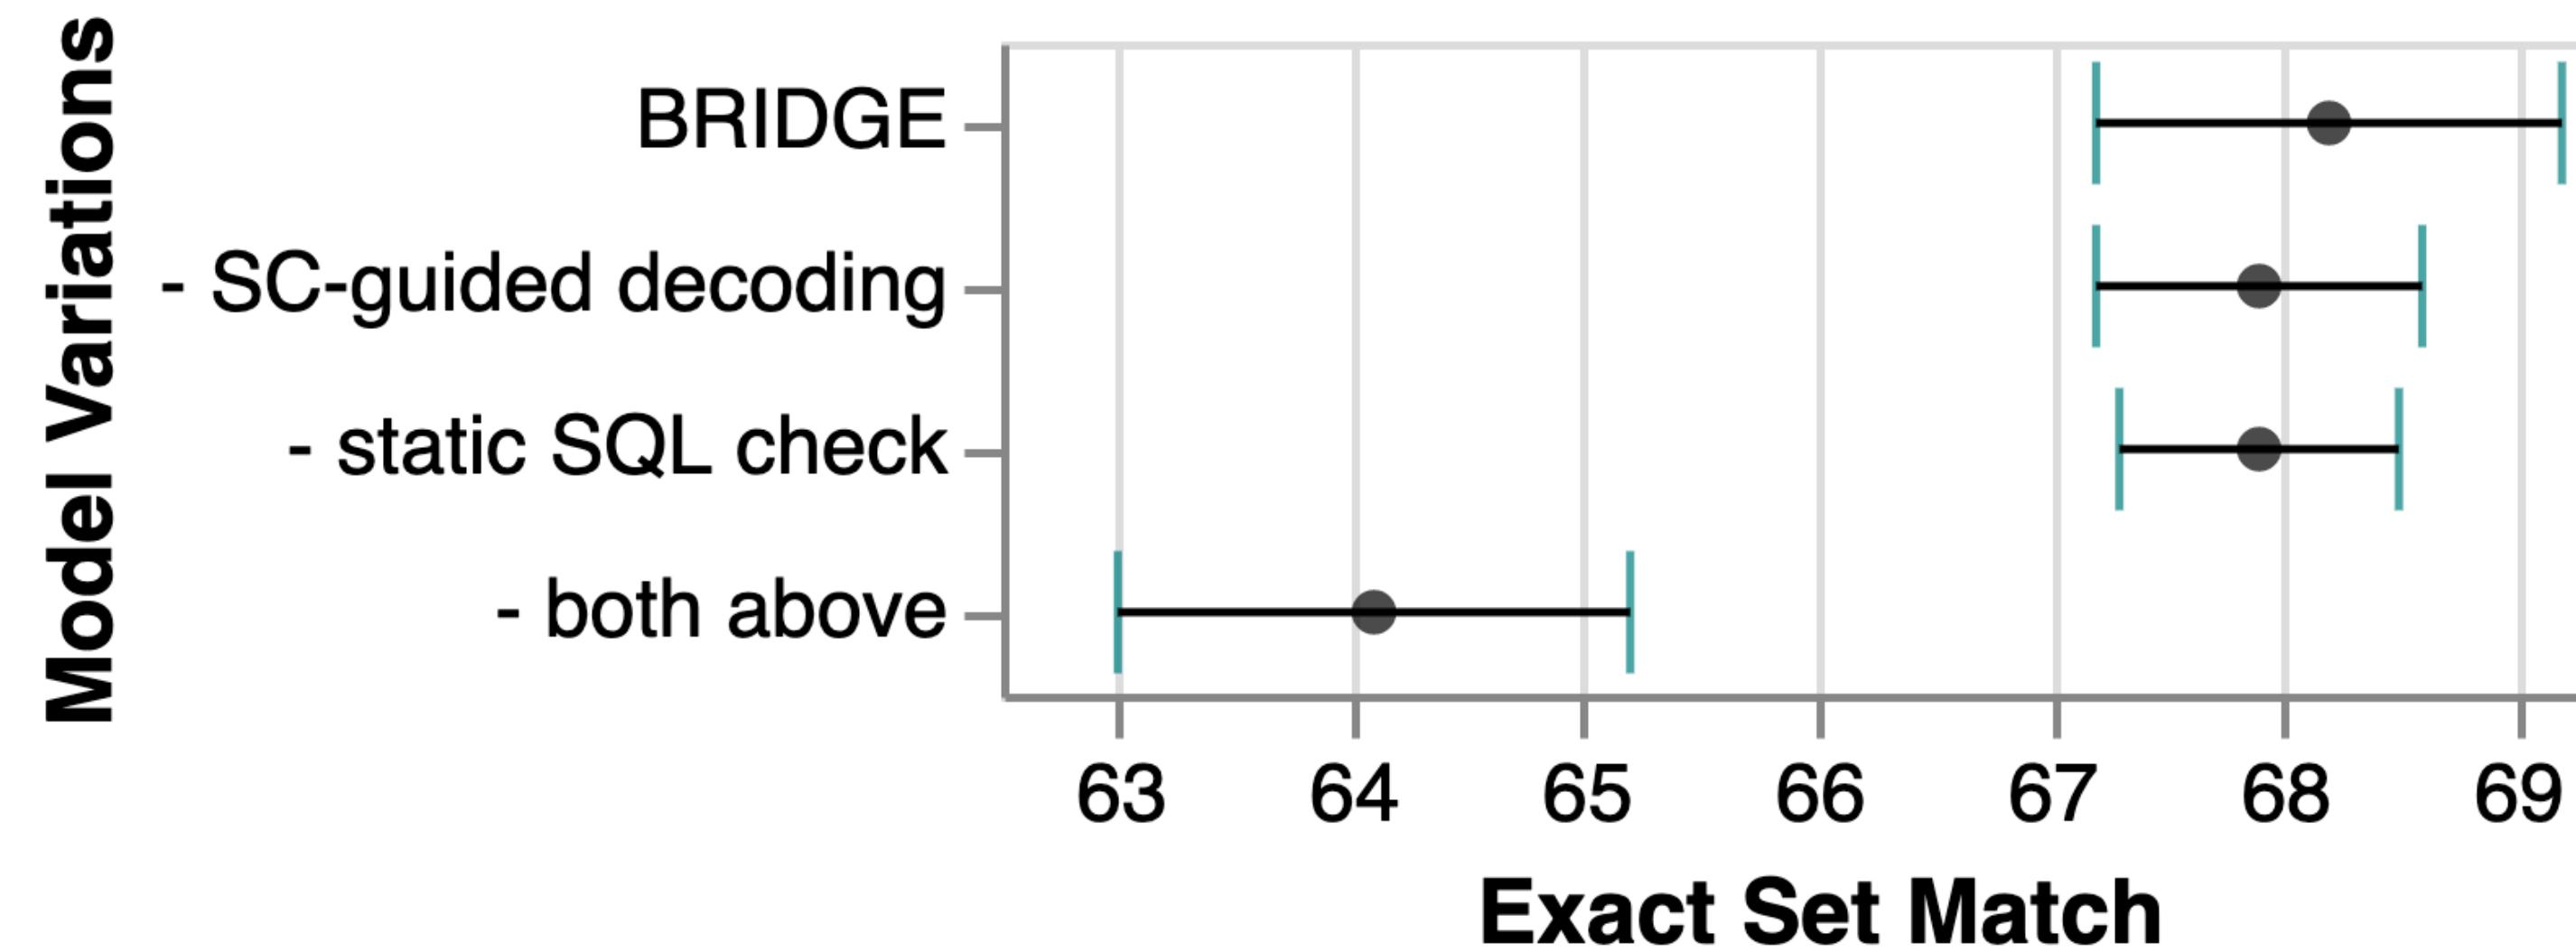
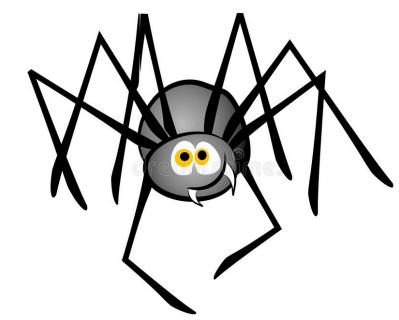


Figure 1. Ablation study of BRIDGE decoding strategies on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Ablation Study

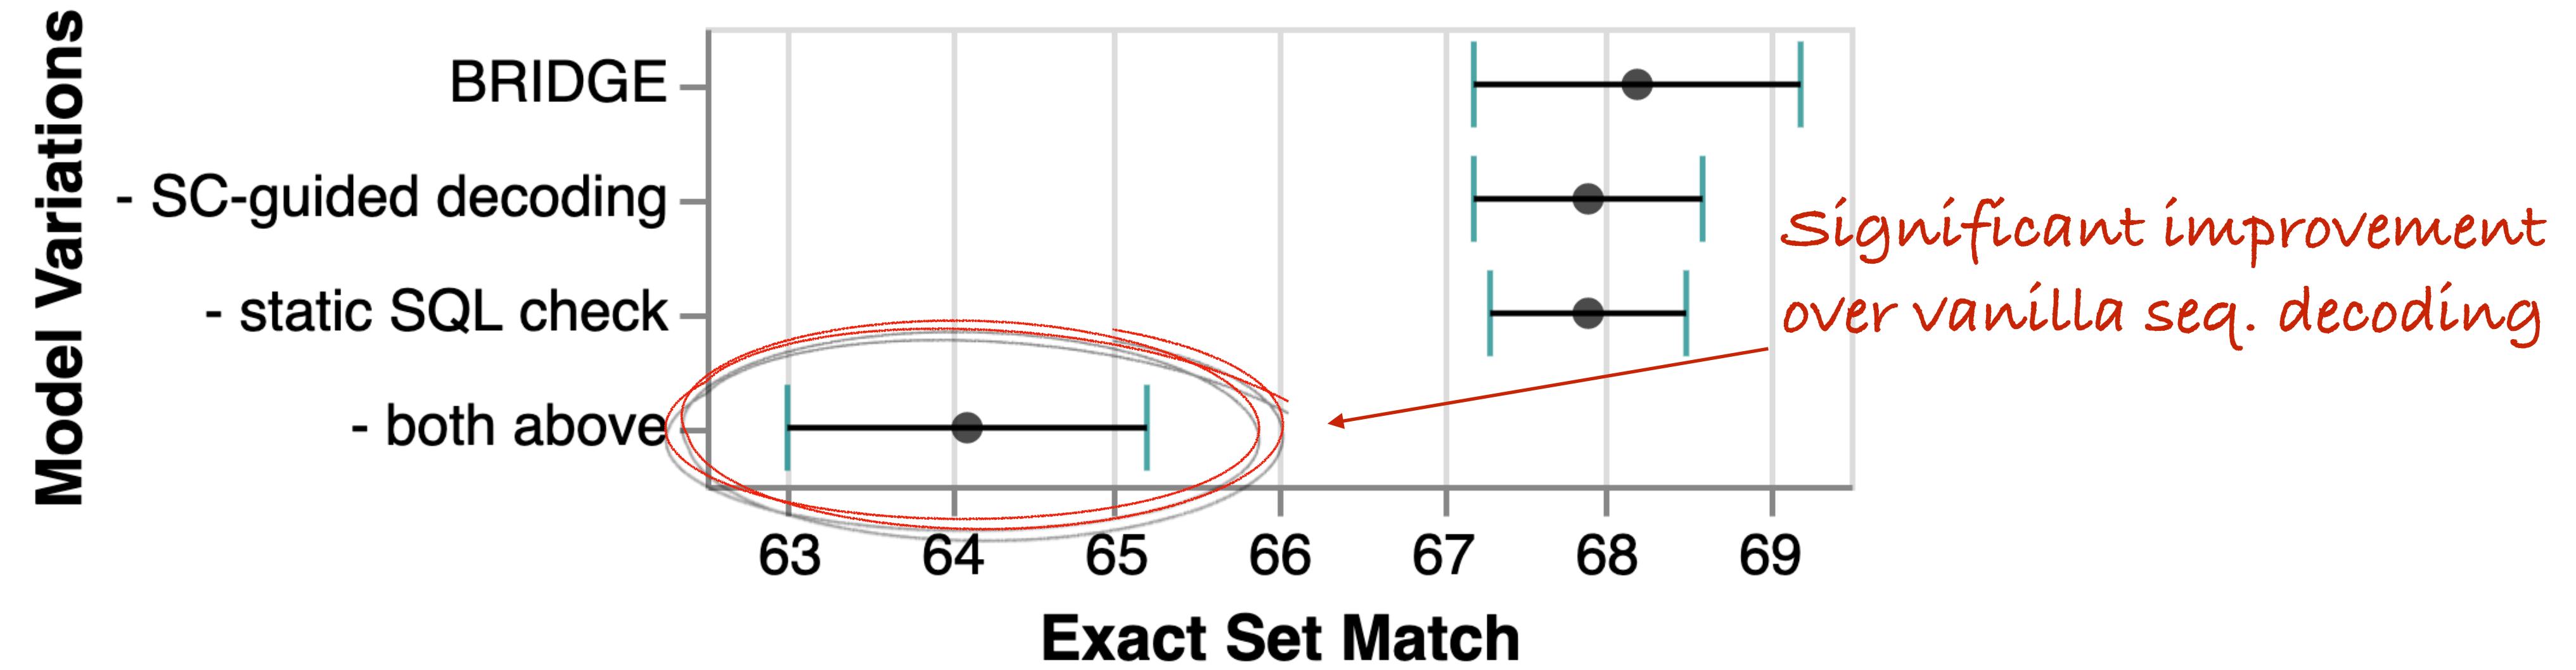
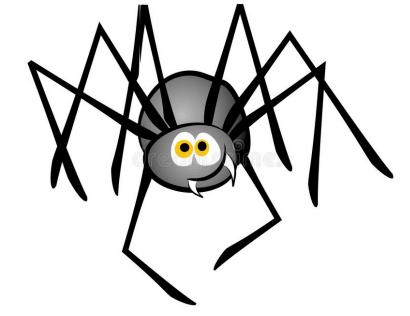


Figure 1. Ablation study of BRIDGE decoding strategies on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Ablation Study

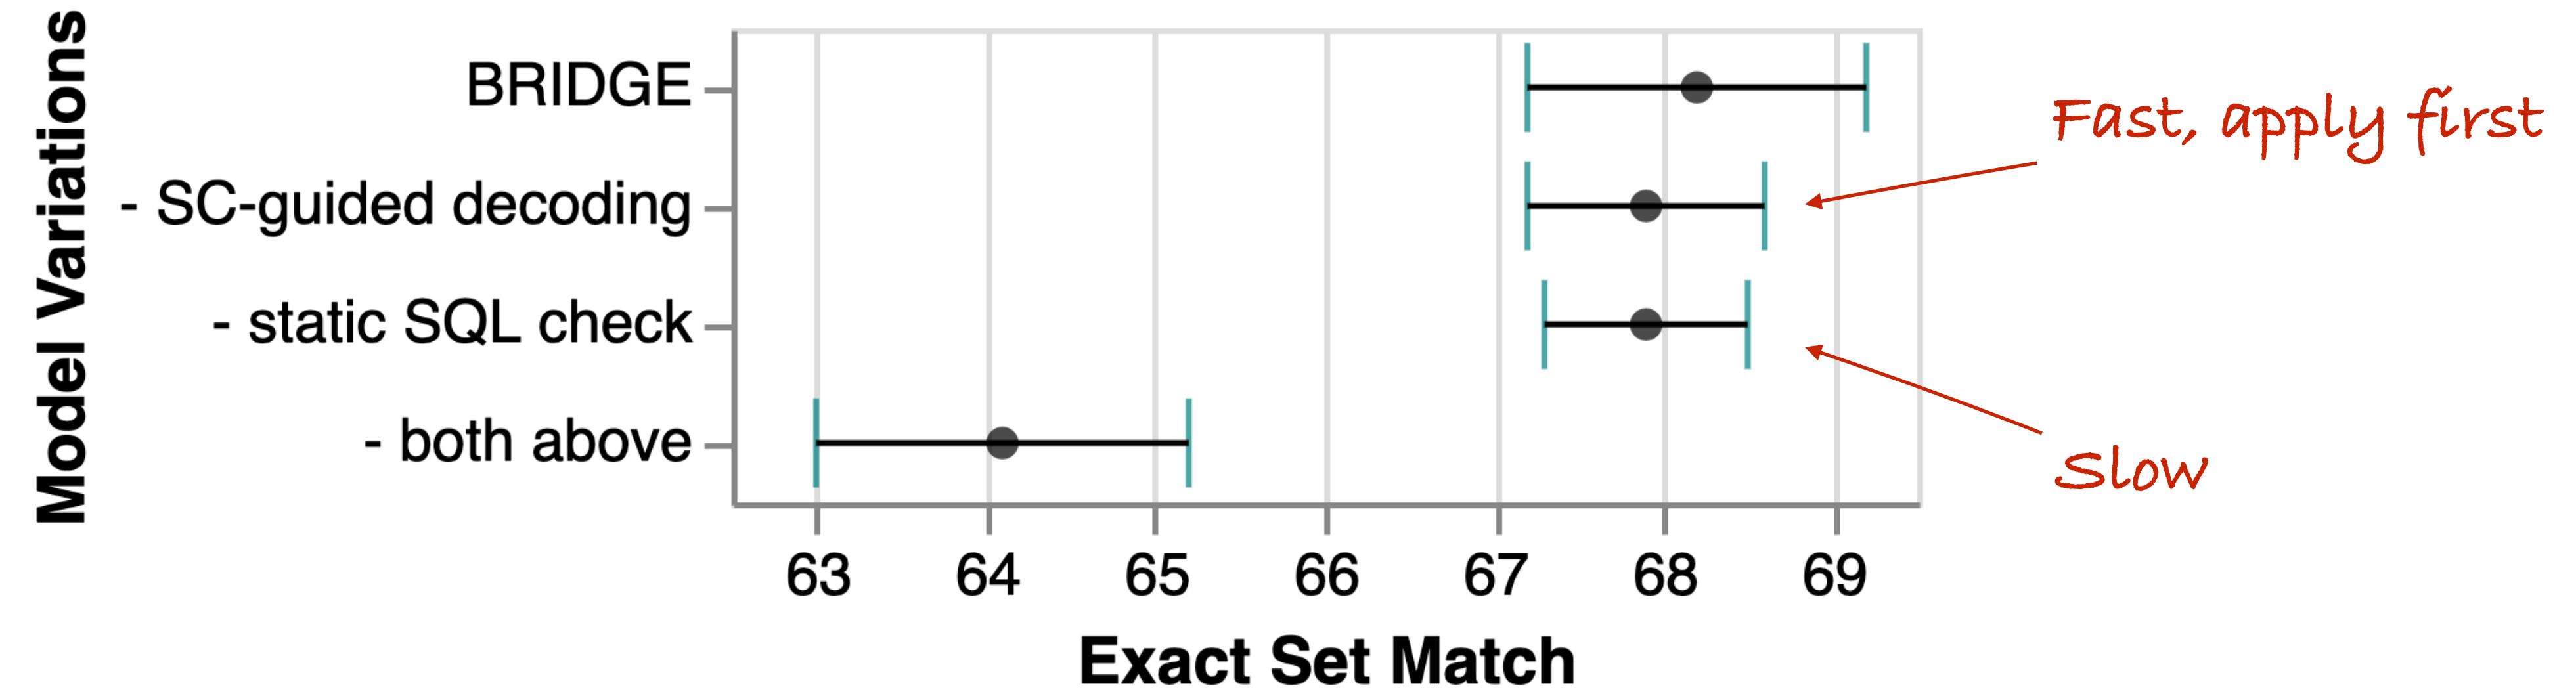
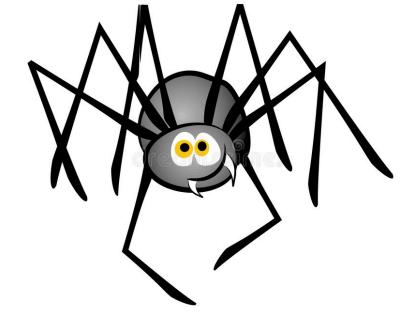


Figure 1. Ablation study of BRIDGE decoding strategies on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Ablation Study

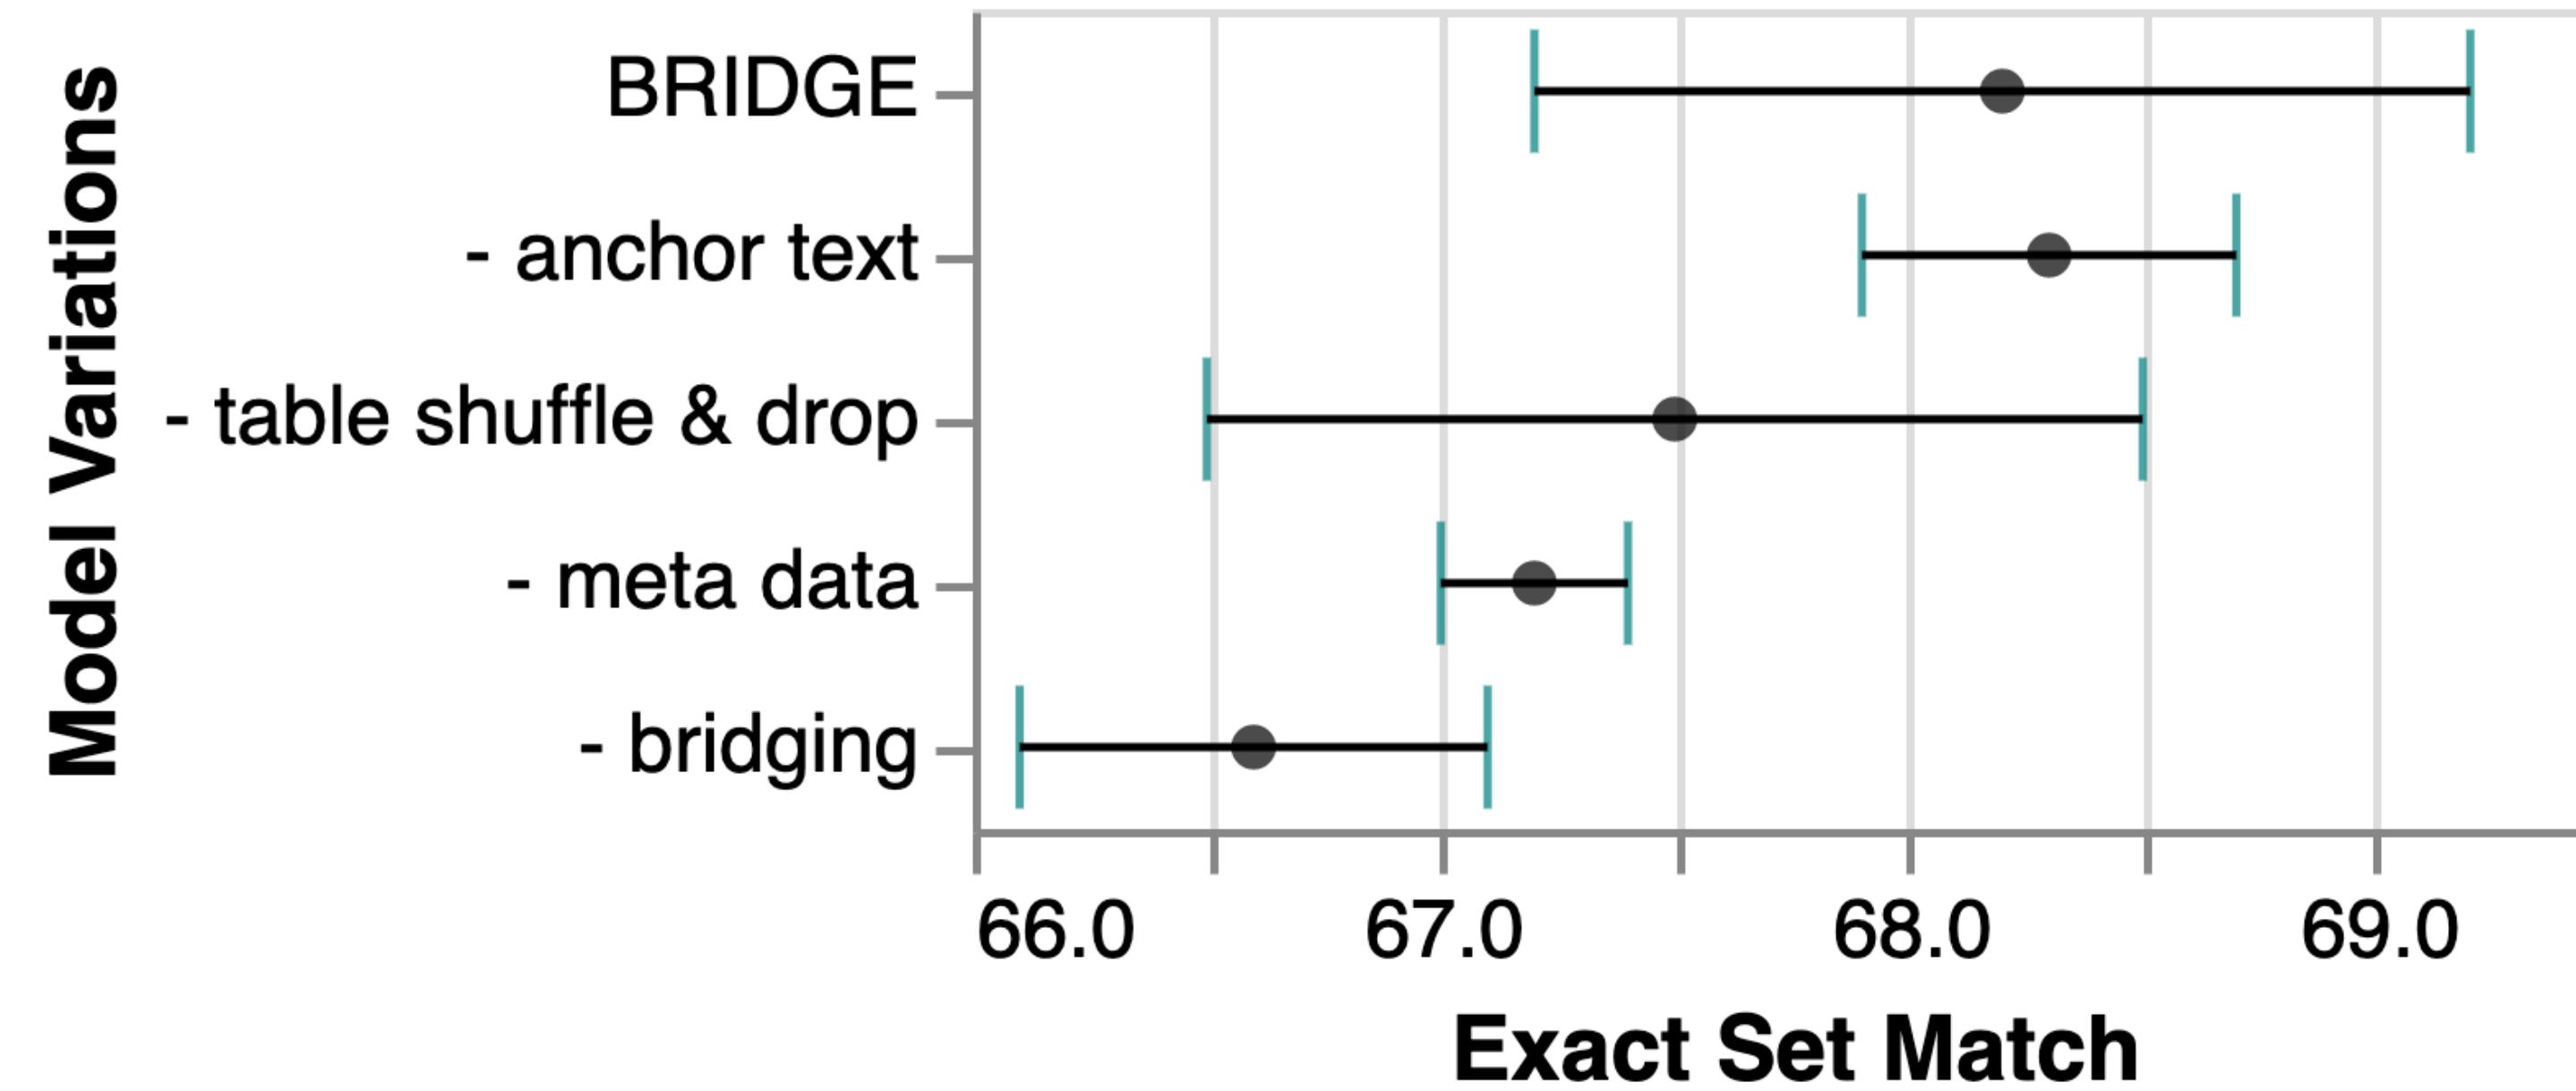
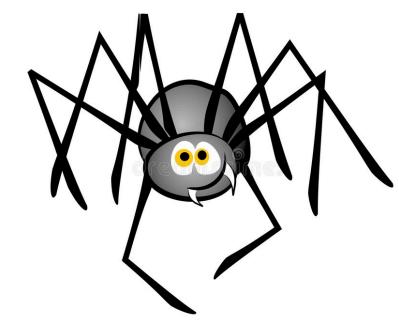


Figure 2. Ablation study of BRIDGE encoding strategies on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Ablation Study

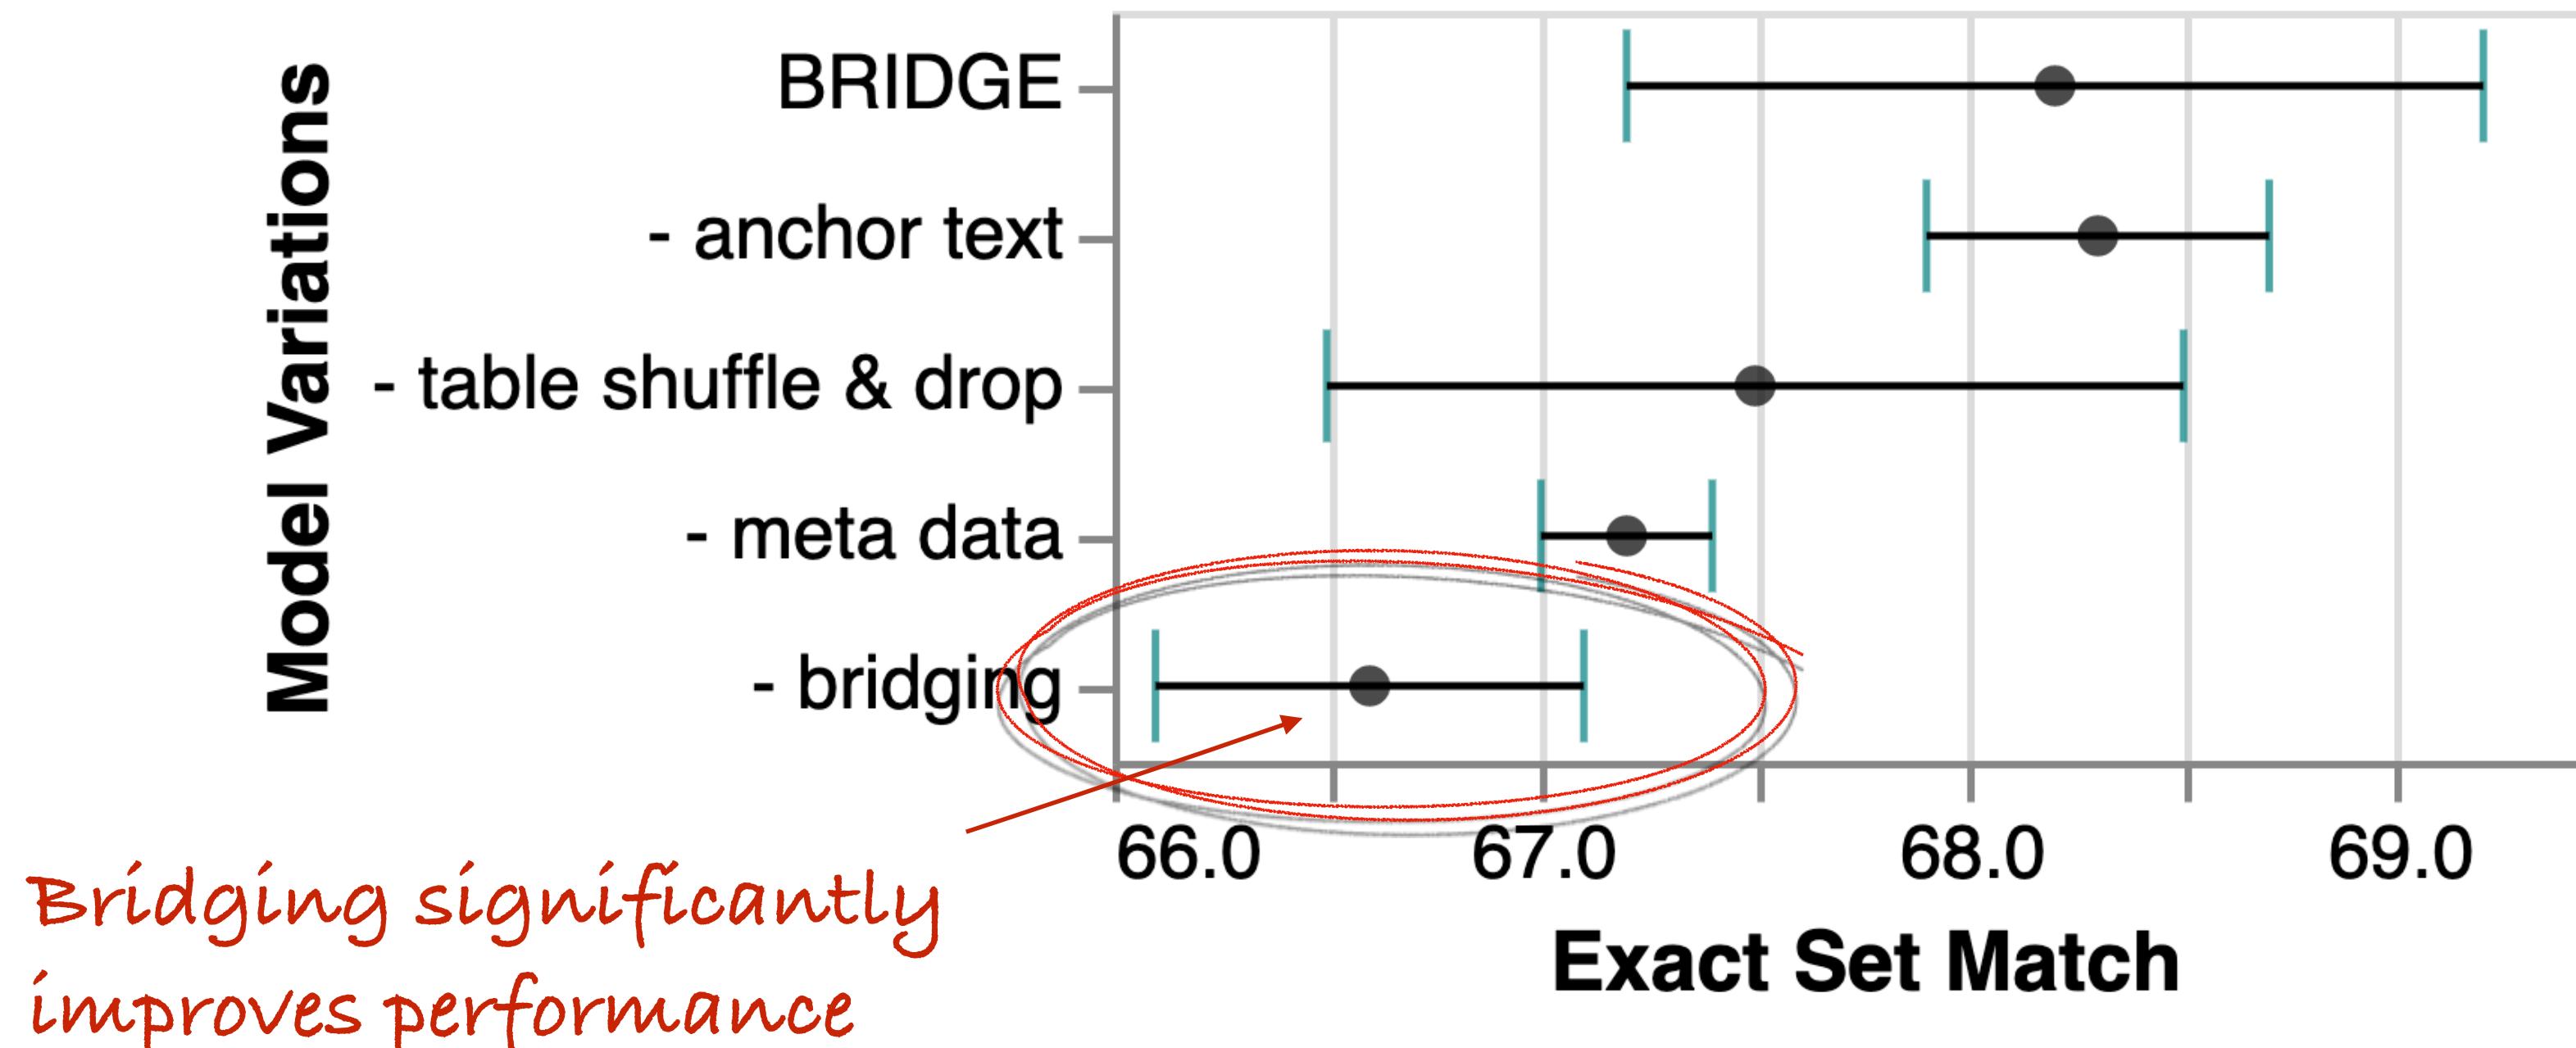
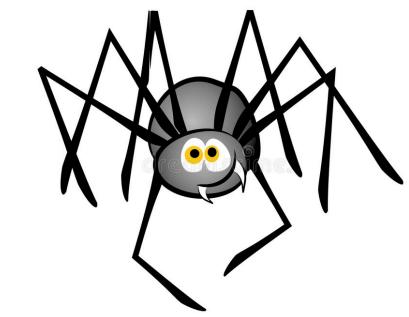


Figure 2. Ablation study of BRIDGE encoding strategies on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Bridging Ablation Performance by Difficulty Level

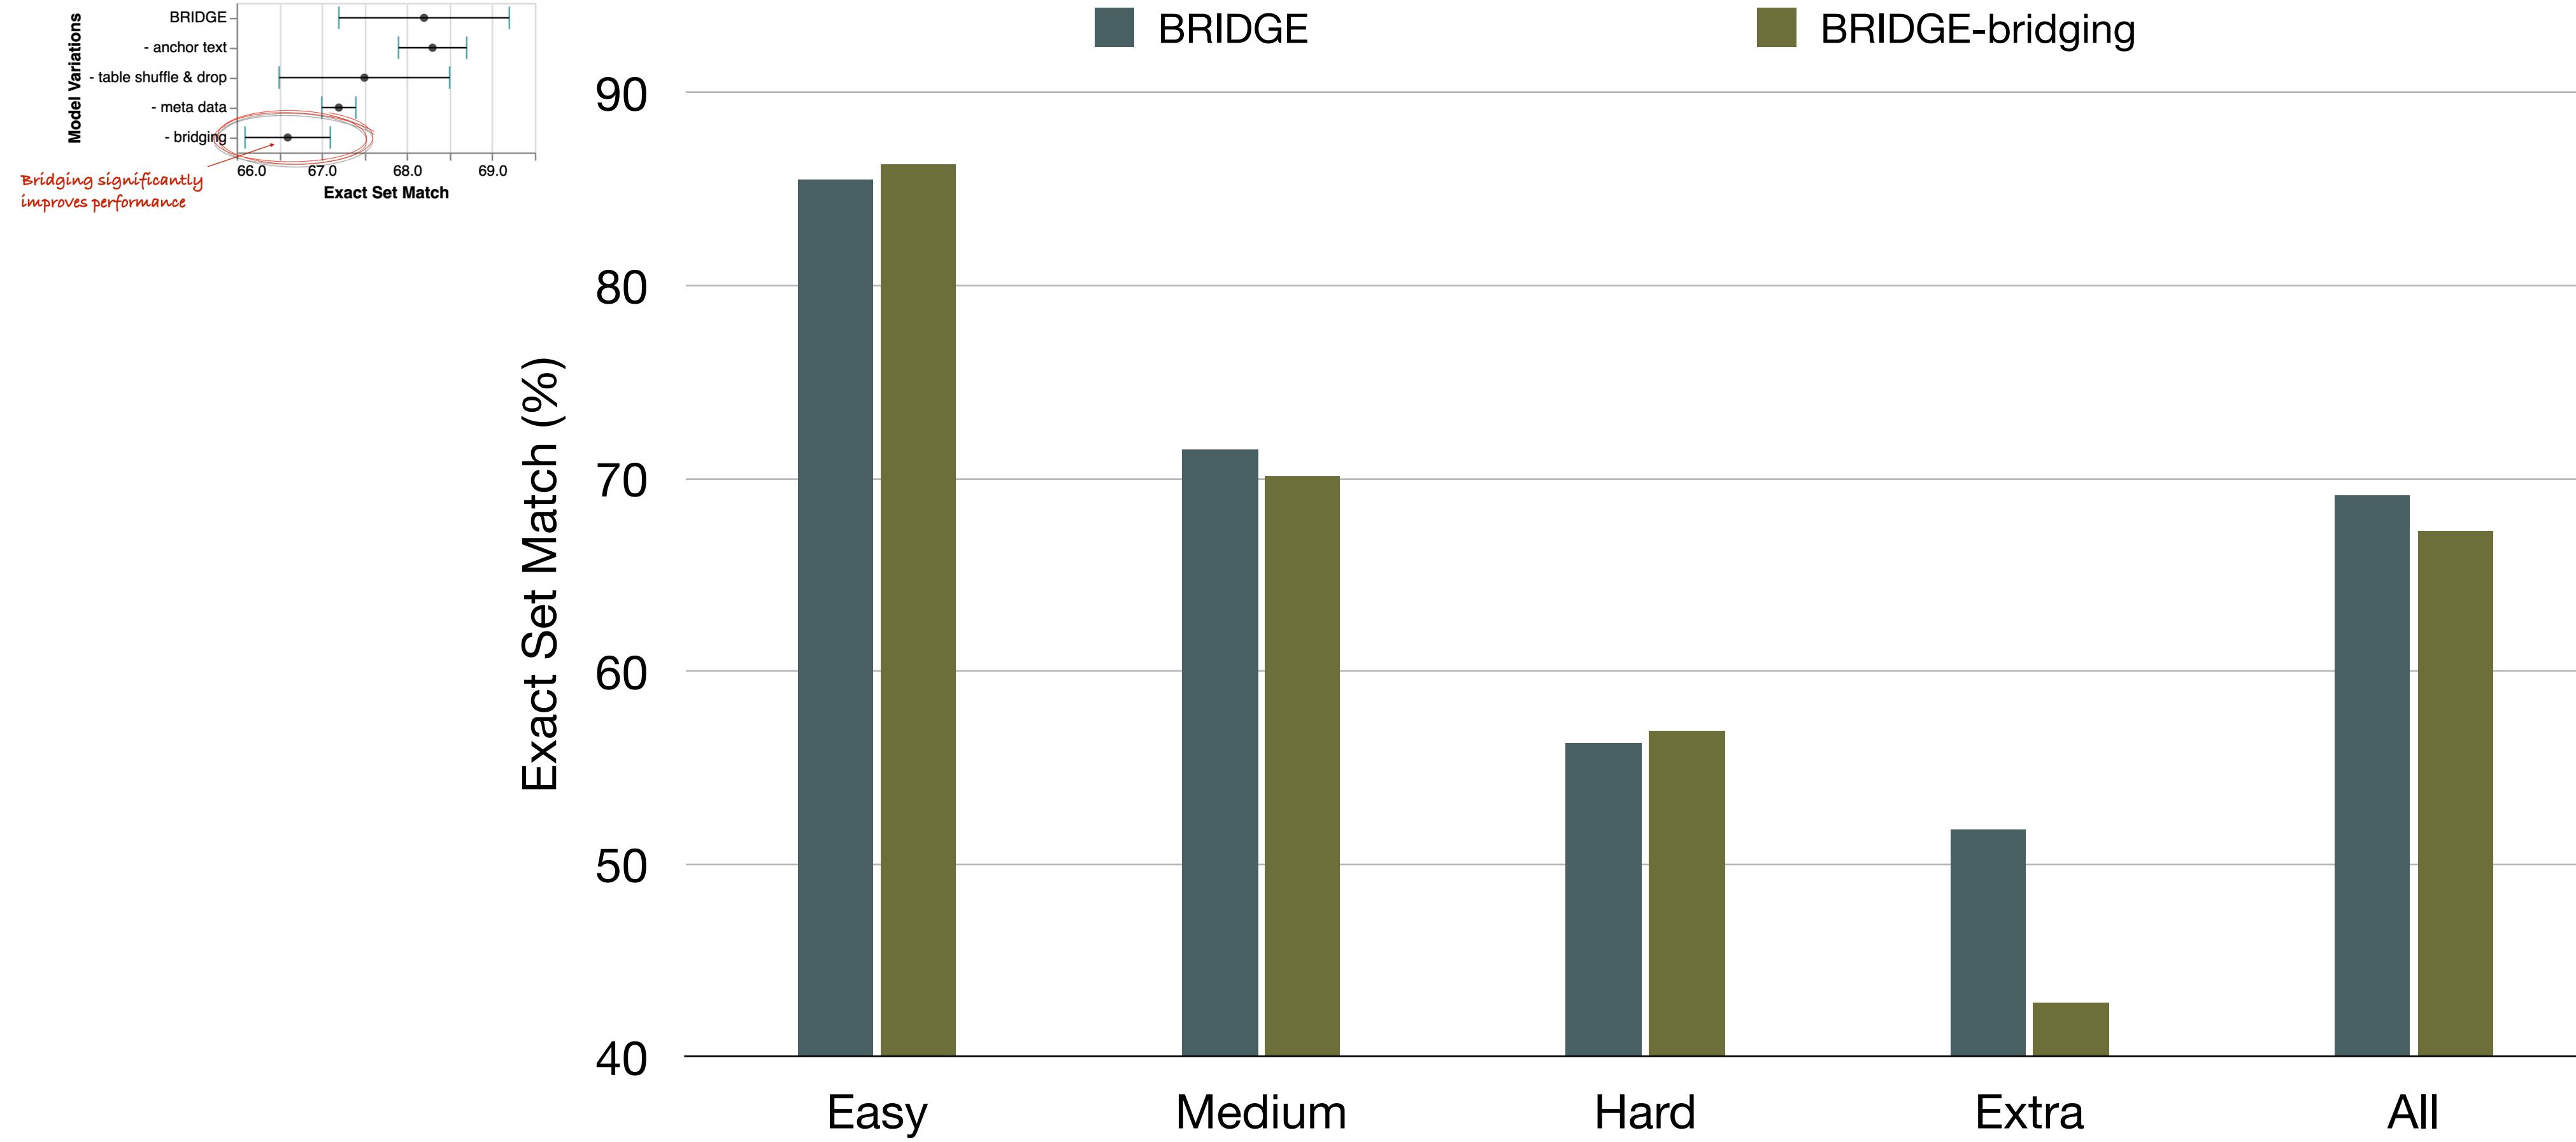
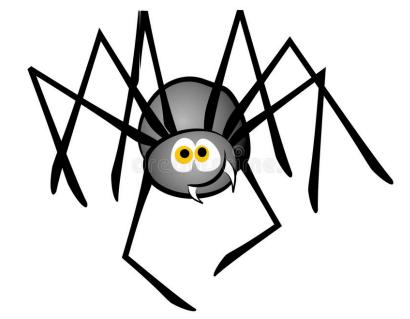


Figure 2.1. Performance of BRIDGE vs. BRIDGE - bridging on the Spider dev set.

Ablation Study

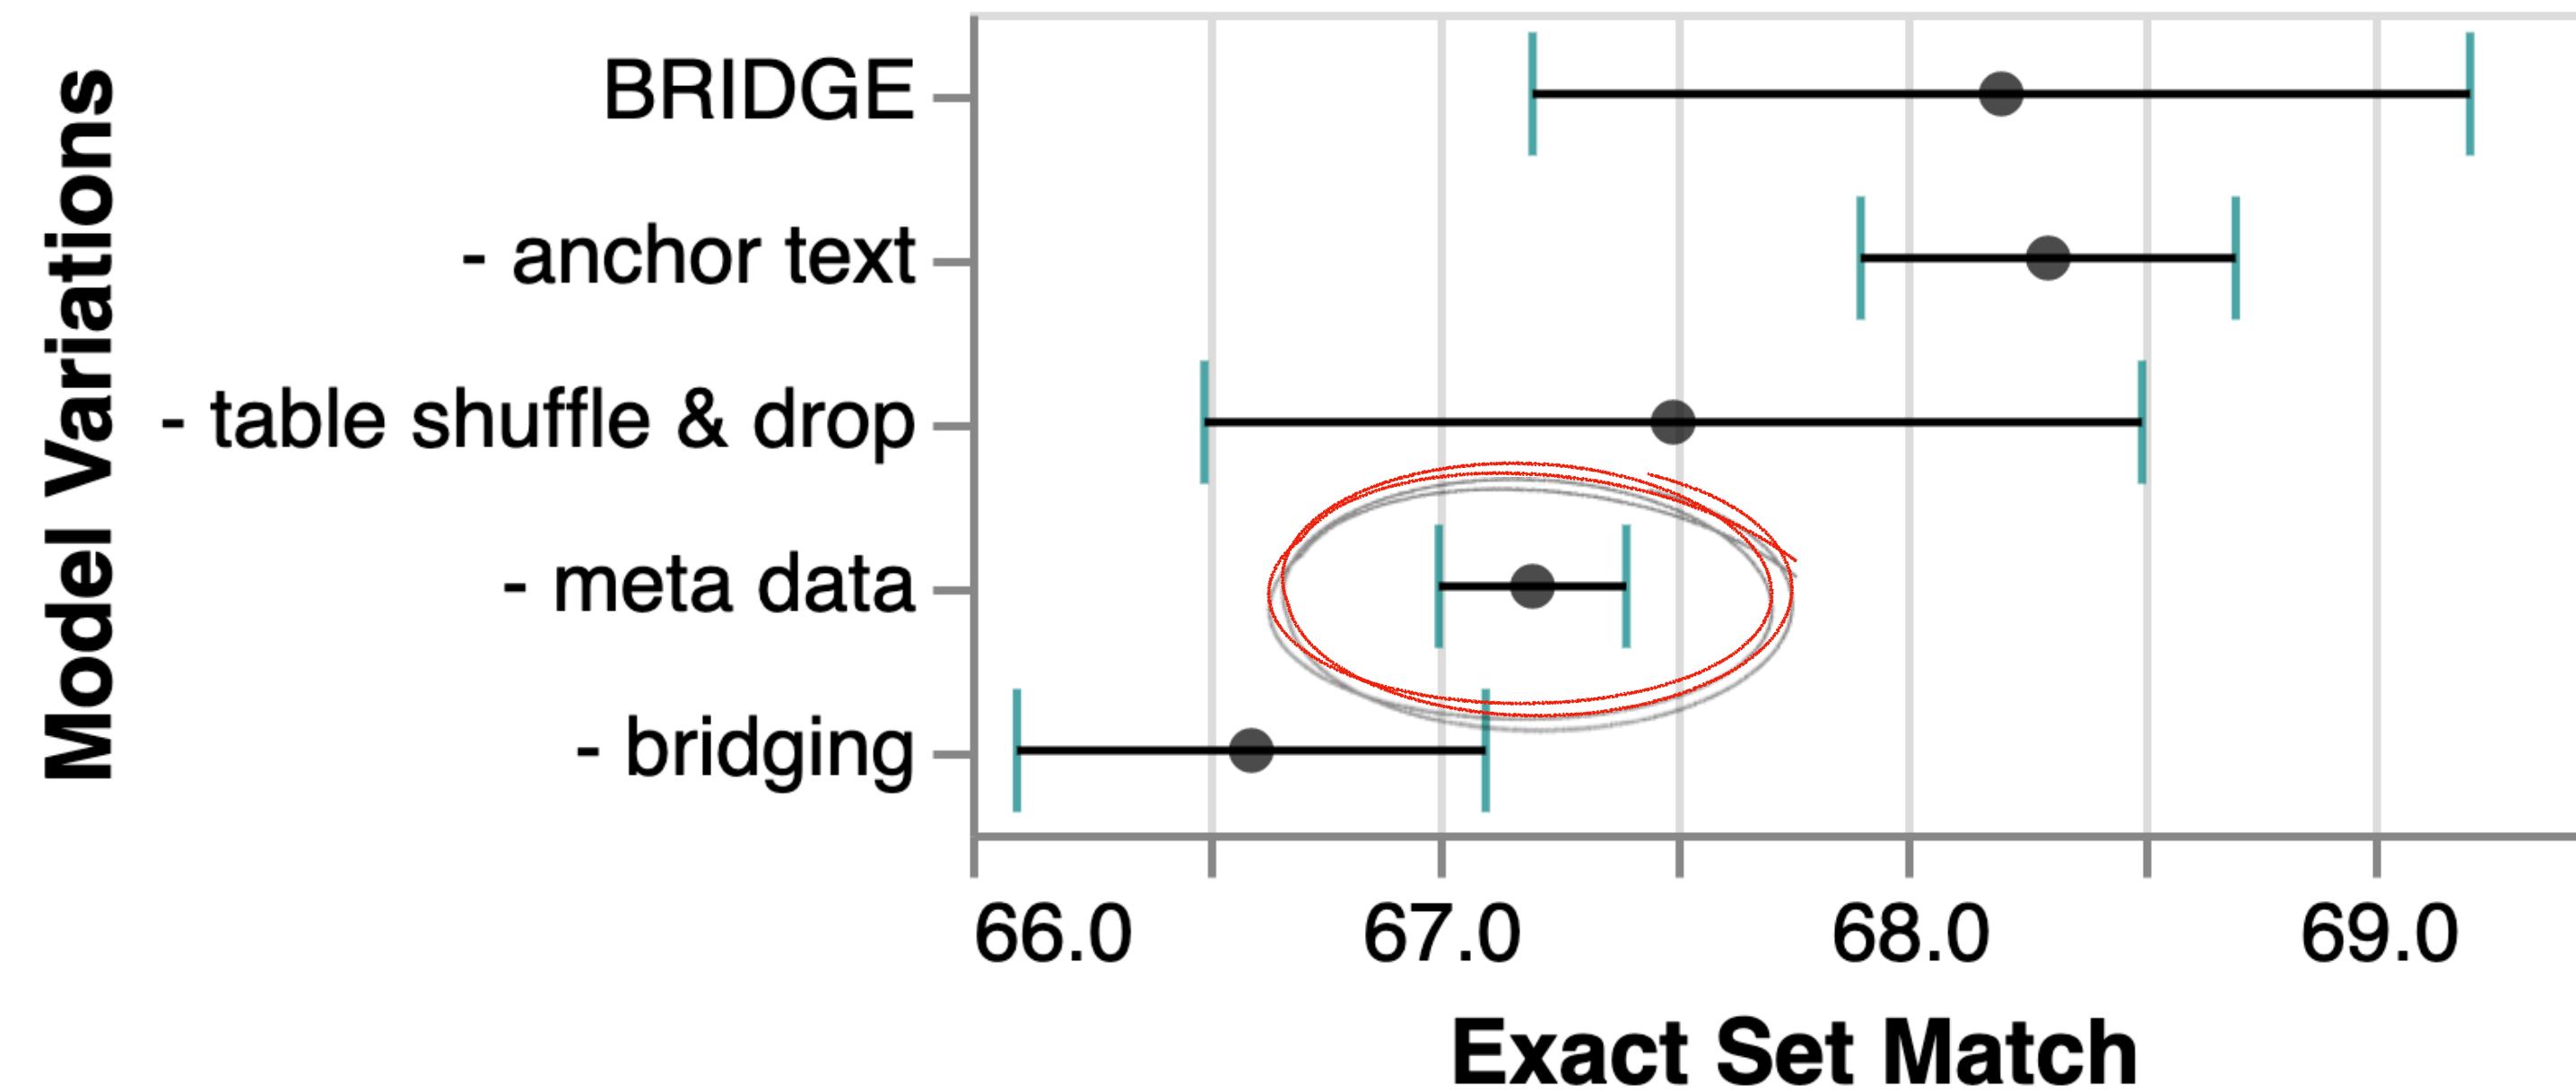
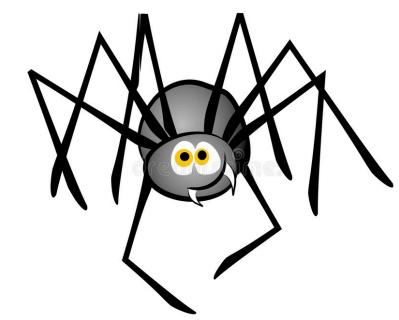


Figure 2. Ablation study of BRIDGE encoding strategies on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Ablation Study

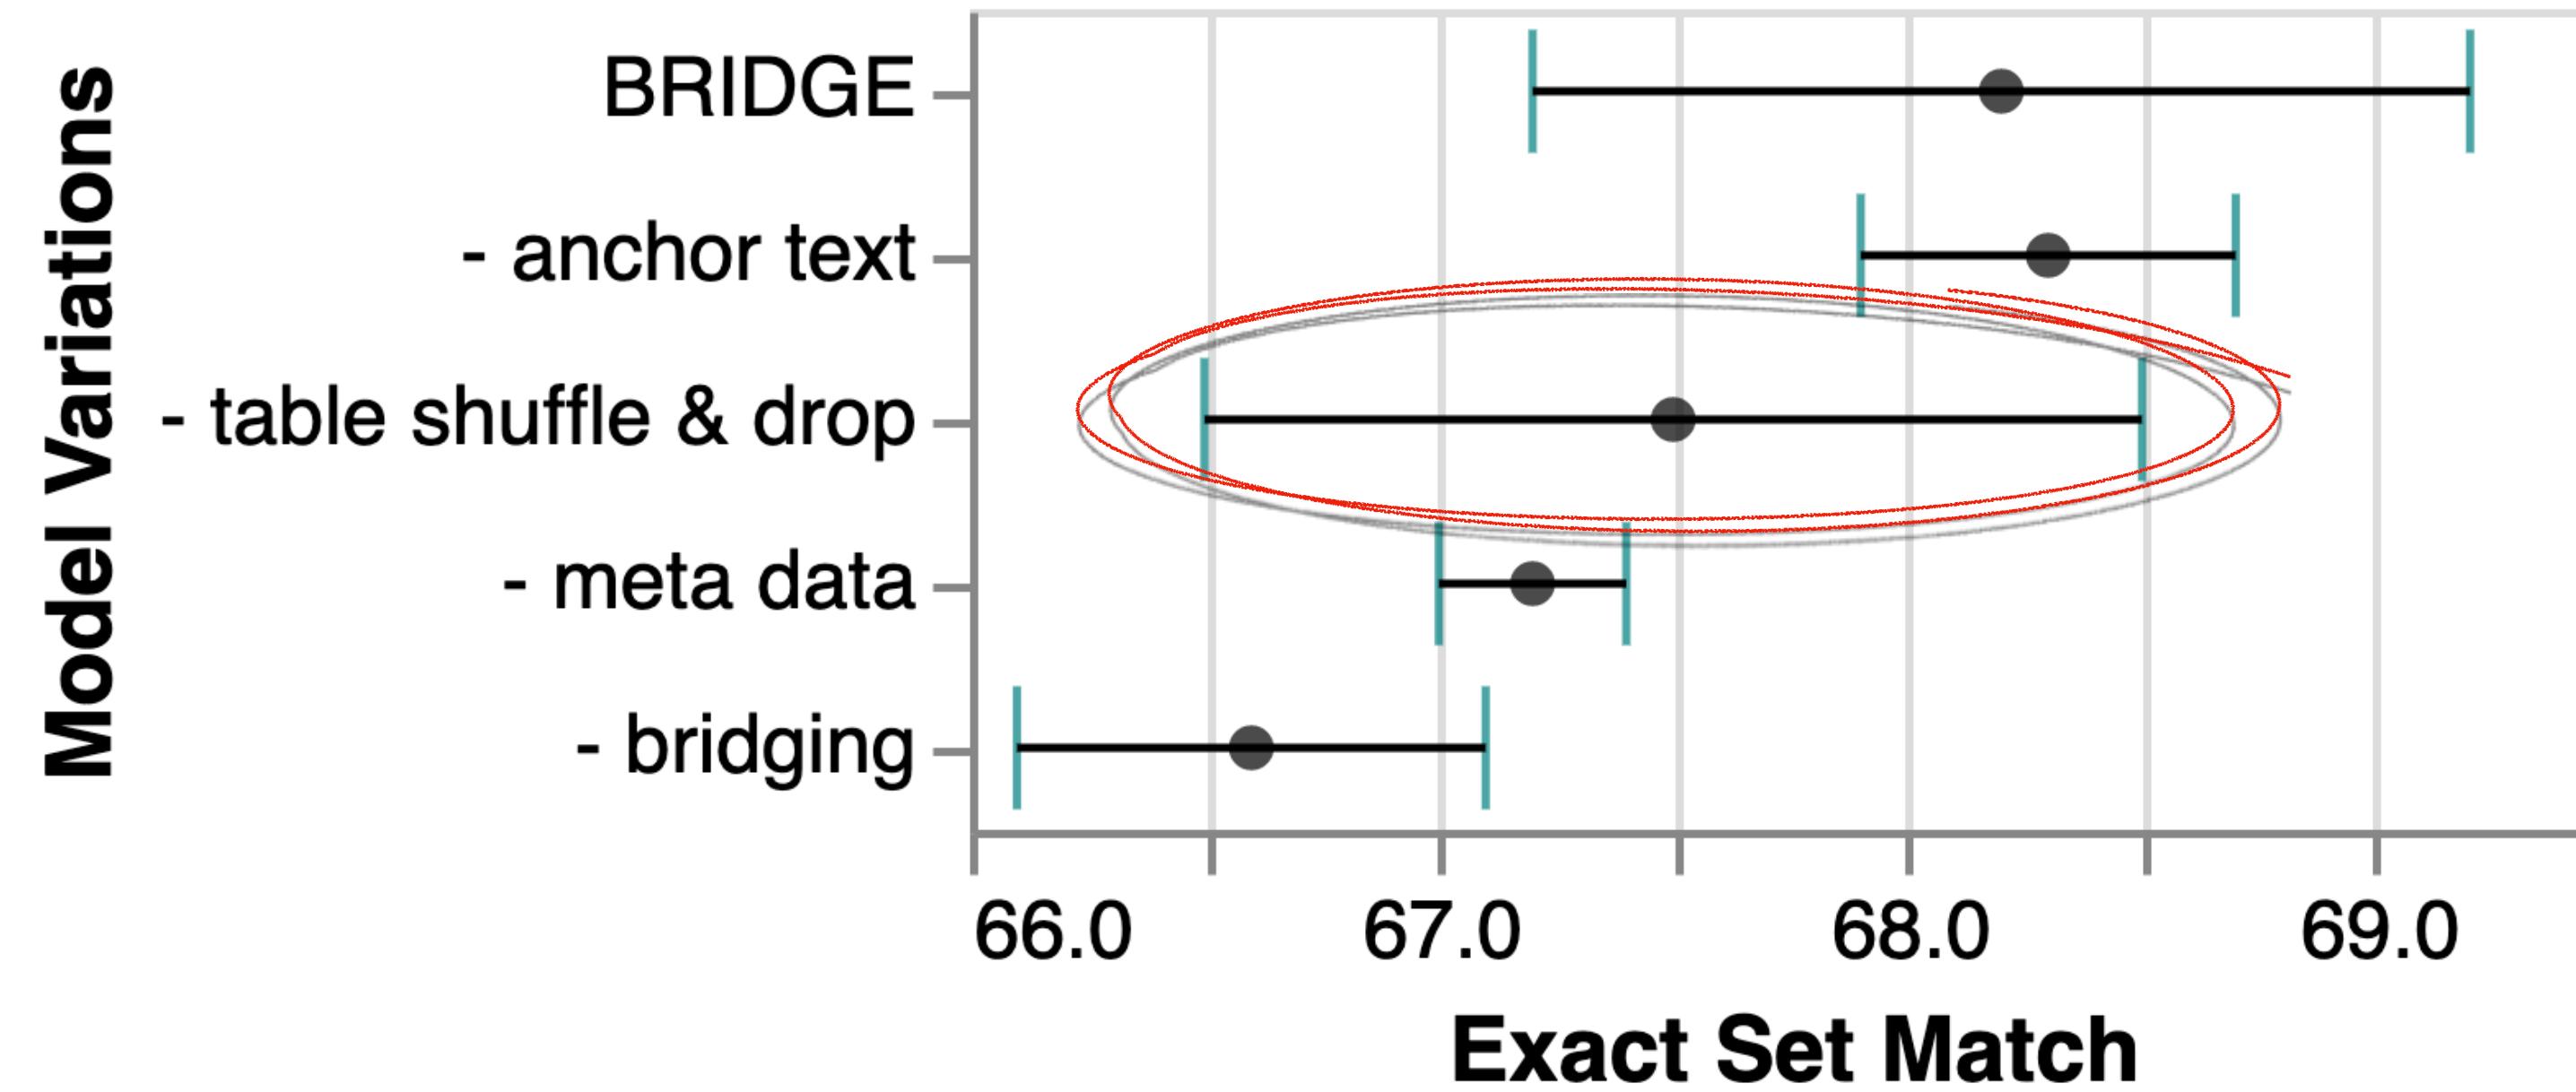
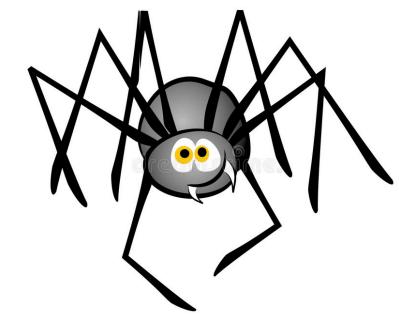


Figure 2. Ablation study of BRIDGE encoding strategies on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Ablation Study

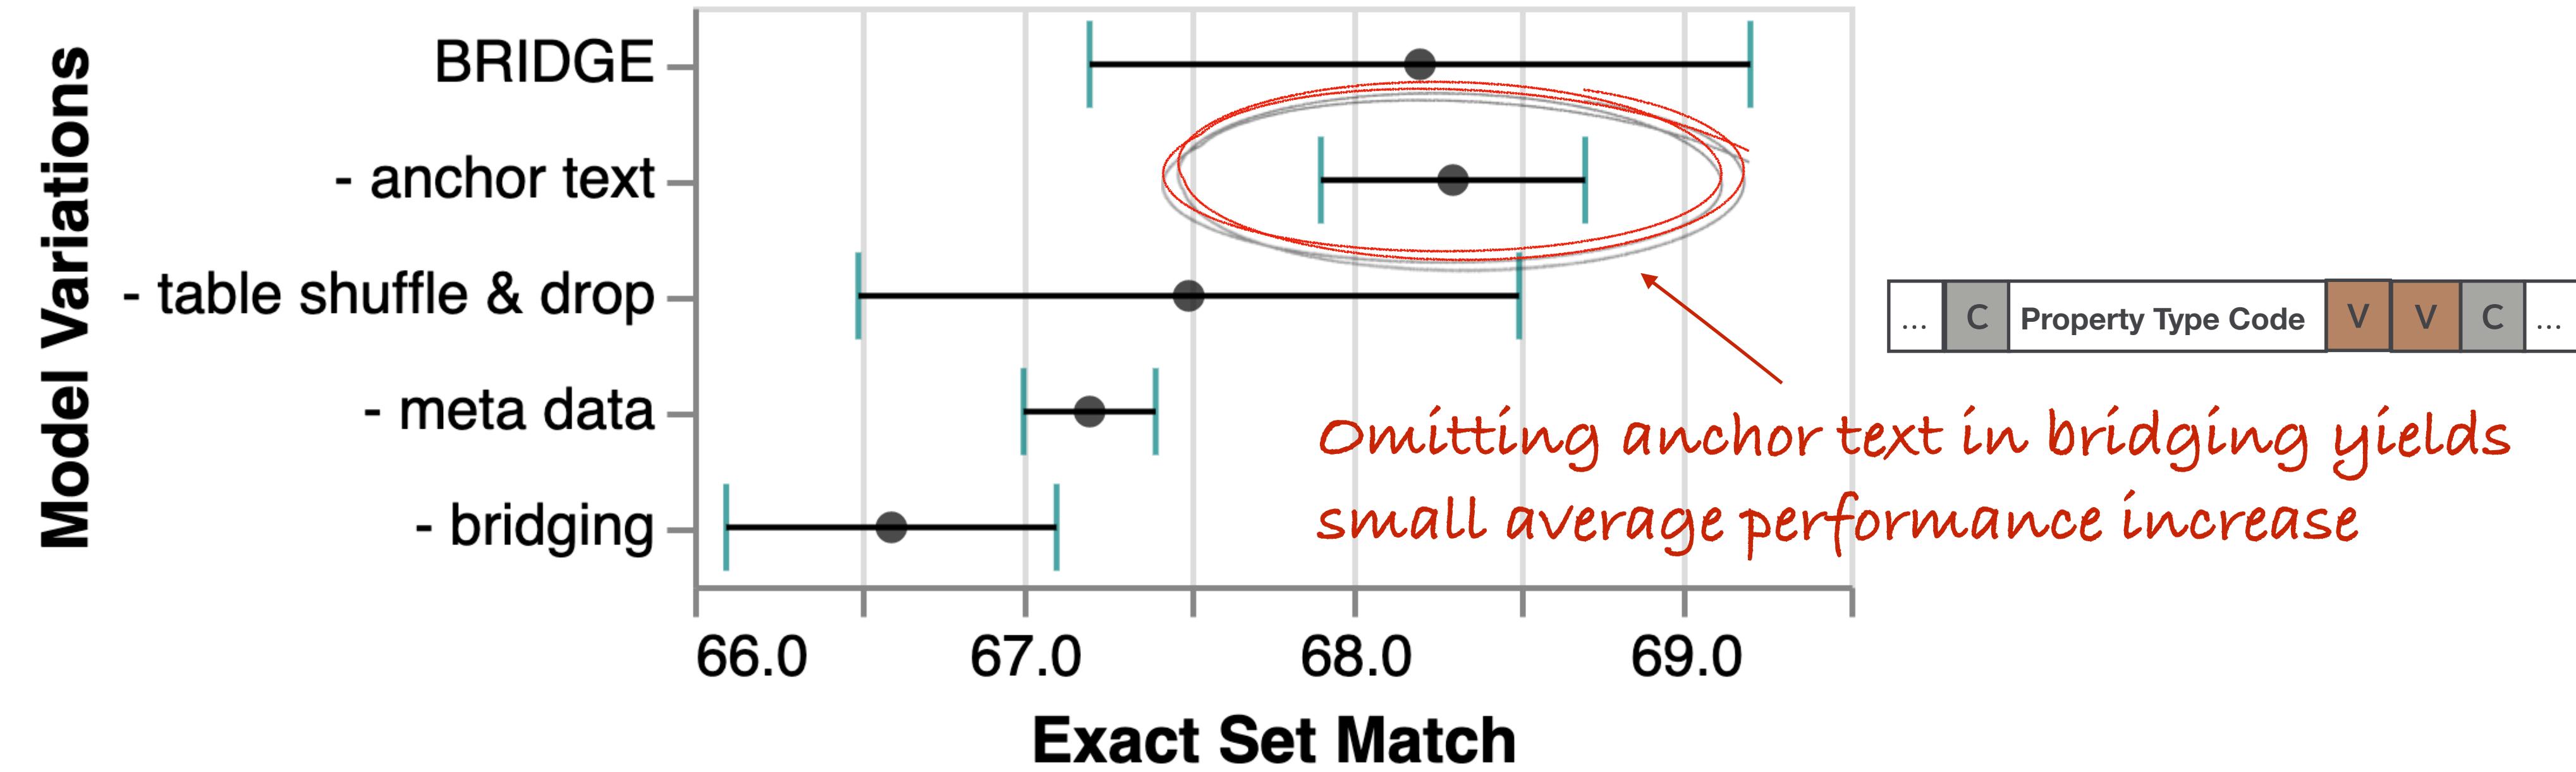
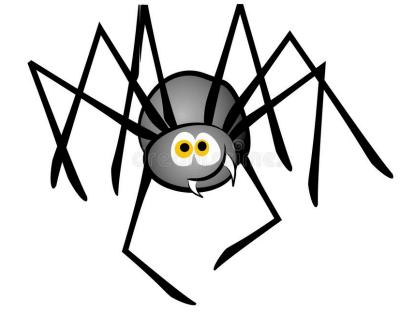


Figure 2. Ablation study of BRIDGE encoding strategies on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Ablation Study

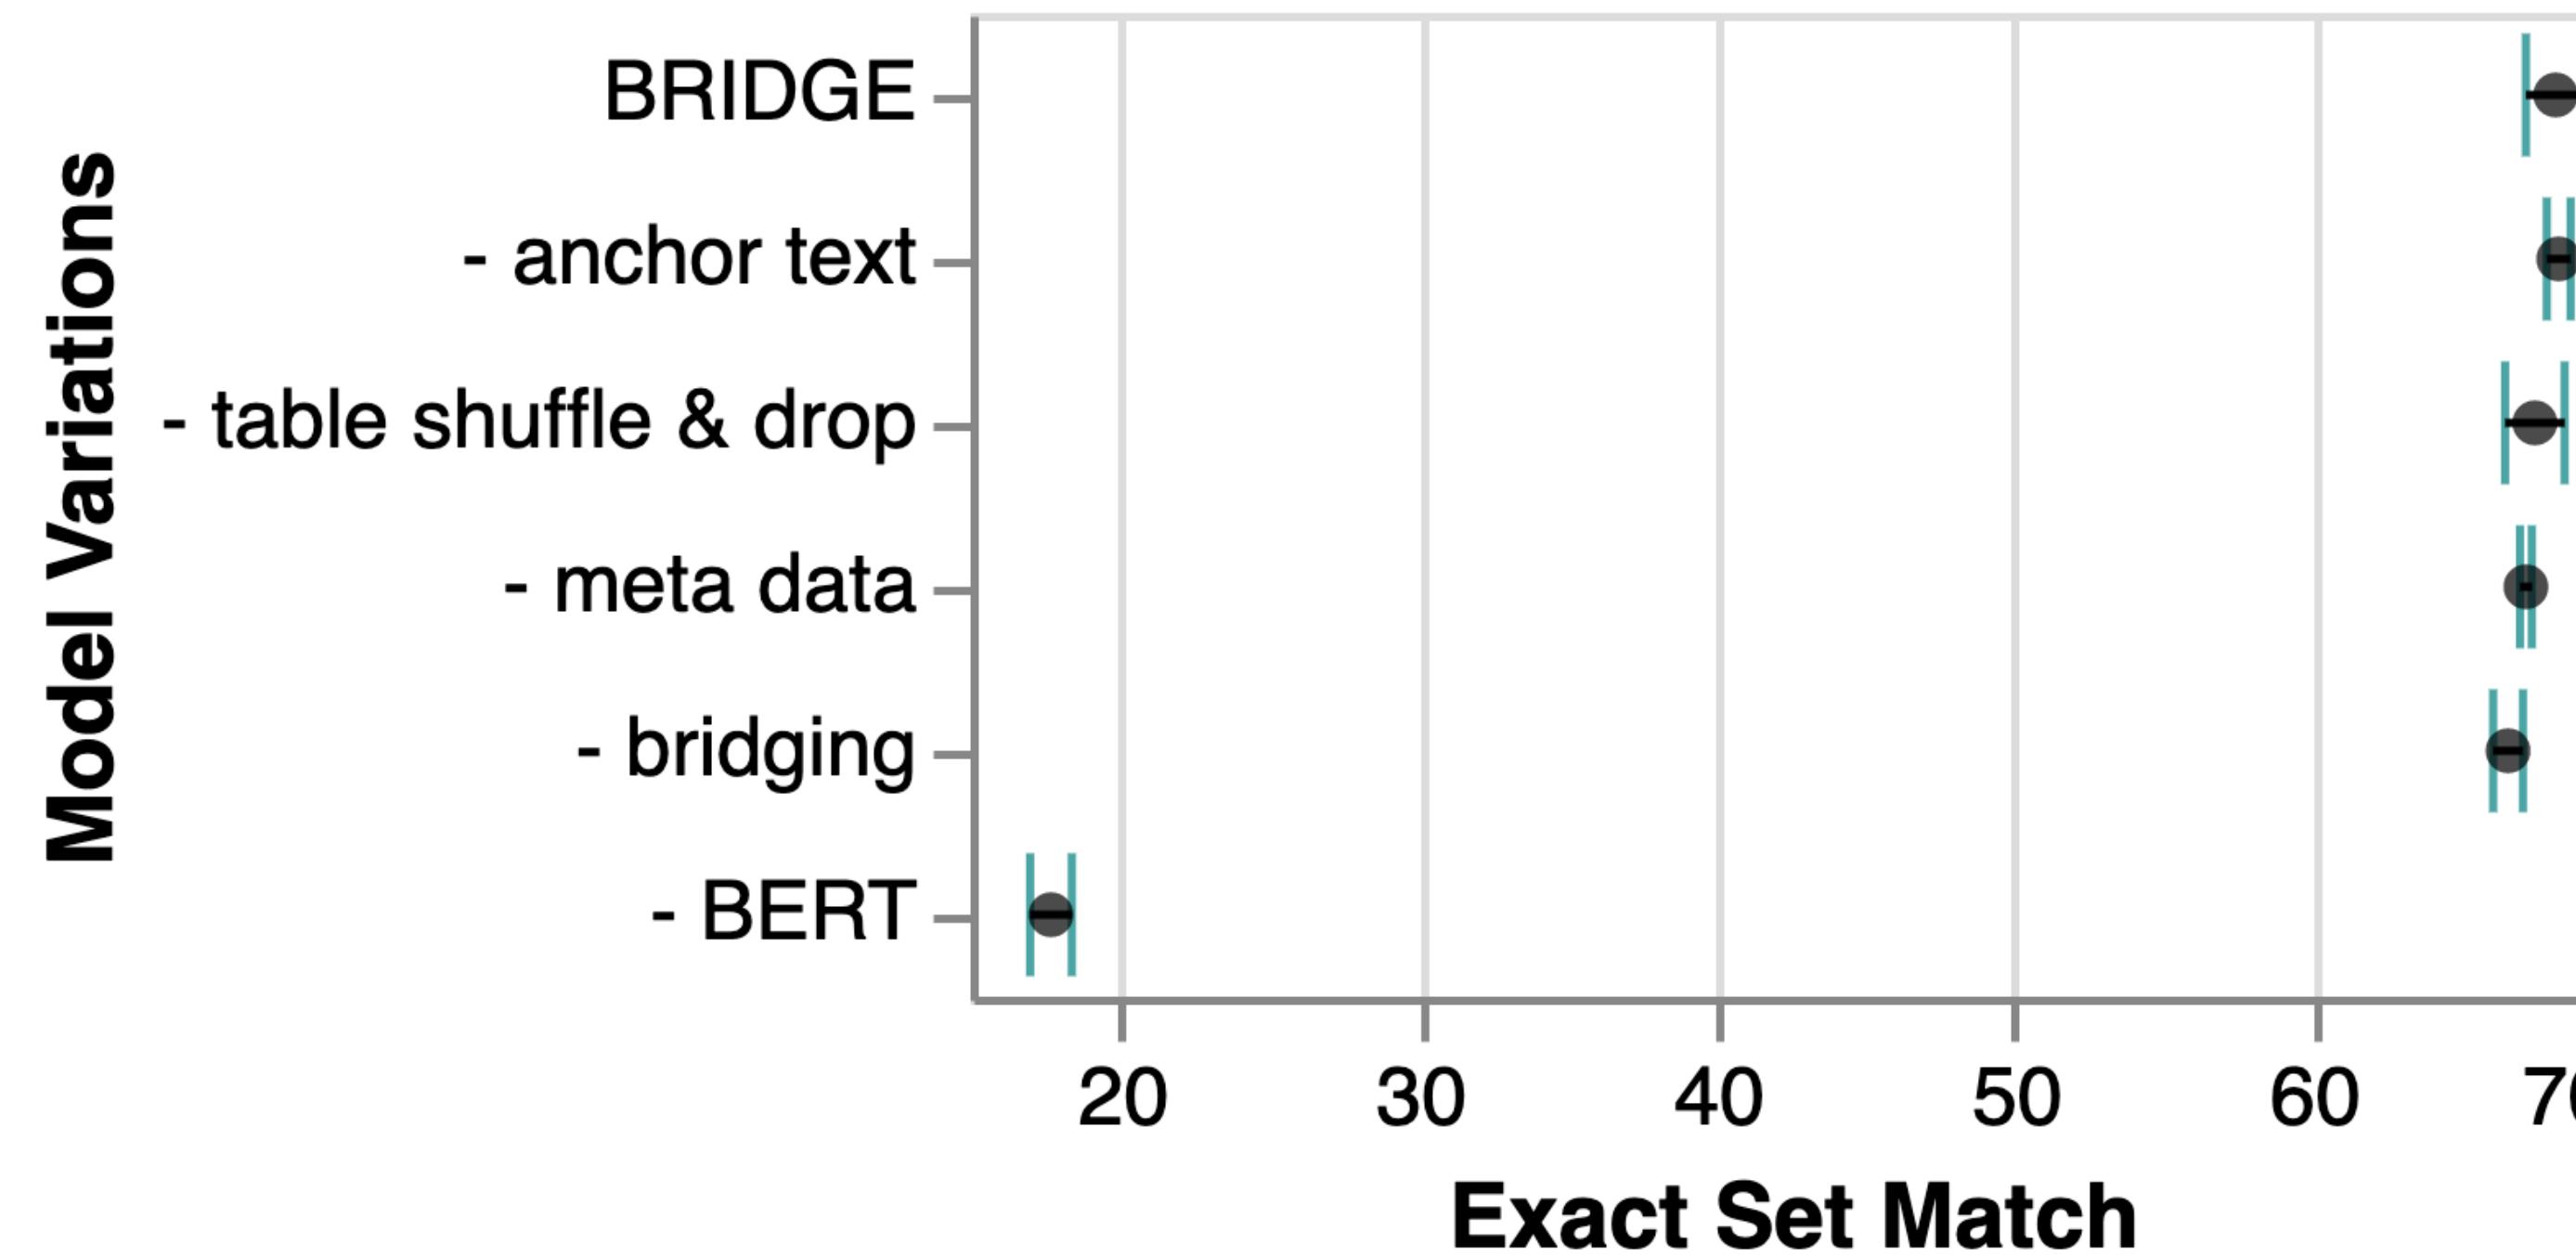
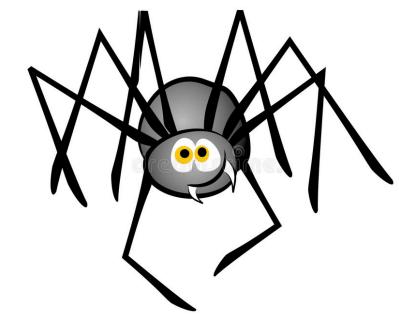
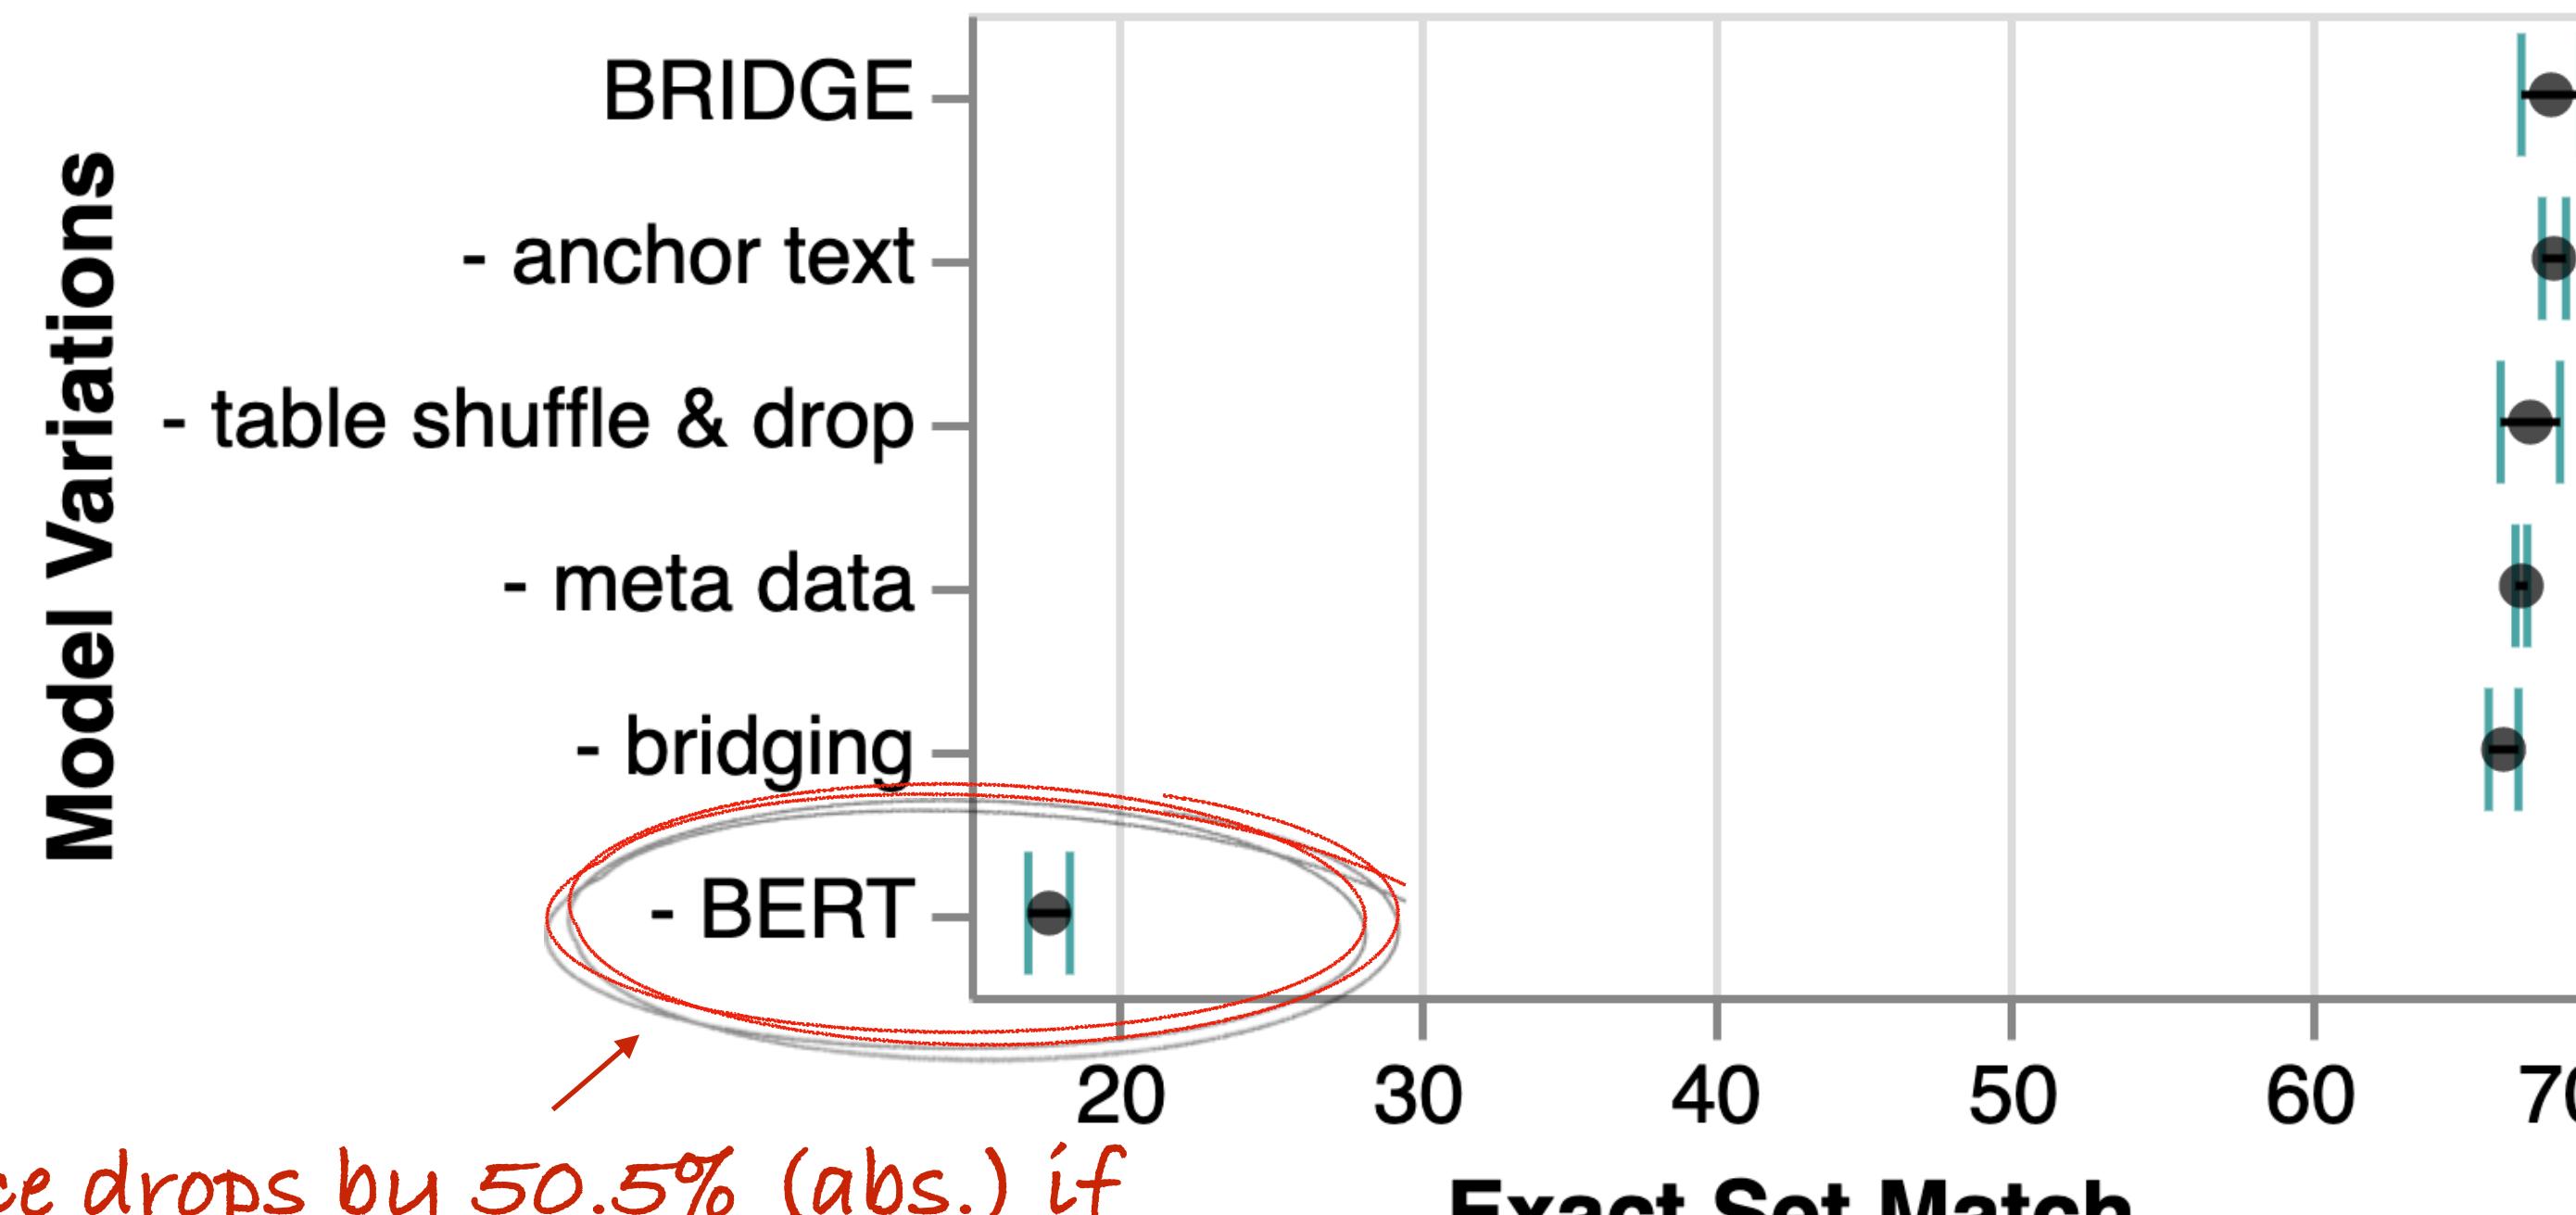
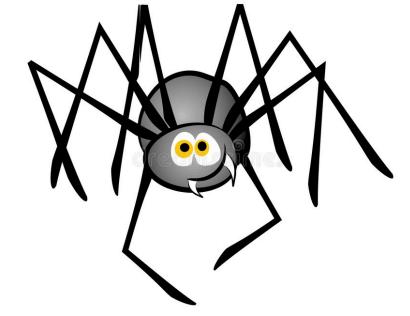


Figure 3. Ablation study of BRIDGE vs. BRIDGE - BERT on the Spider Dev set. We train 3 models using different random seeds for each model variation and average the performances.

Ablation Study



Performance drops by 50.5% (abs.) if
BERT is removed

Figure 3. Ablation study of BRIDGE vs. BRIDGE - BERT on the Spider Dev set.
We train 3 models using different random seeds for each model variation
and average the performances.



Ablation Study

Model	w/o EG		w/ EG	
	EM	EX	EM	EX
BRIDGE _L	86.2	91.7	86.8	92.6
-anchor text	84.2	90.0	85.2	91.3
-bridging	82.6	88.5	84.5	90.8

Figure 4. Ablation study of BRIDGE WikiSQL Dev set. We train only 1 model for each model variation since model variation on WikiSQL is very small. EG refers to “execution guided decoding”.

Ensemble Model

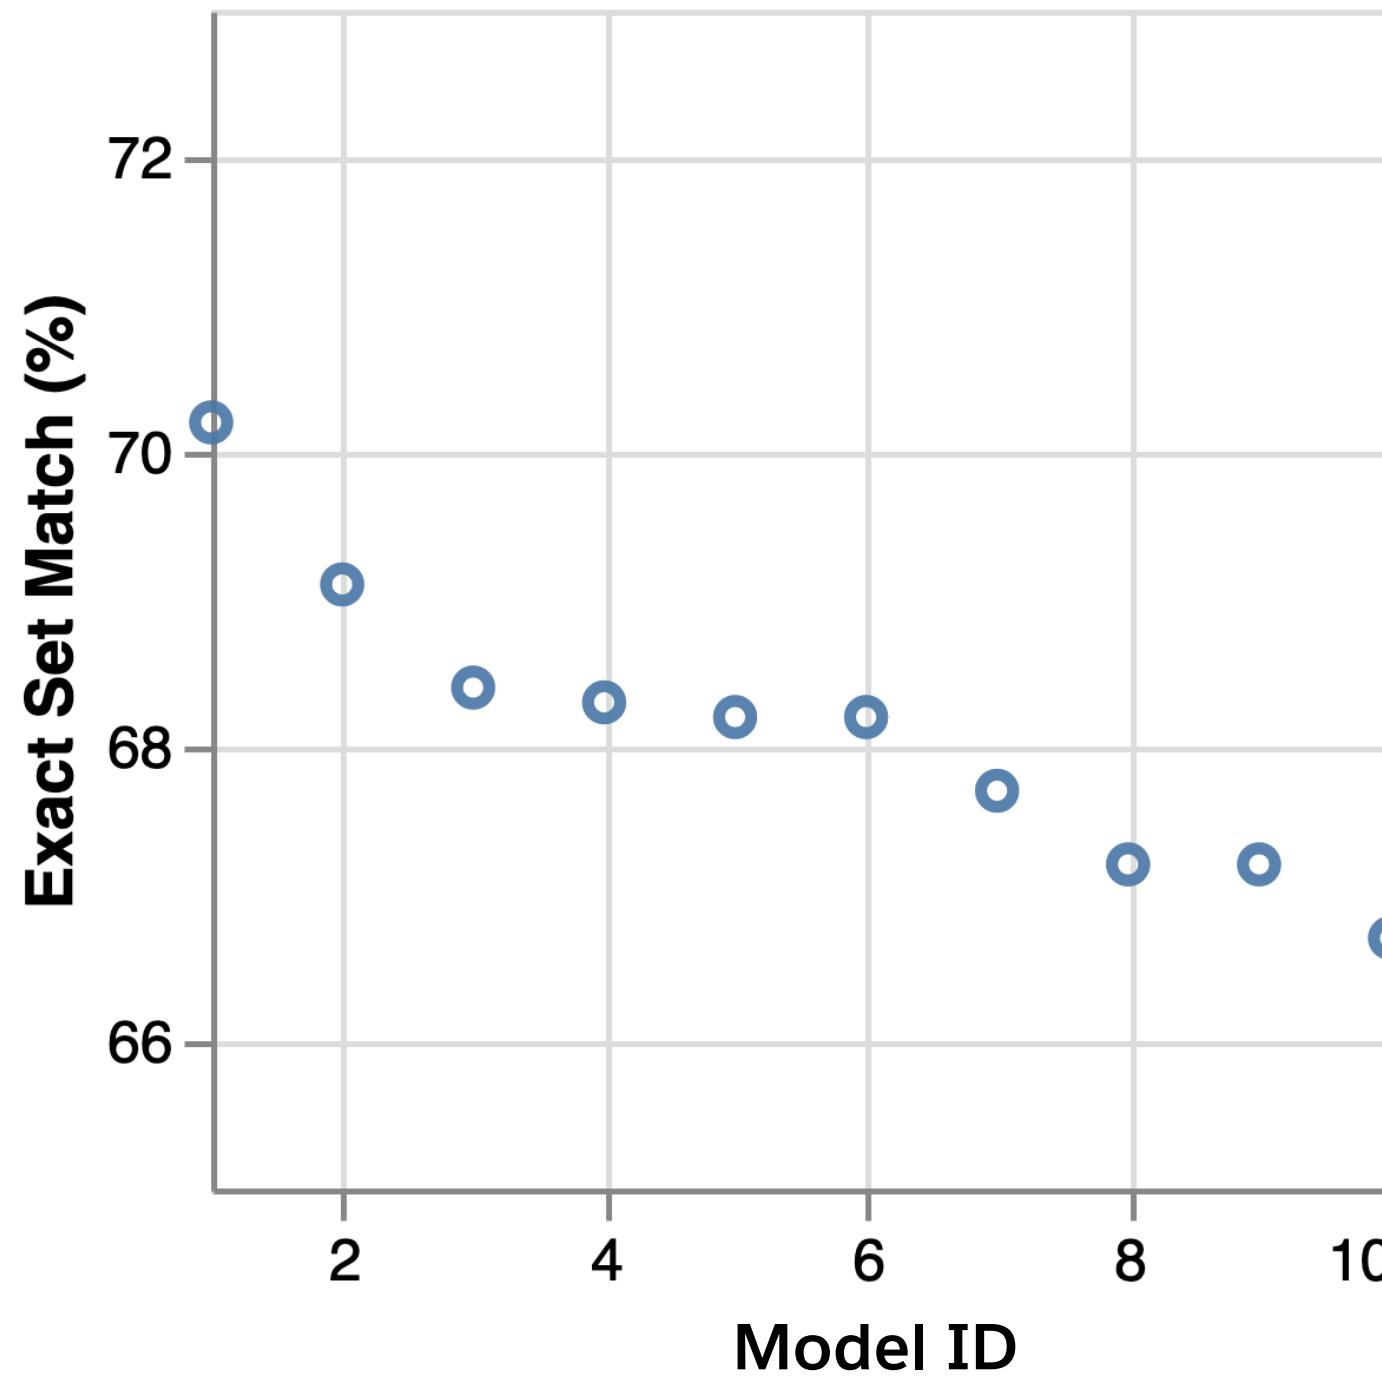
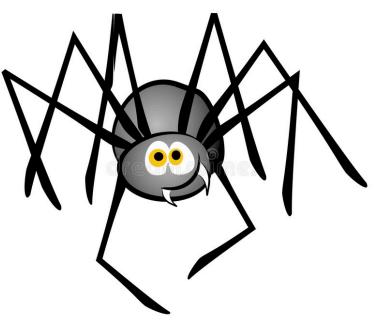


Figure 5. Performance comparison of BRIDGE model trained using 10 different random seeds on the Spider dev set.

Ensemble Model

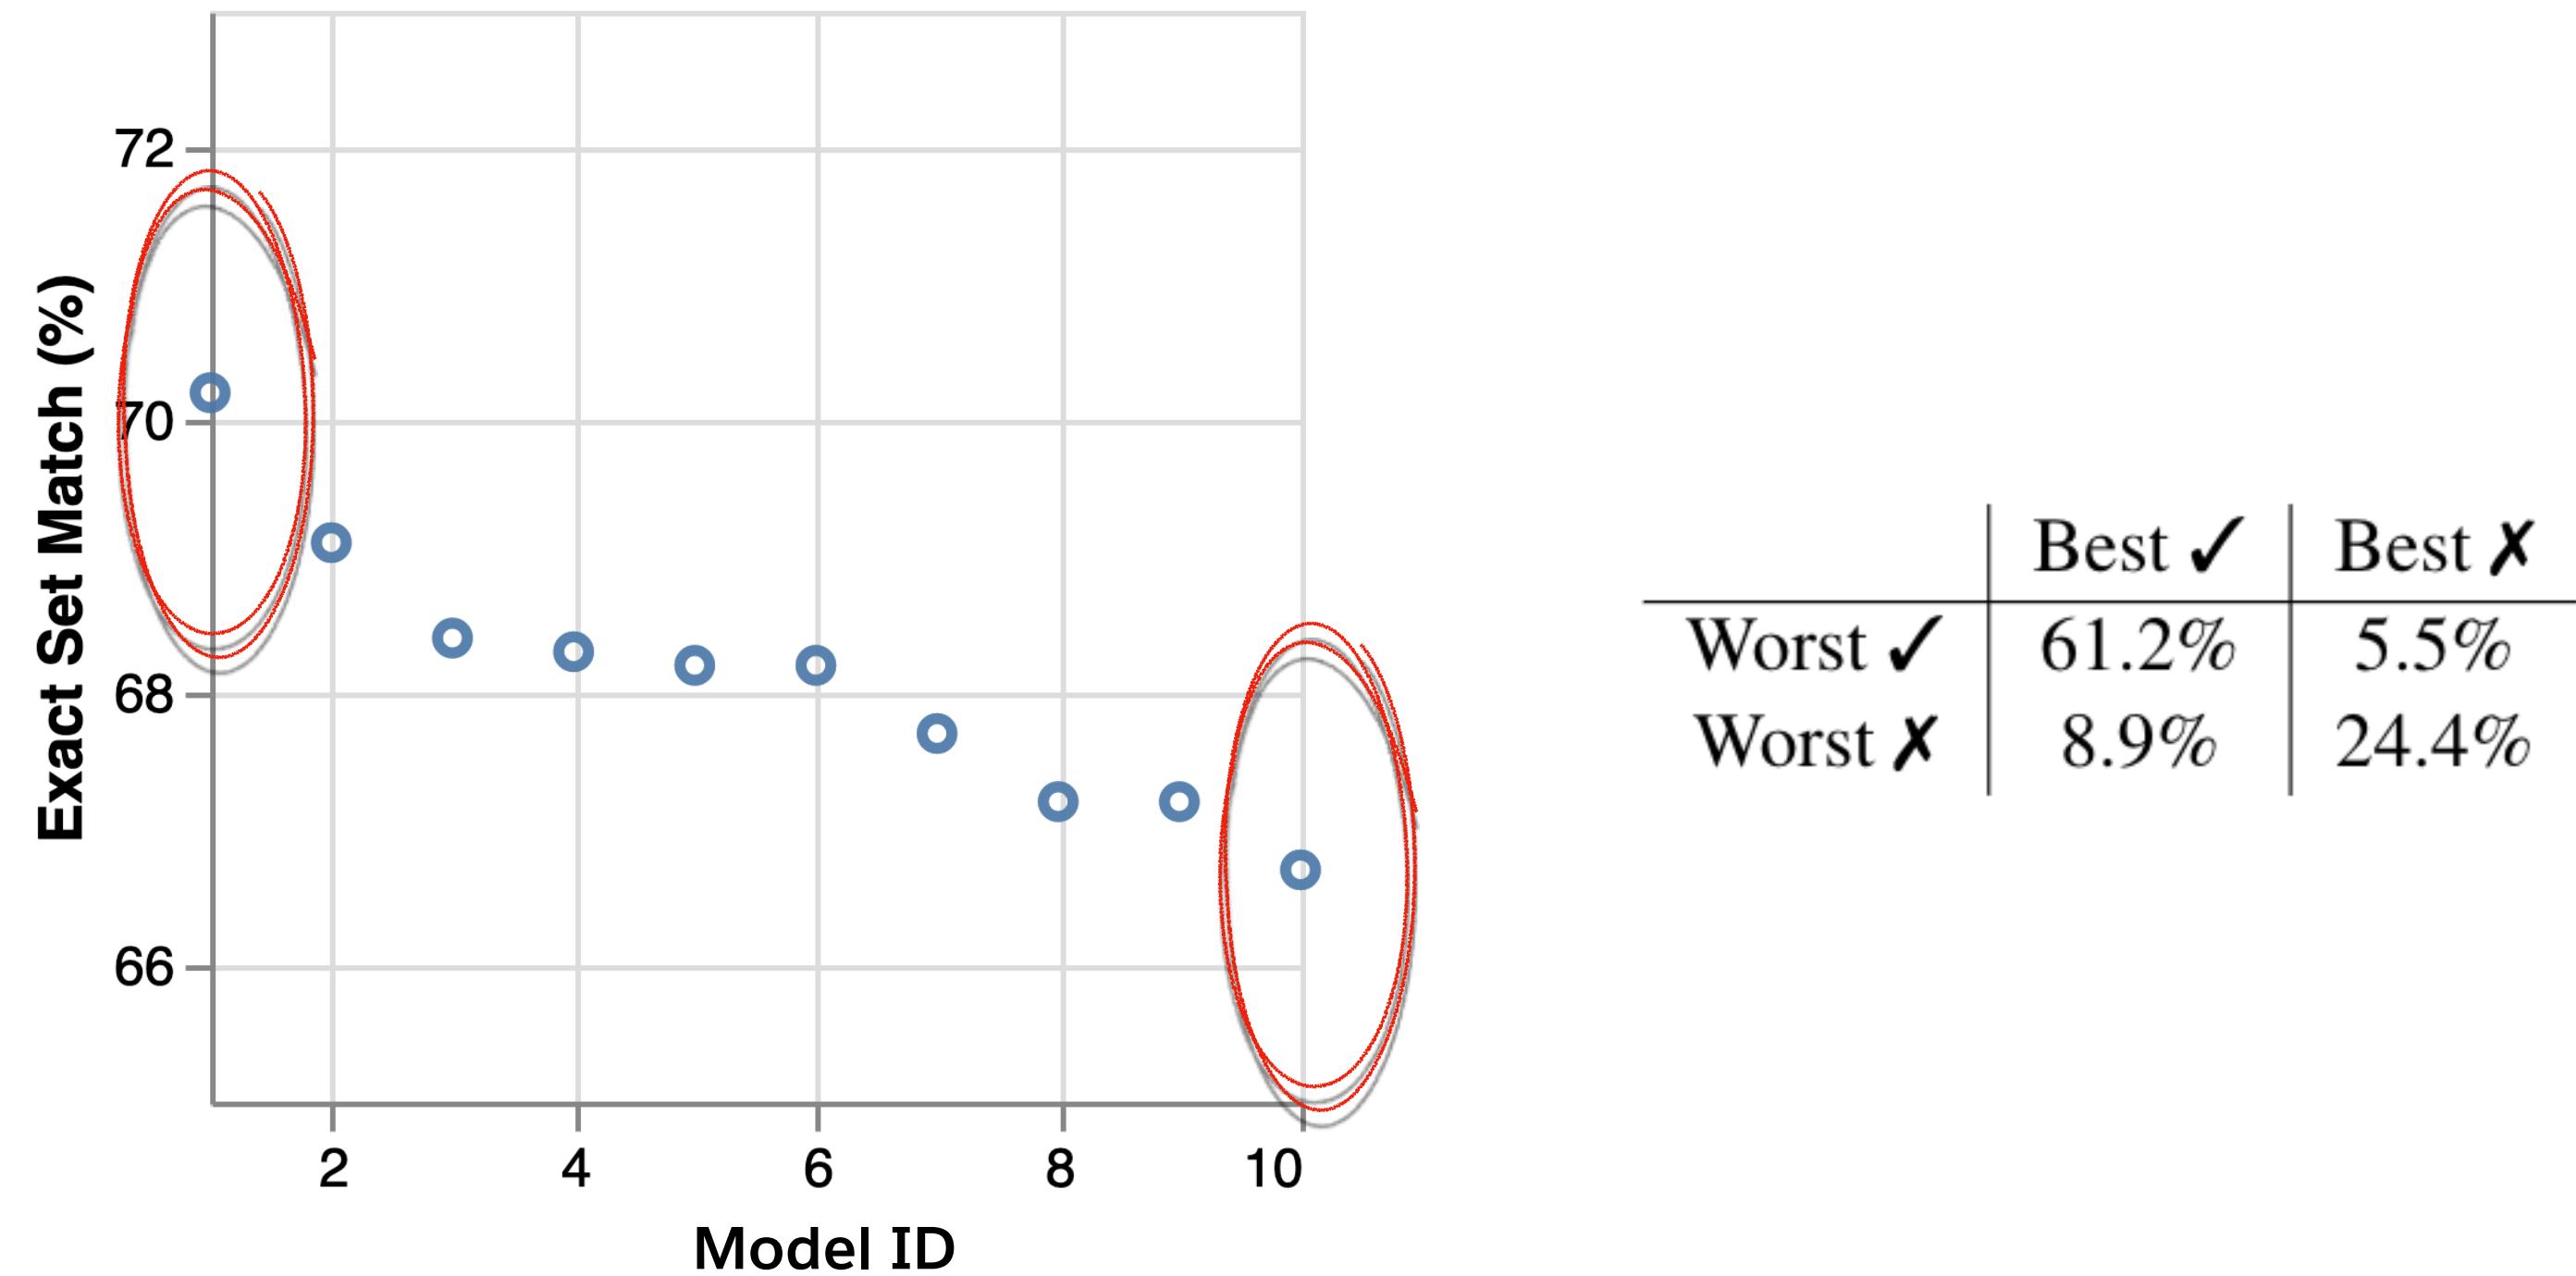
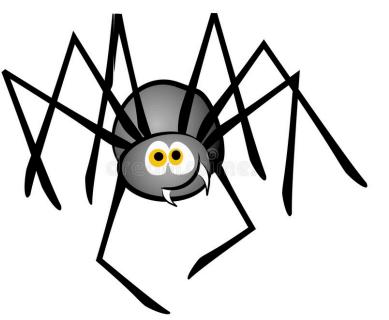


Figure 5. Performance comparison of BRIDGE model trained using 10 different random seeds on the Spider dev set.

Ensemble Model

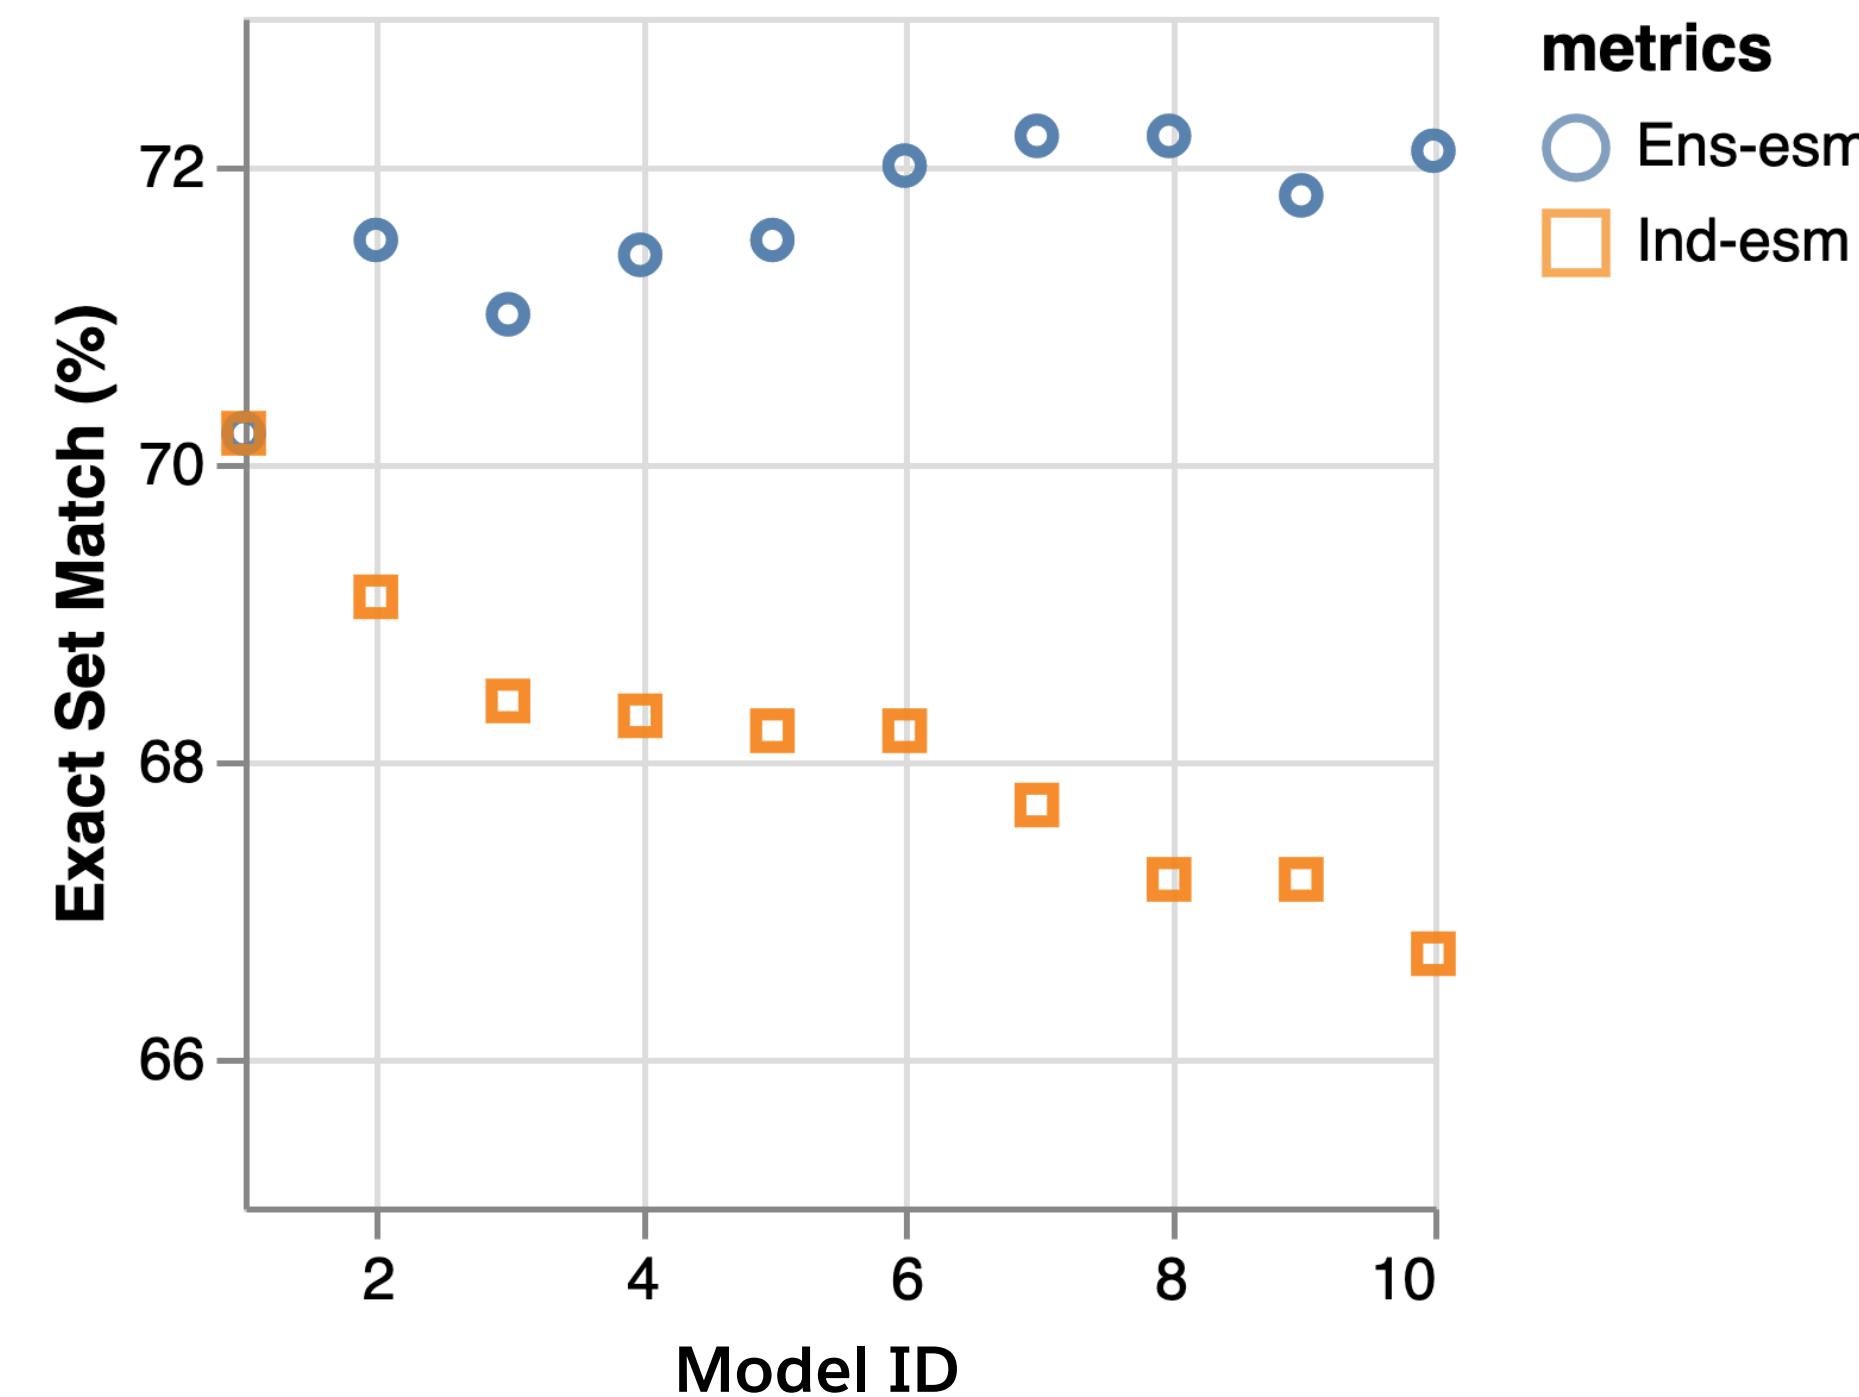
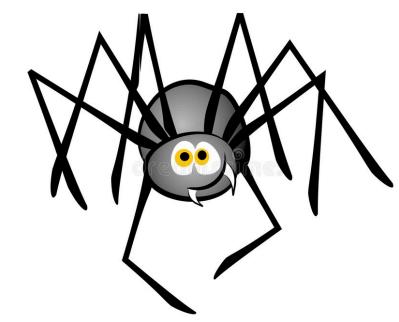
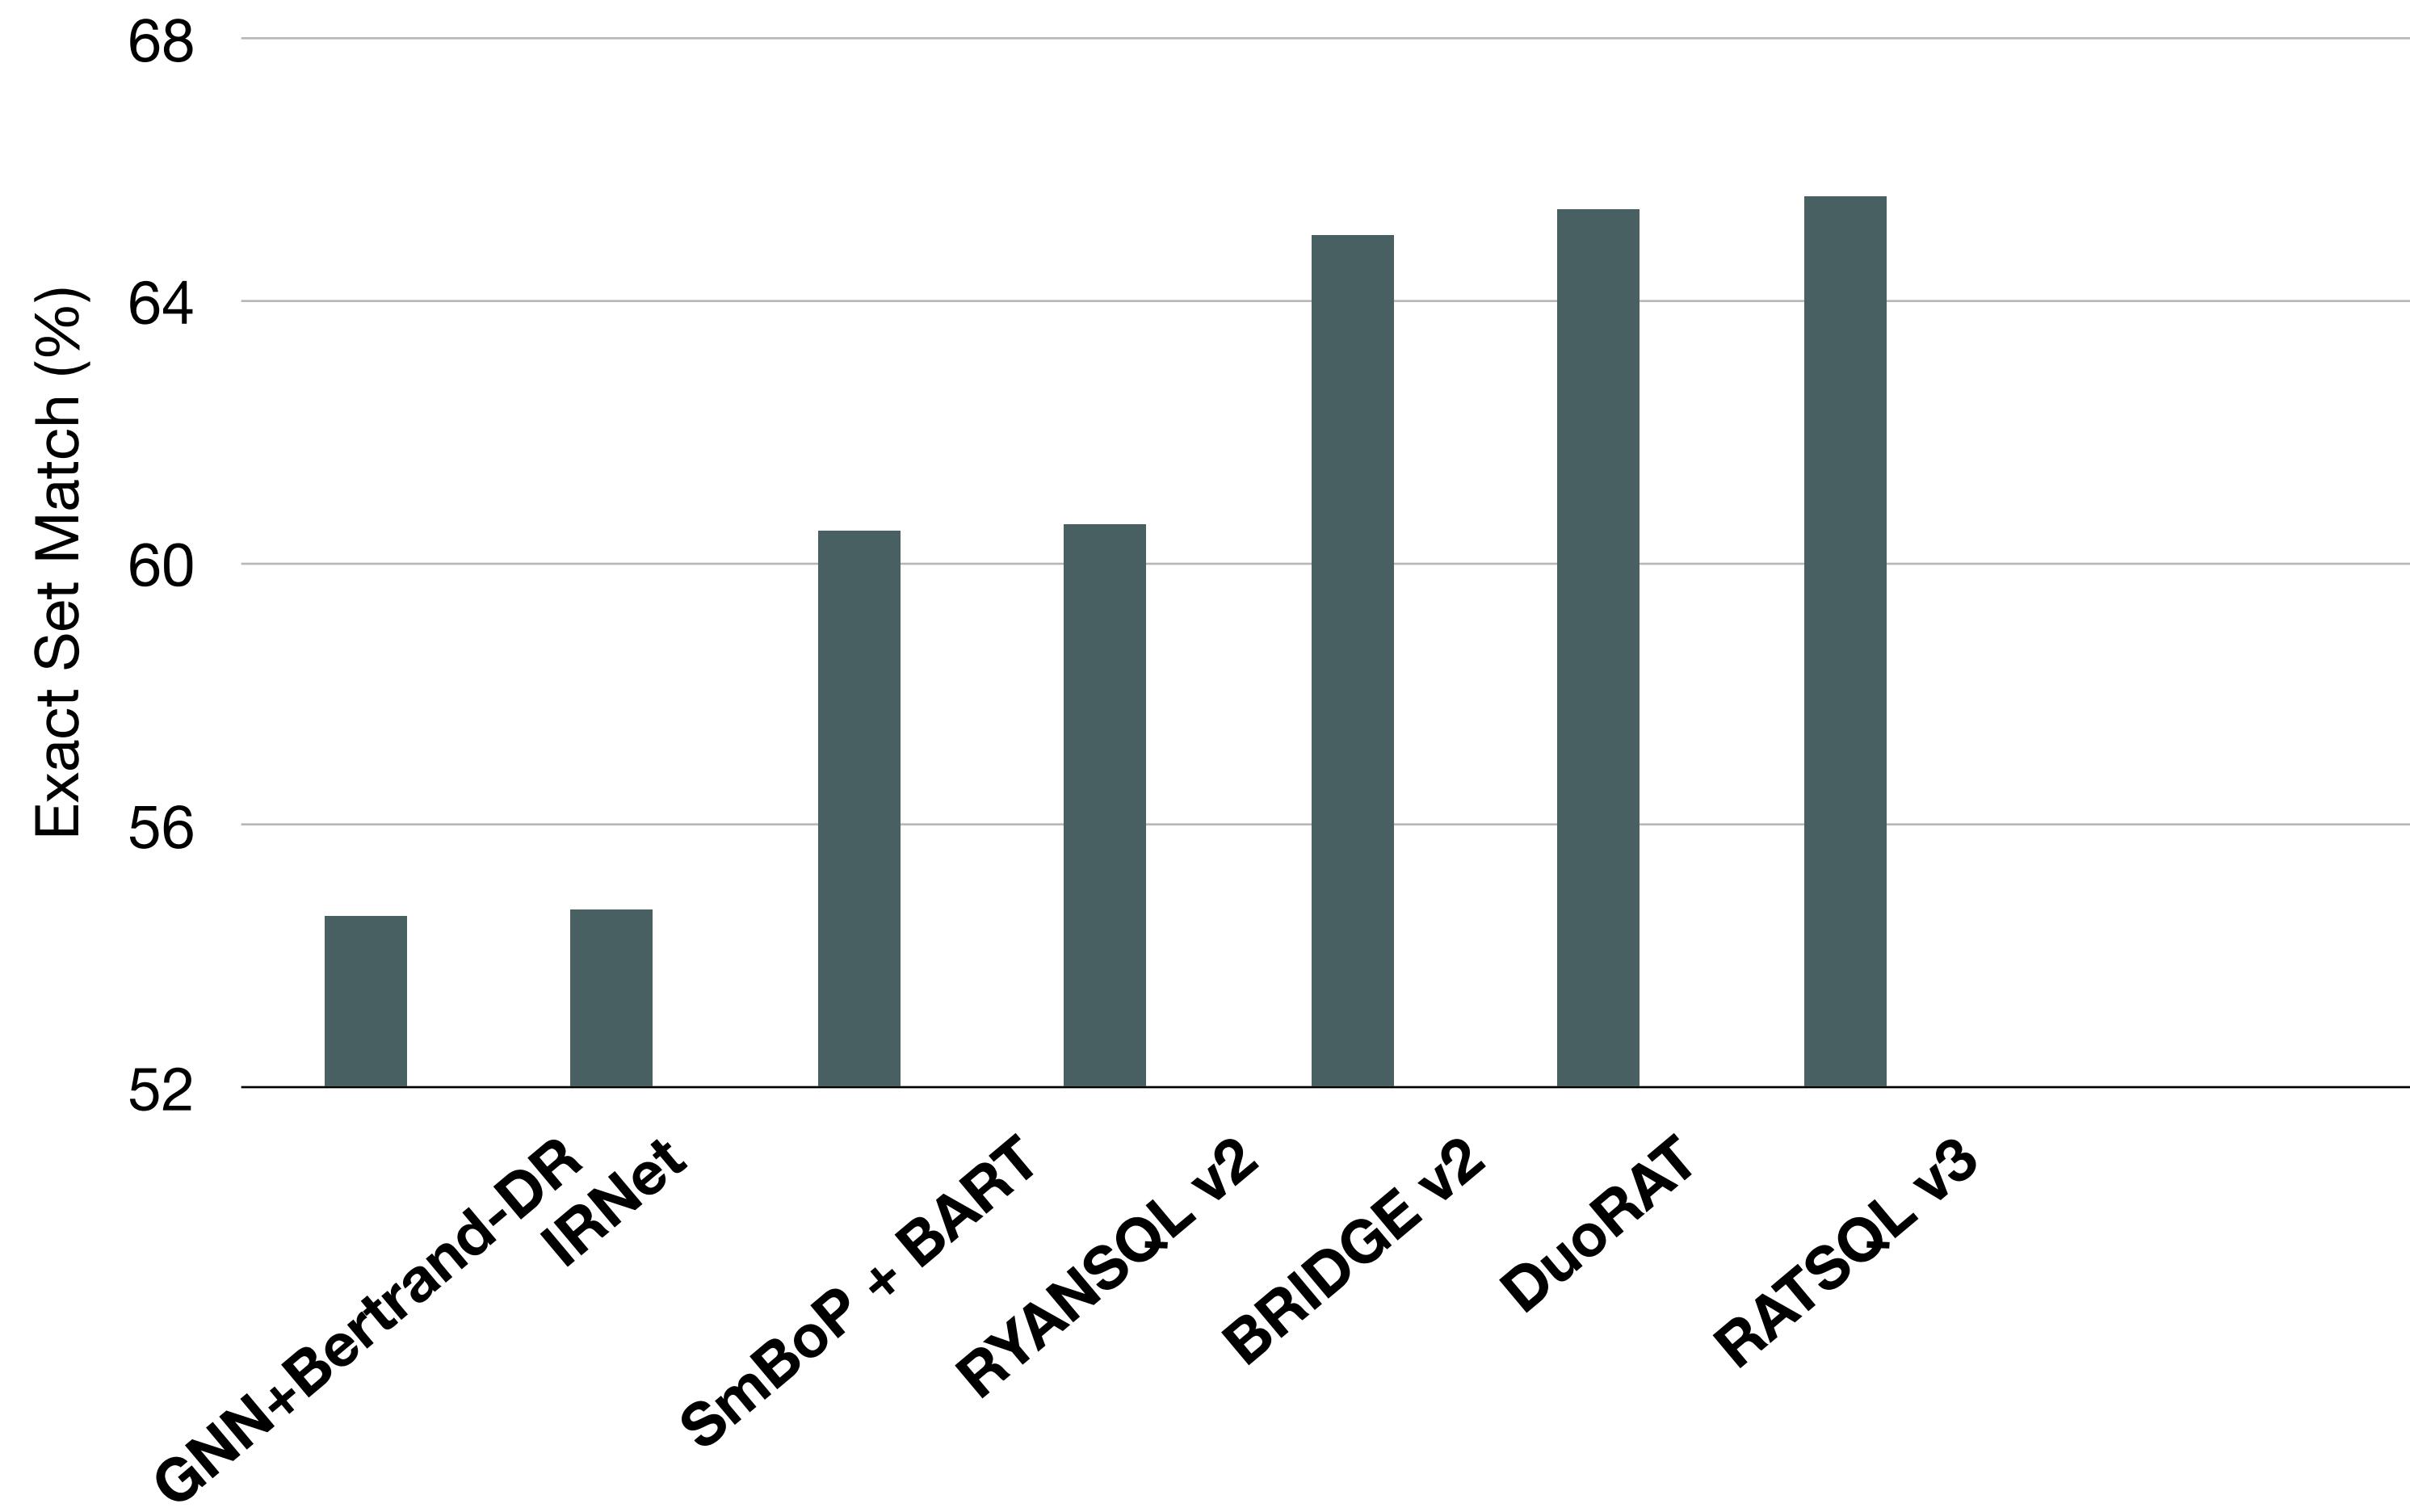


Figure 6. Performance of model ensemble (using step-wise output distribution average) on the Spider dev set.

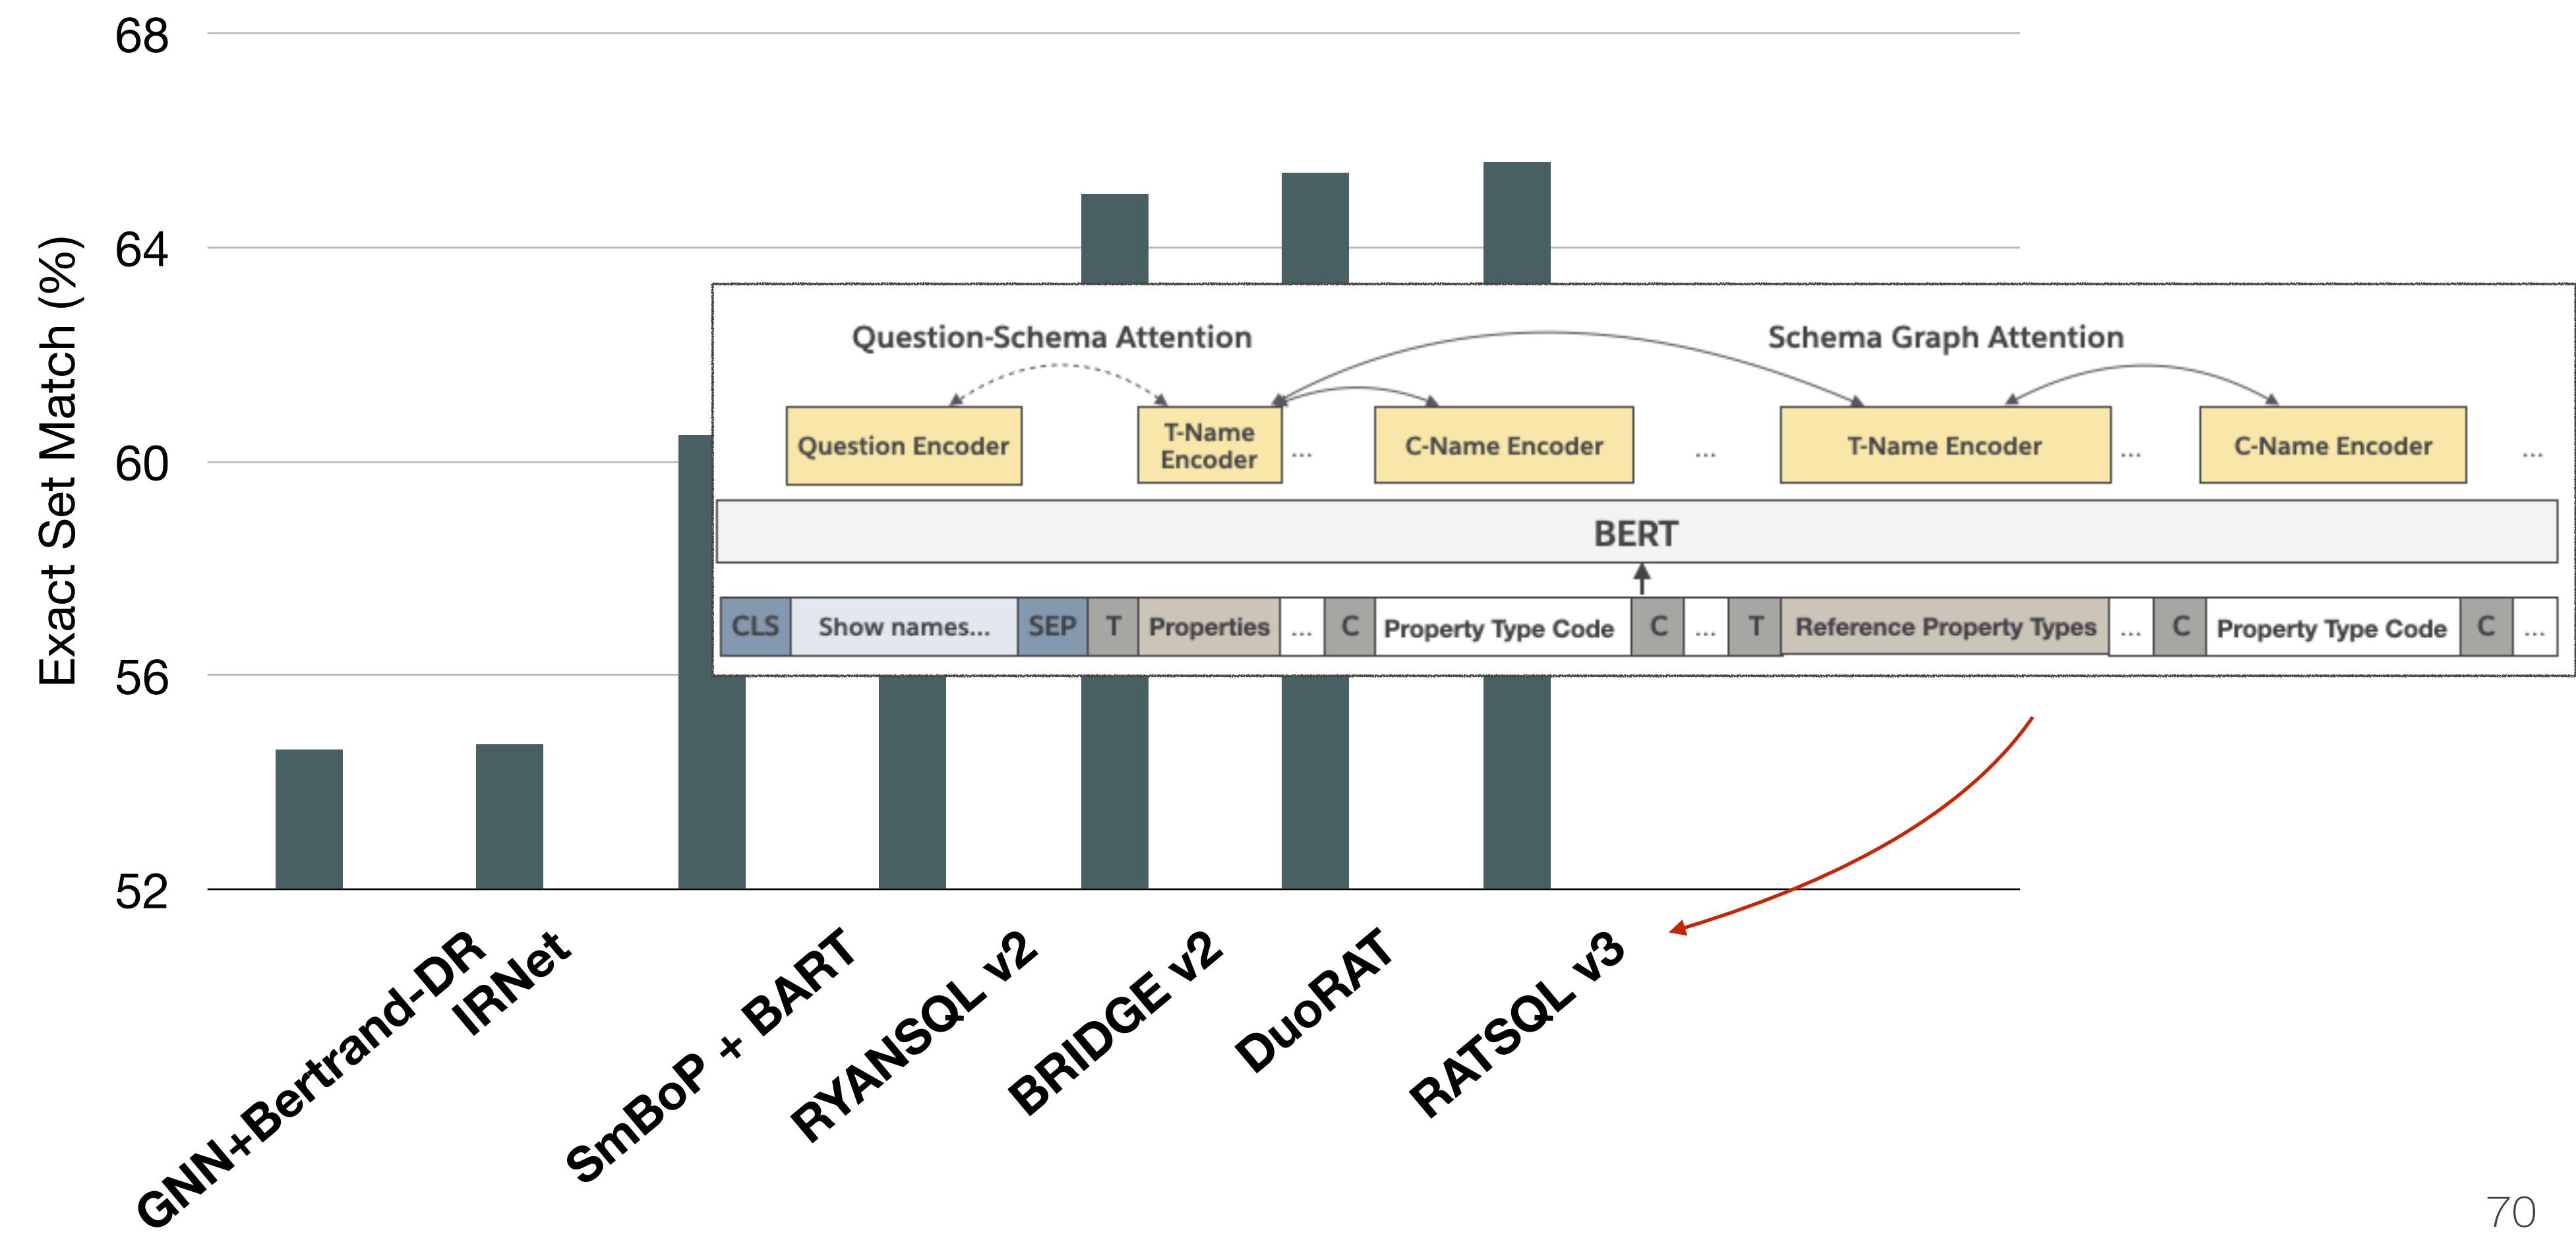
Performance on Spider Leaderboard

Figure 7. Comparison to other top-performing text-to-SQL models on the Spider leaderboard (Jan 31, 2021).



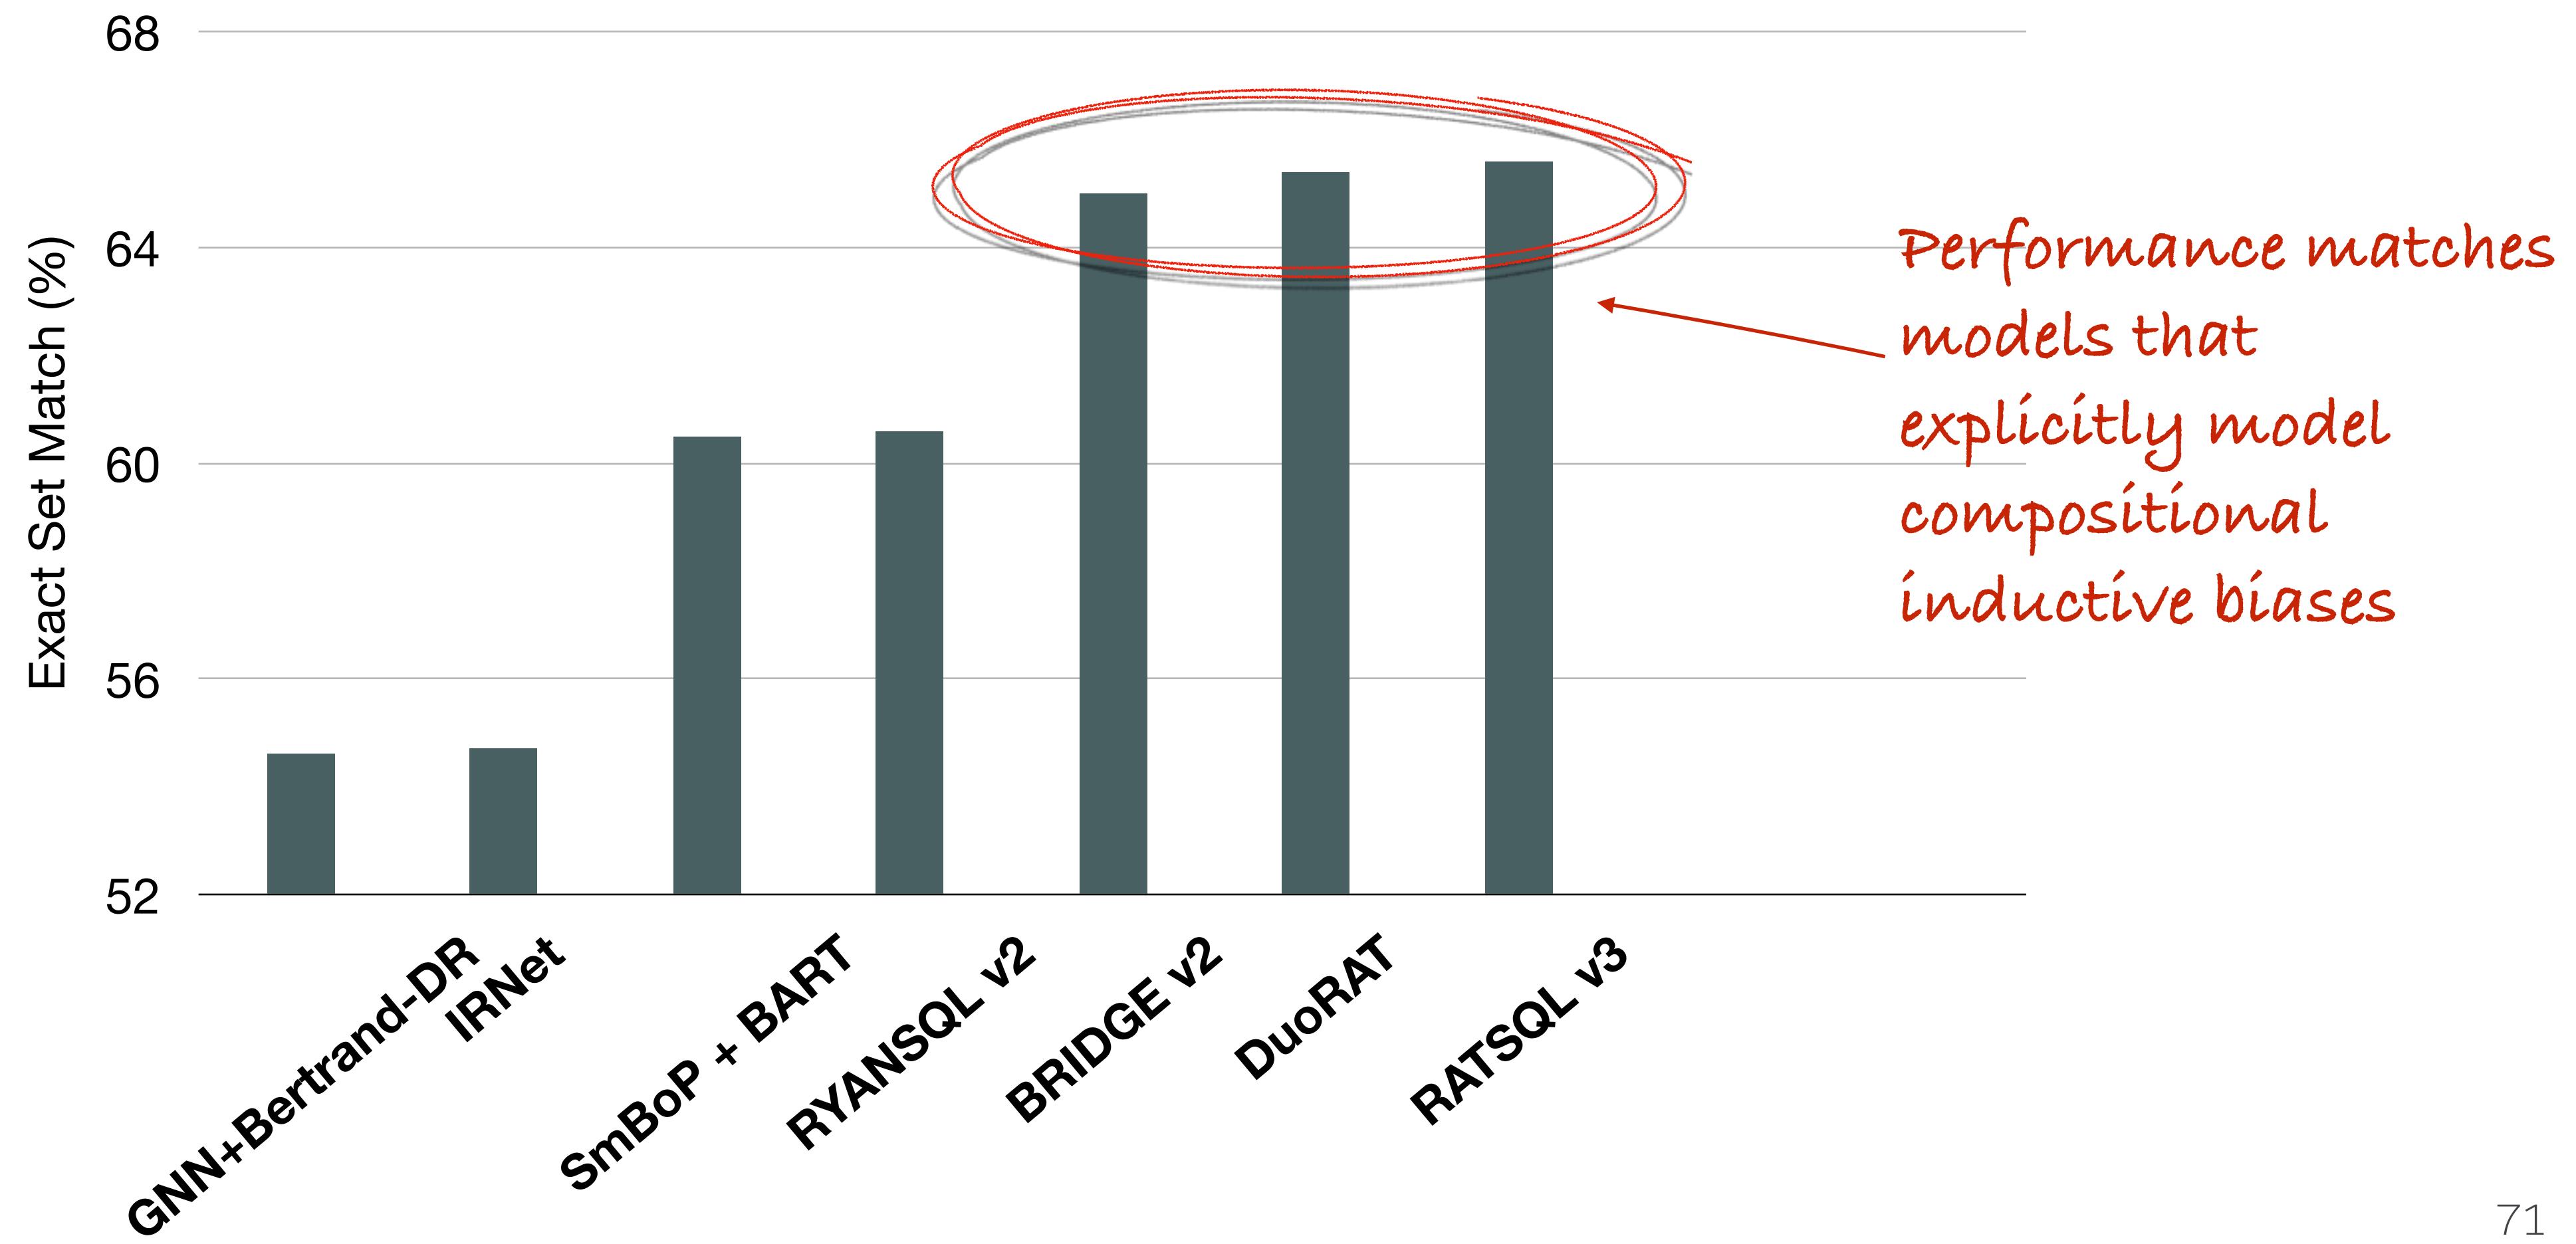
Performance on Spider Leaderboard

Figure 7. Comparison to other top-performing text-to-SQL models on the Spider leaderboard (Jan 31, 2021).



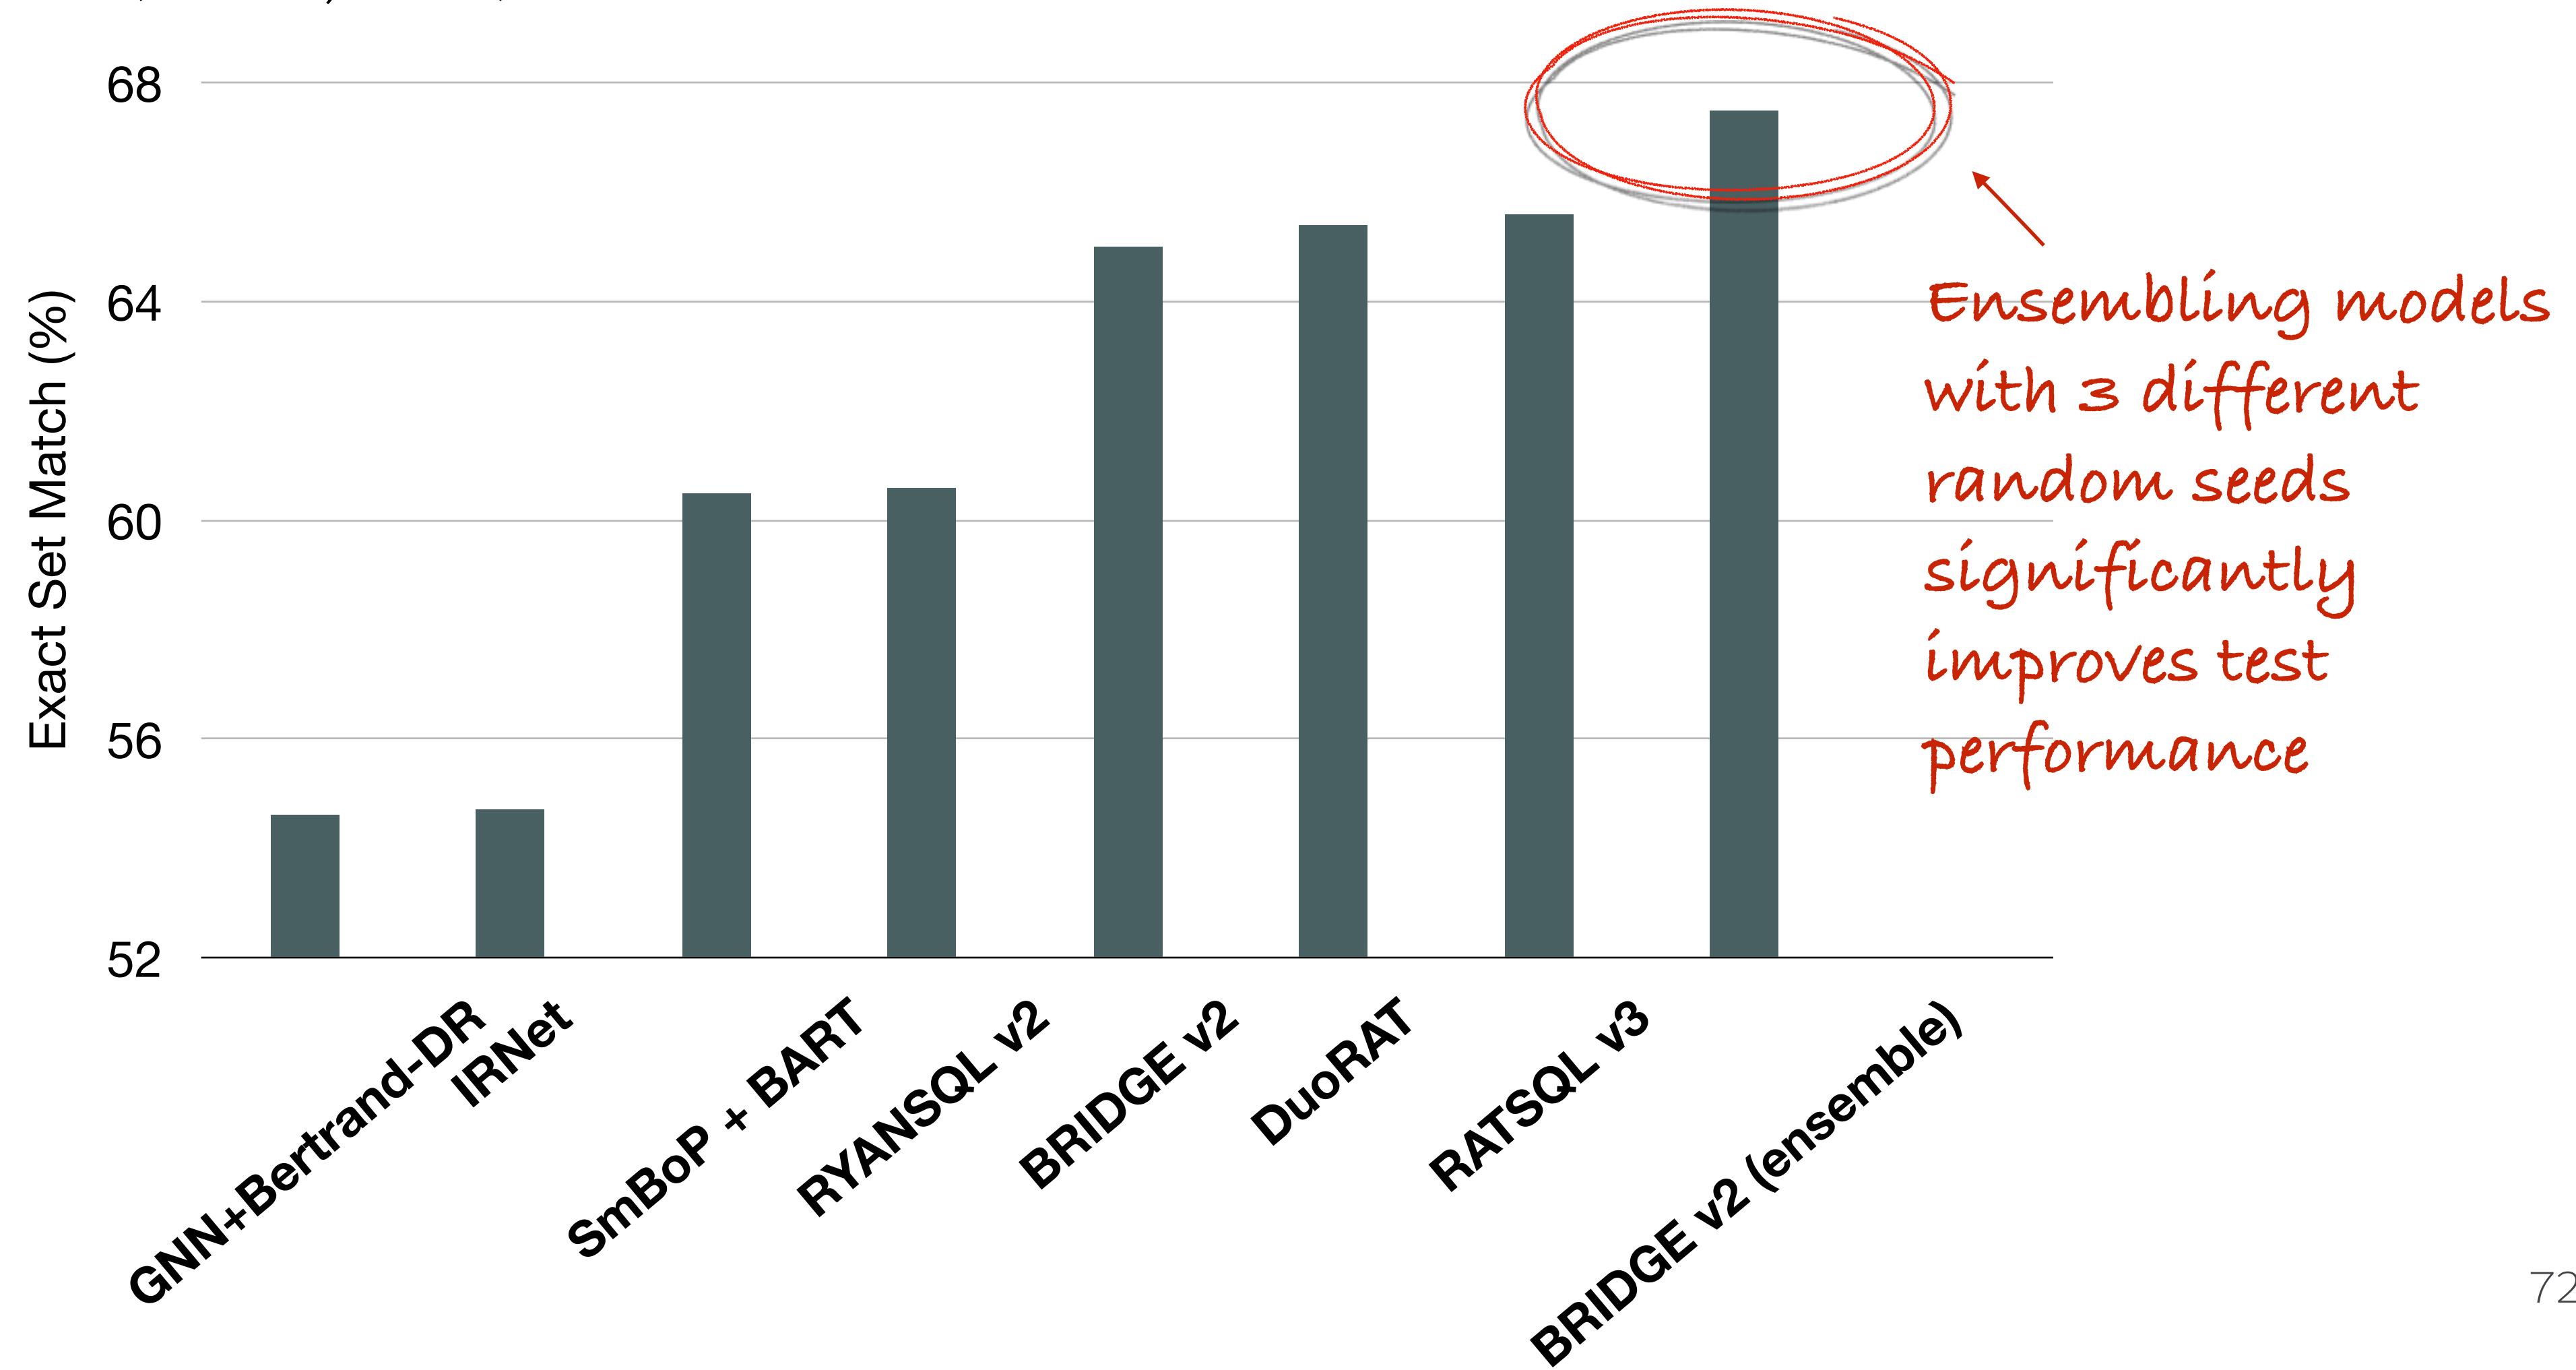
Performance on Spider Leaderboard

Figure 7. Comparison to other top-performing text-to-SQL models on the Spider leaderboard (Jan 31, 2021).



Performance on Spider Leaderboard

Figure 7. Comparison to other top-performing text-to-SQL models on the Spider leaderboard (Jan 31, 2021).



Performance Comparison by Difficulty Level

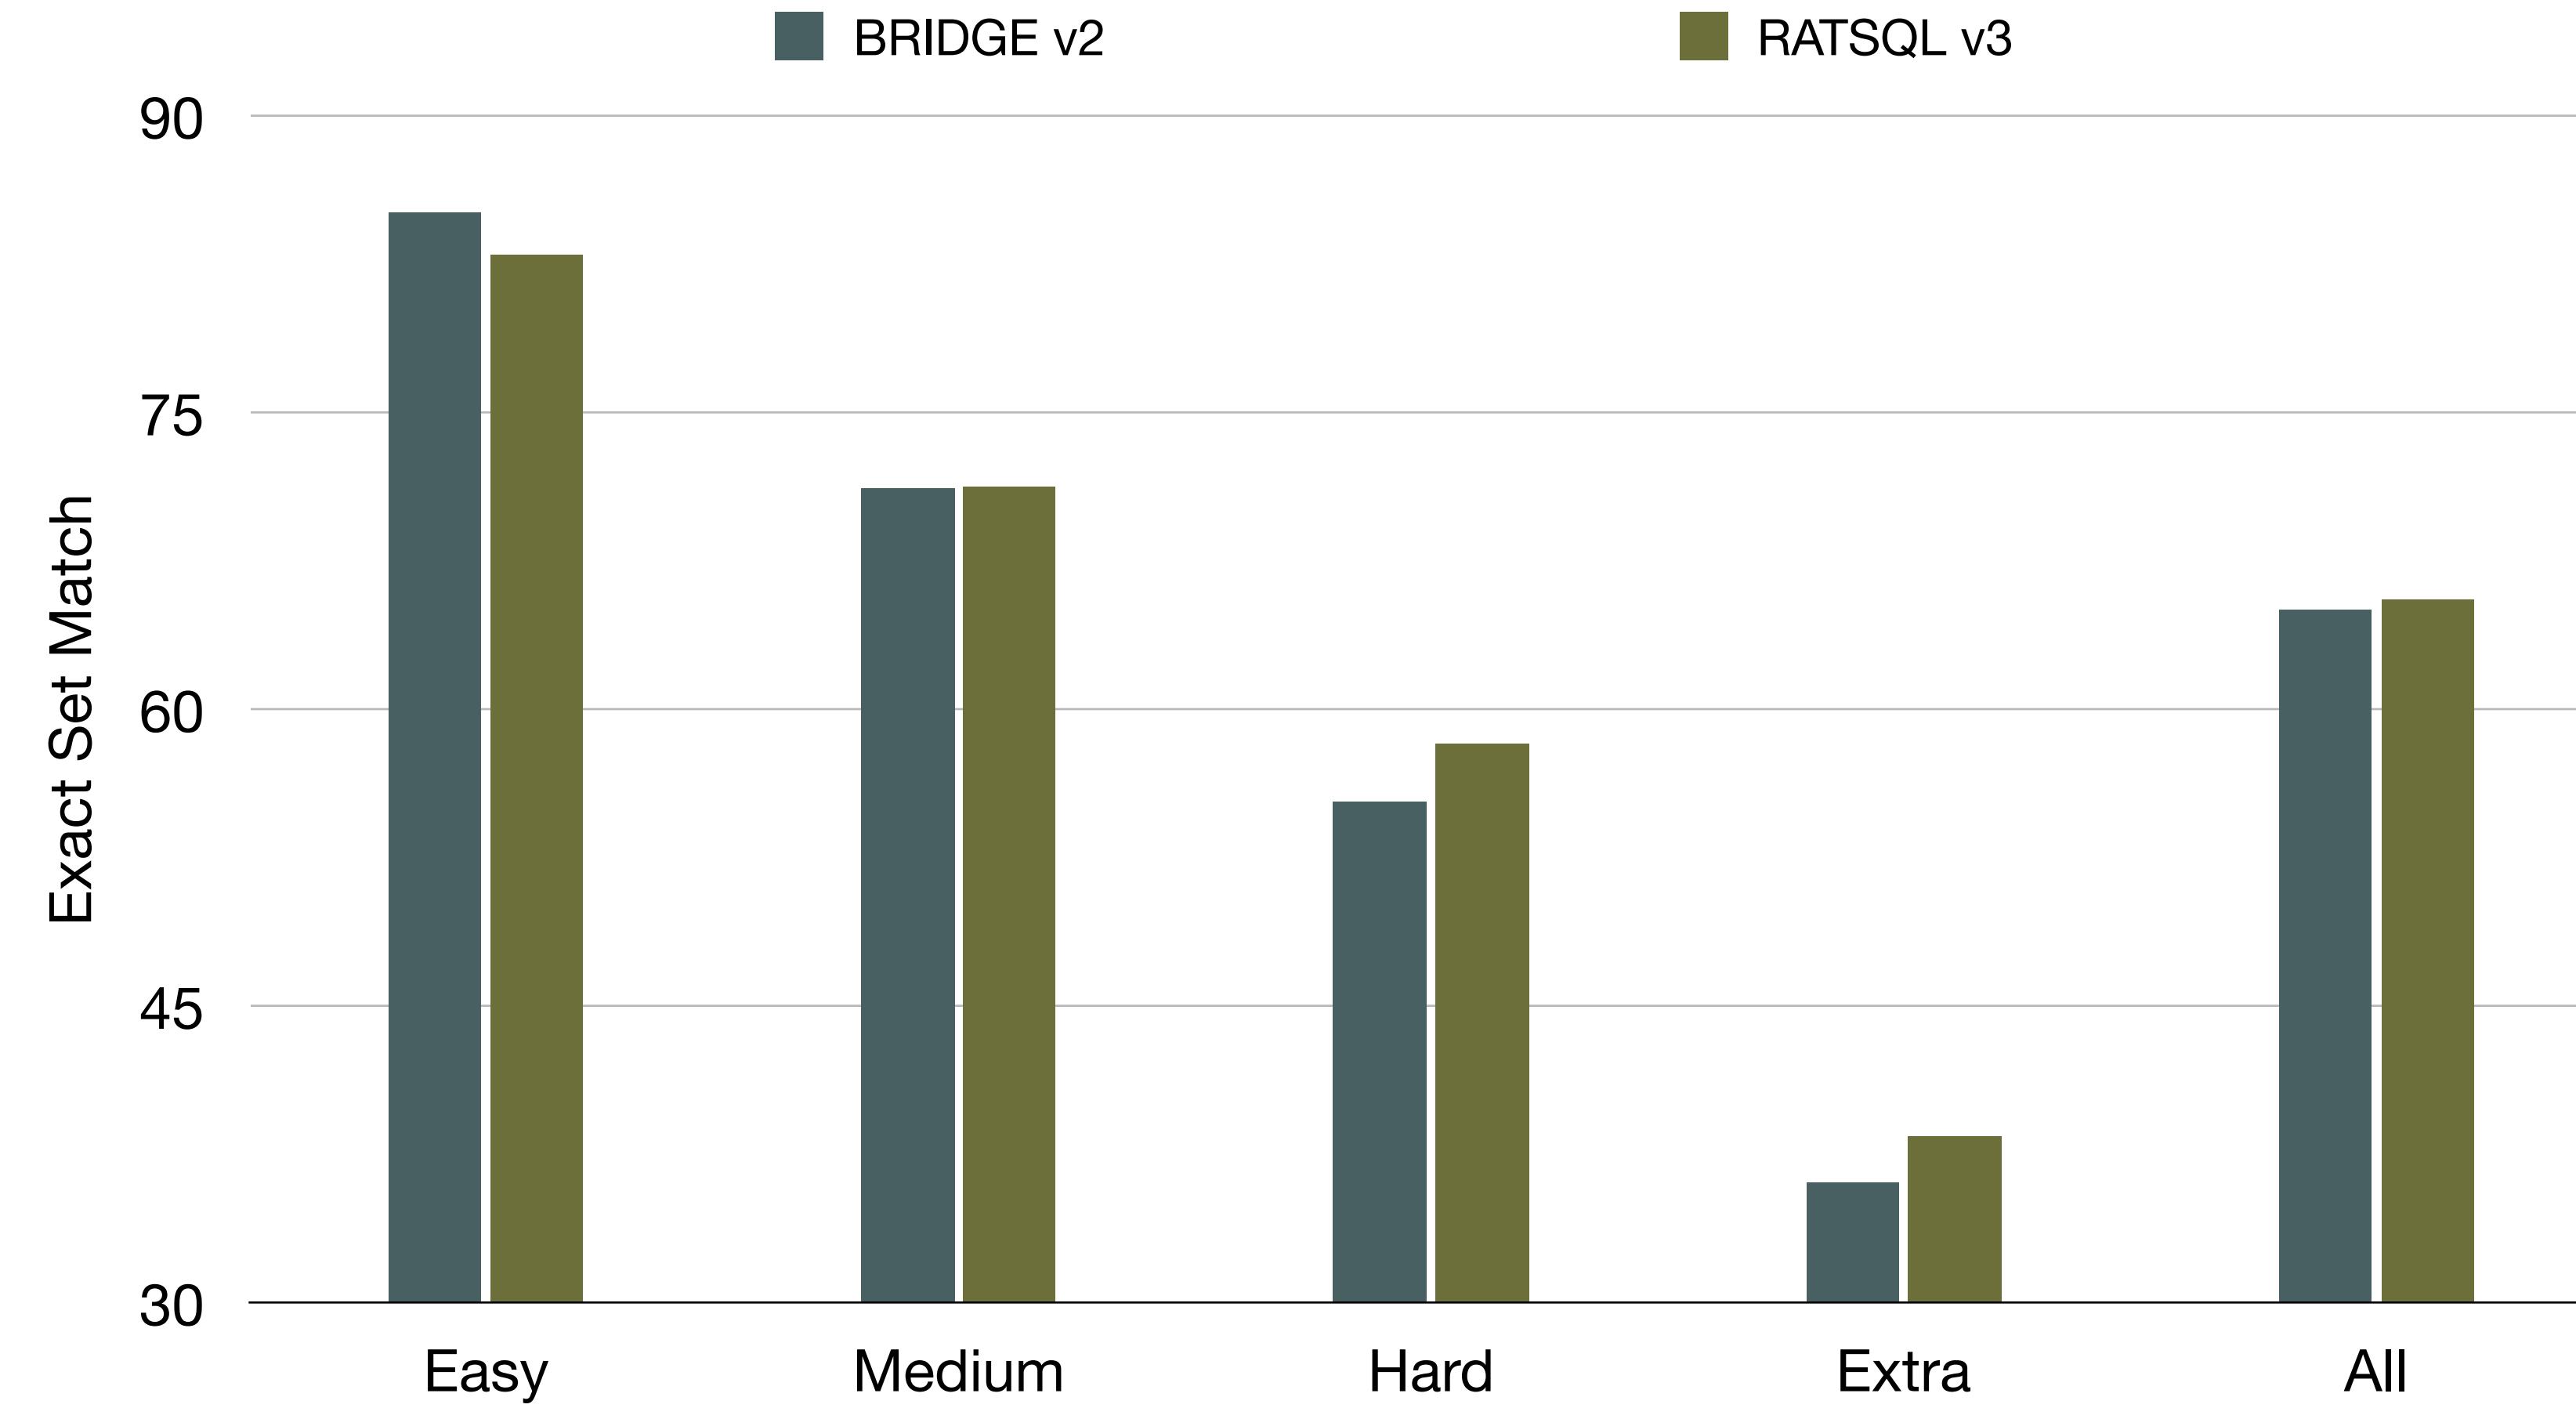
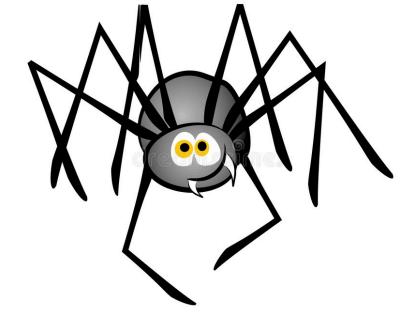


Figure 8. Performance of BRIDGE v2 compared to RATSQ v3 on the Spider Test set.

Performance Comparison by Difficulty Level

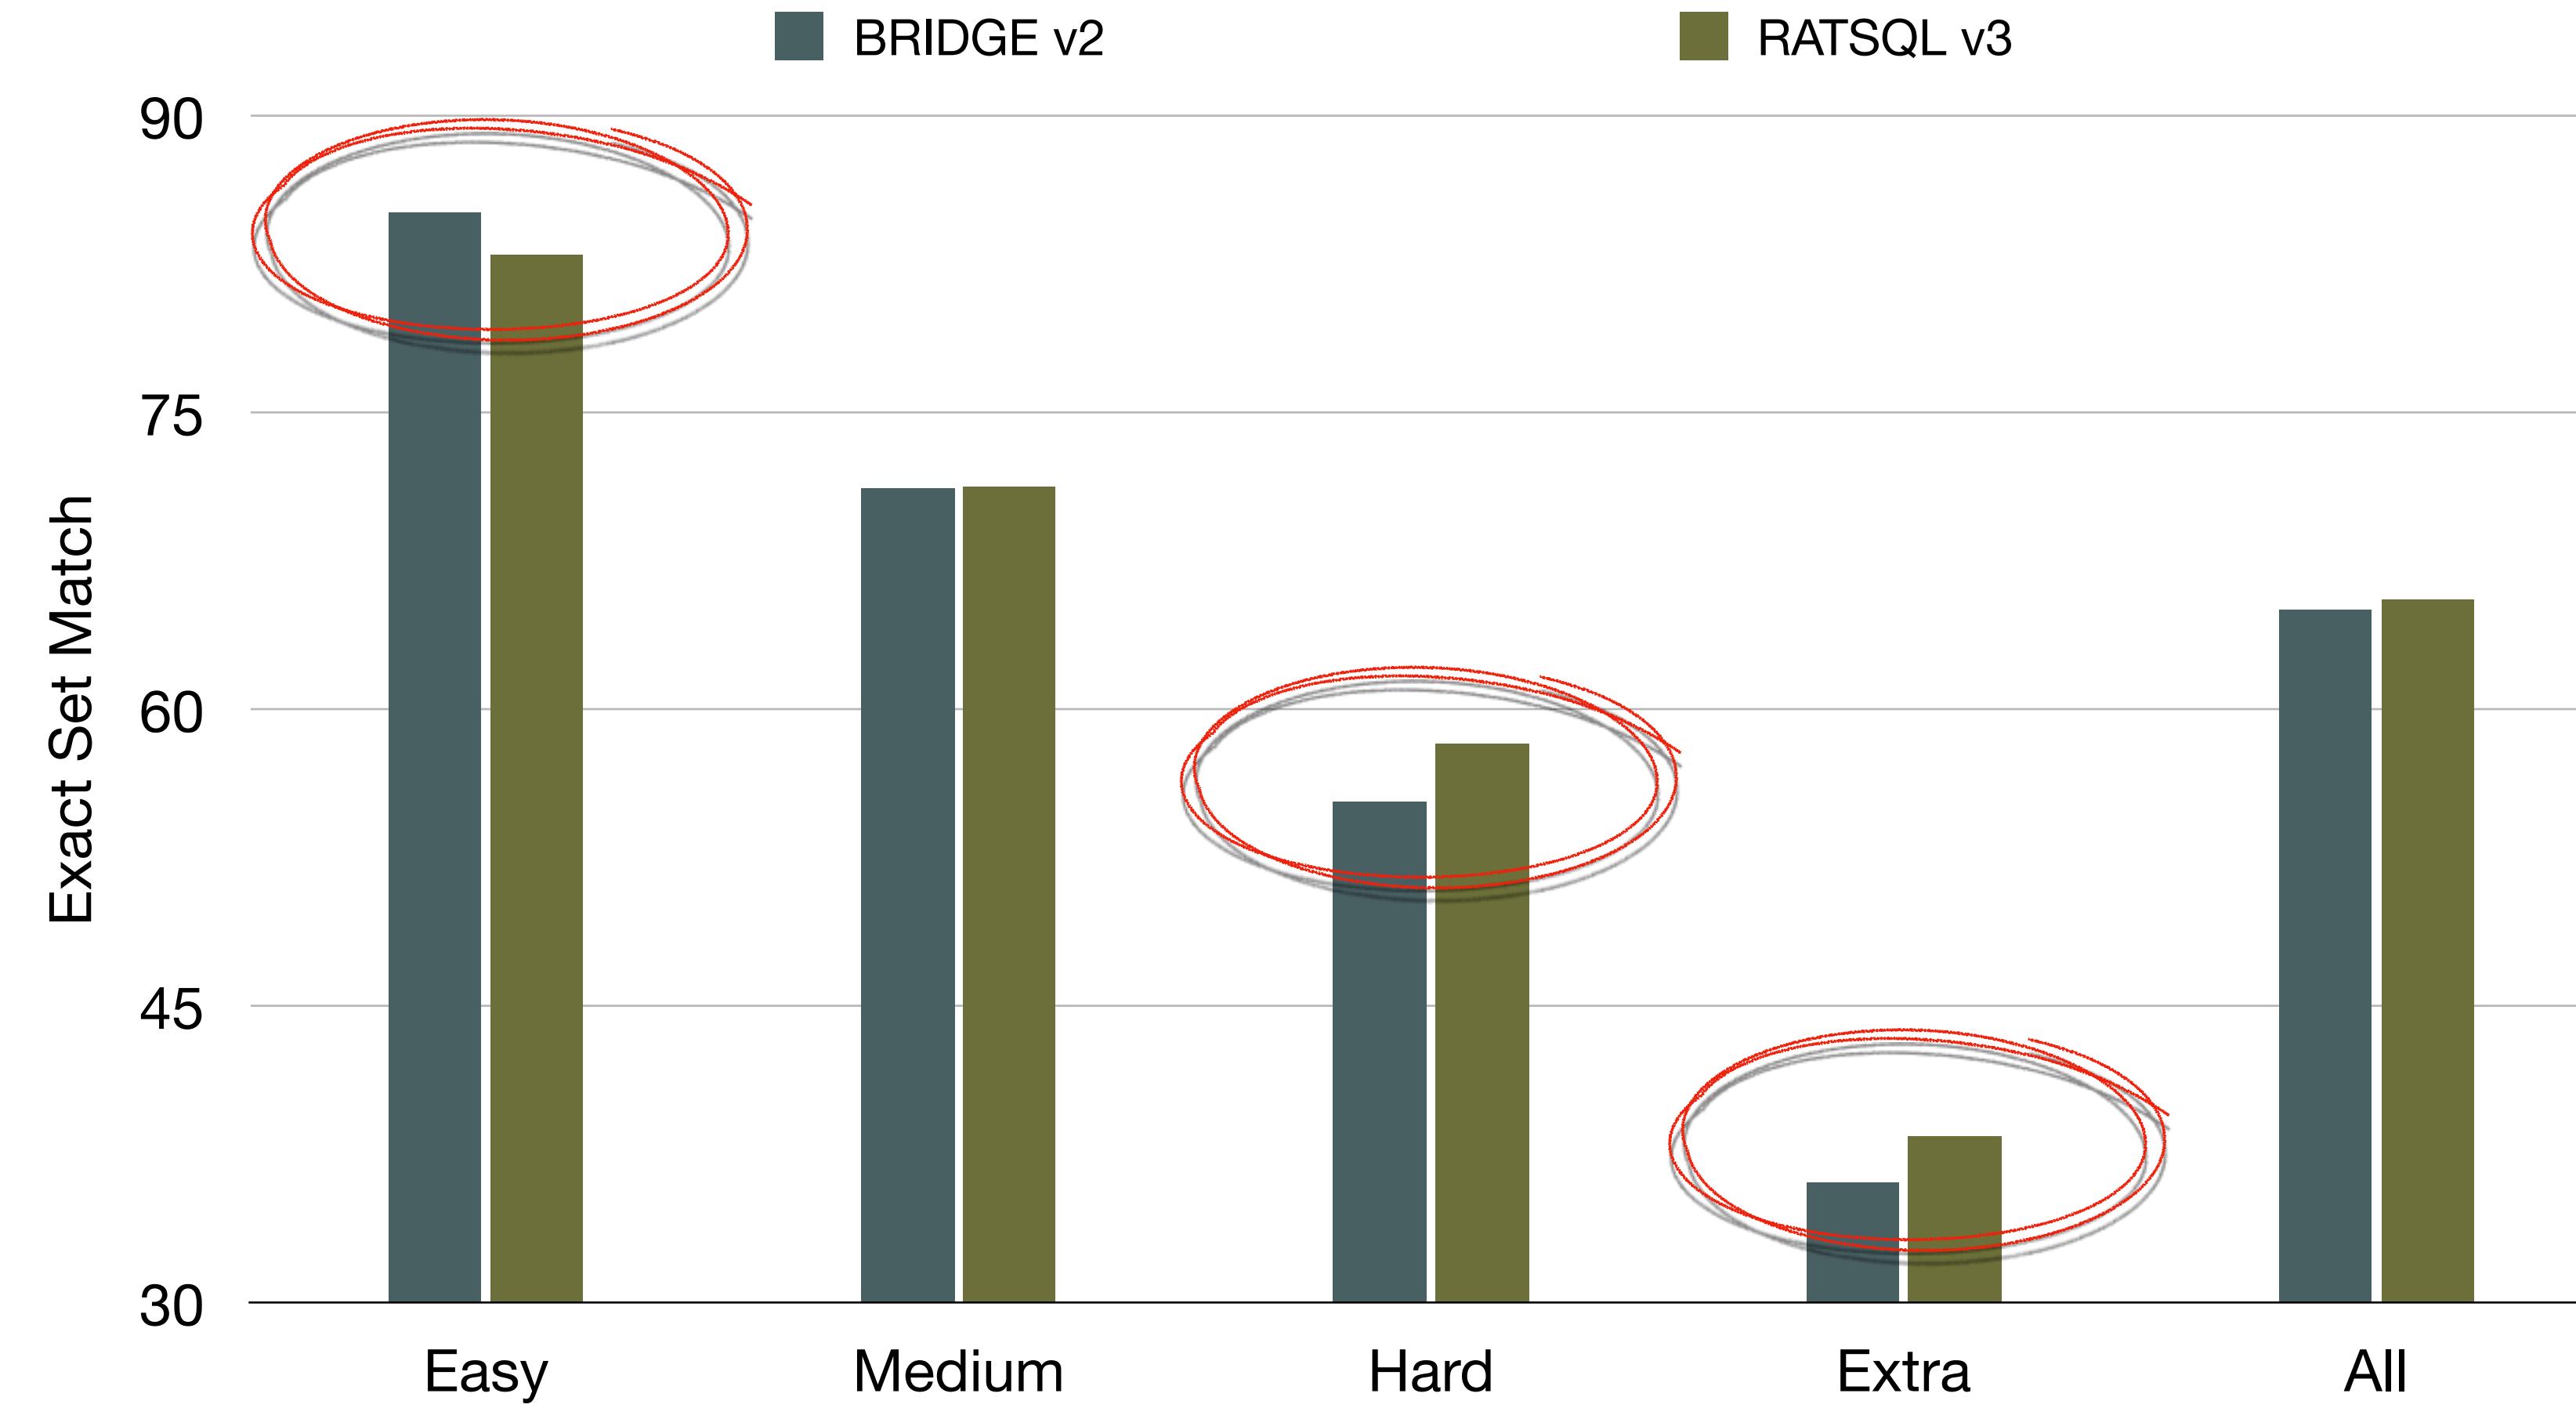
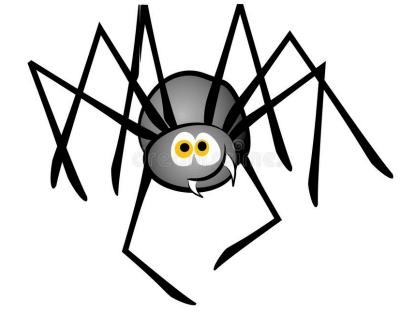
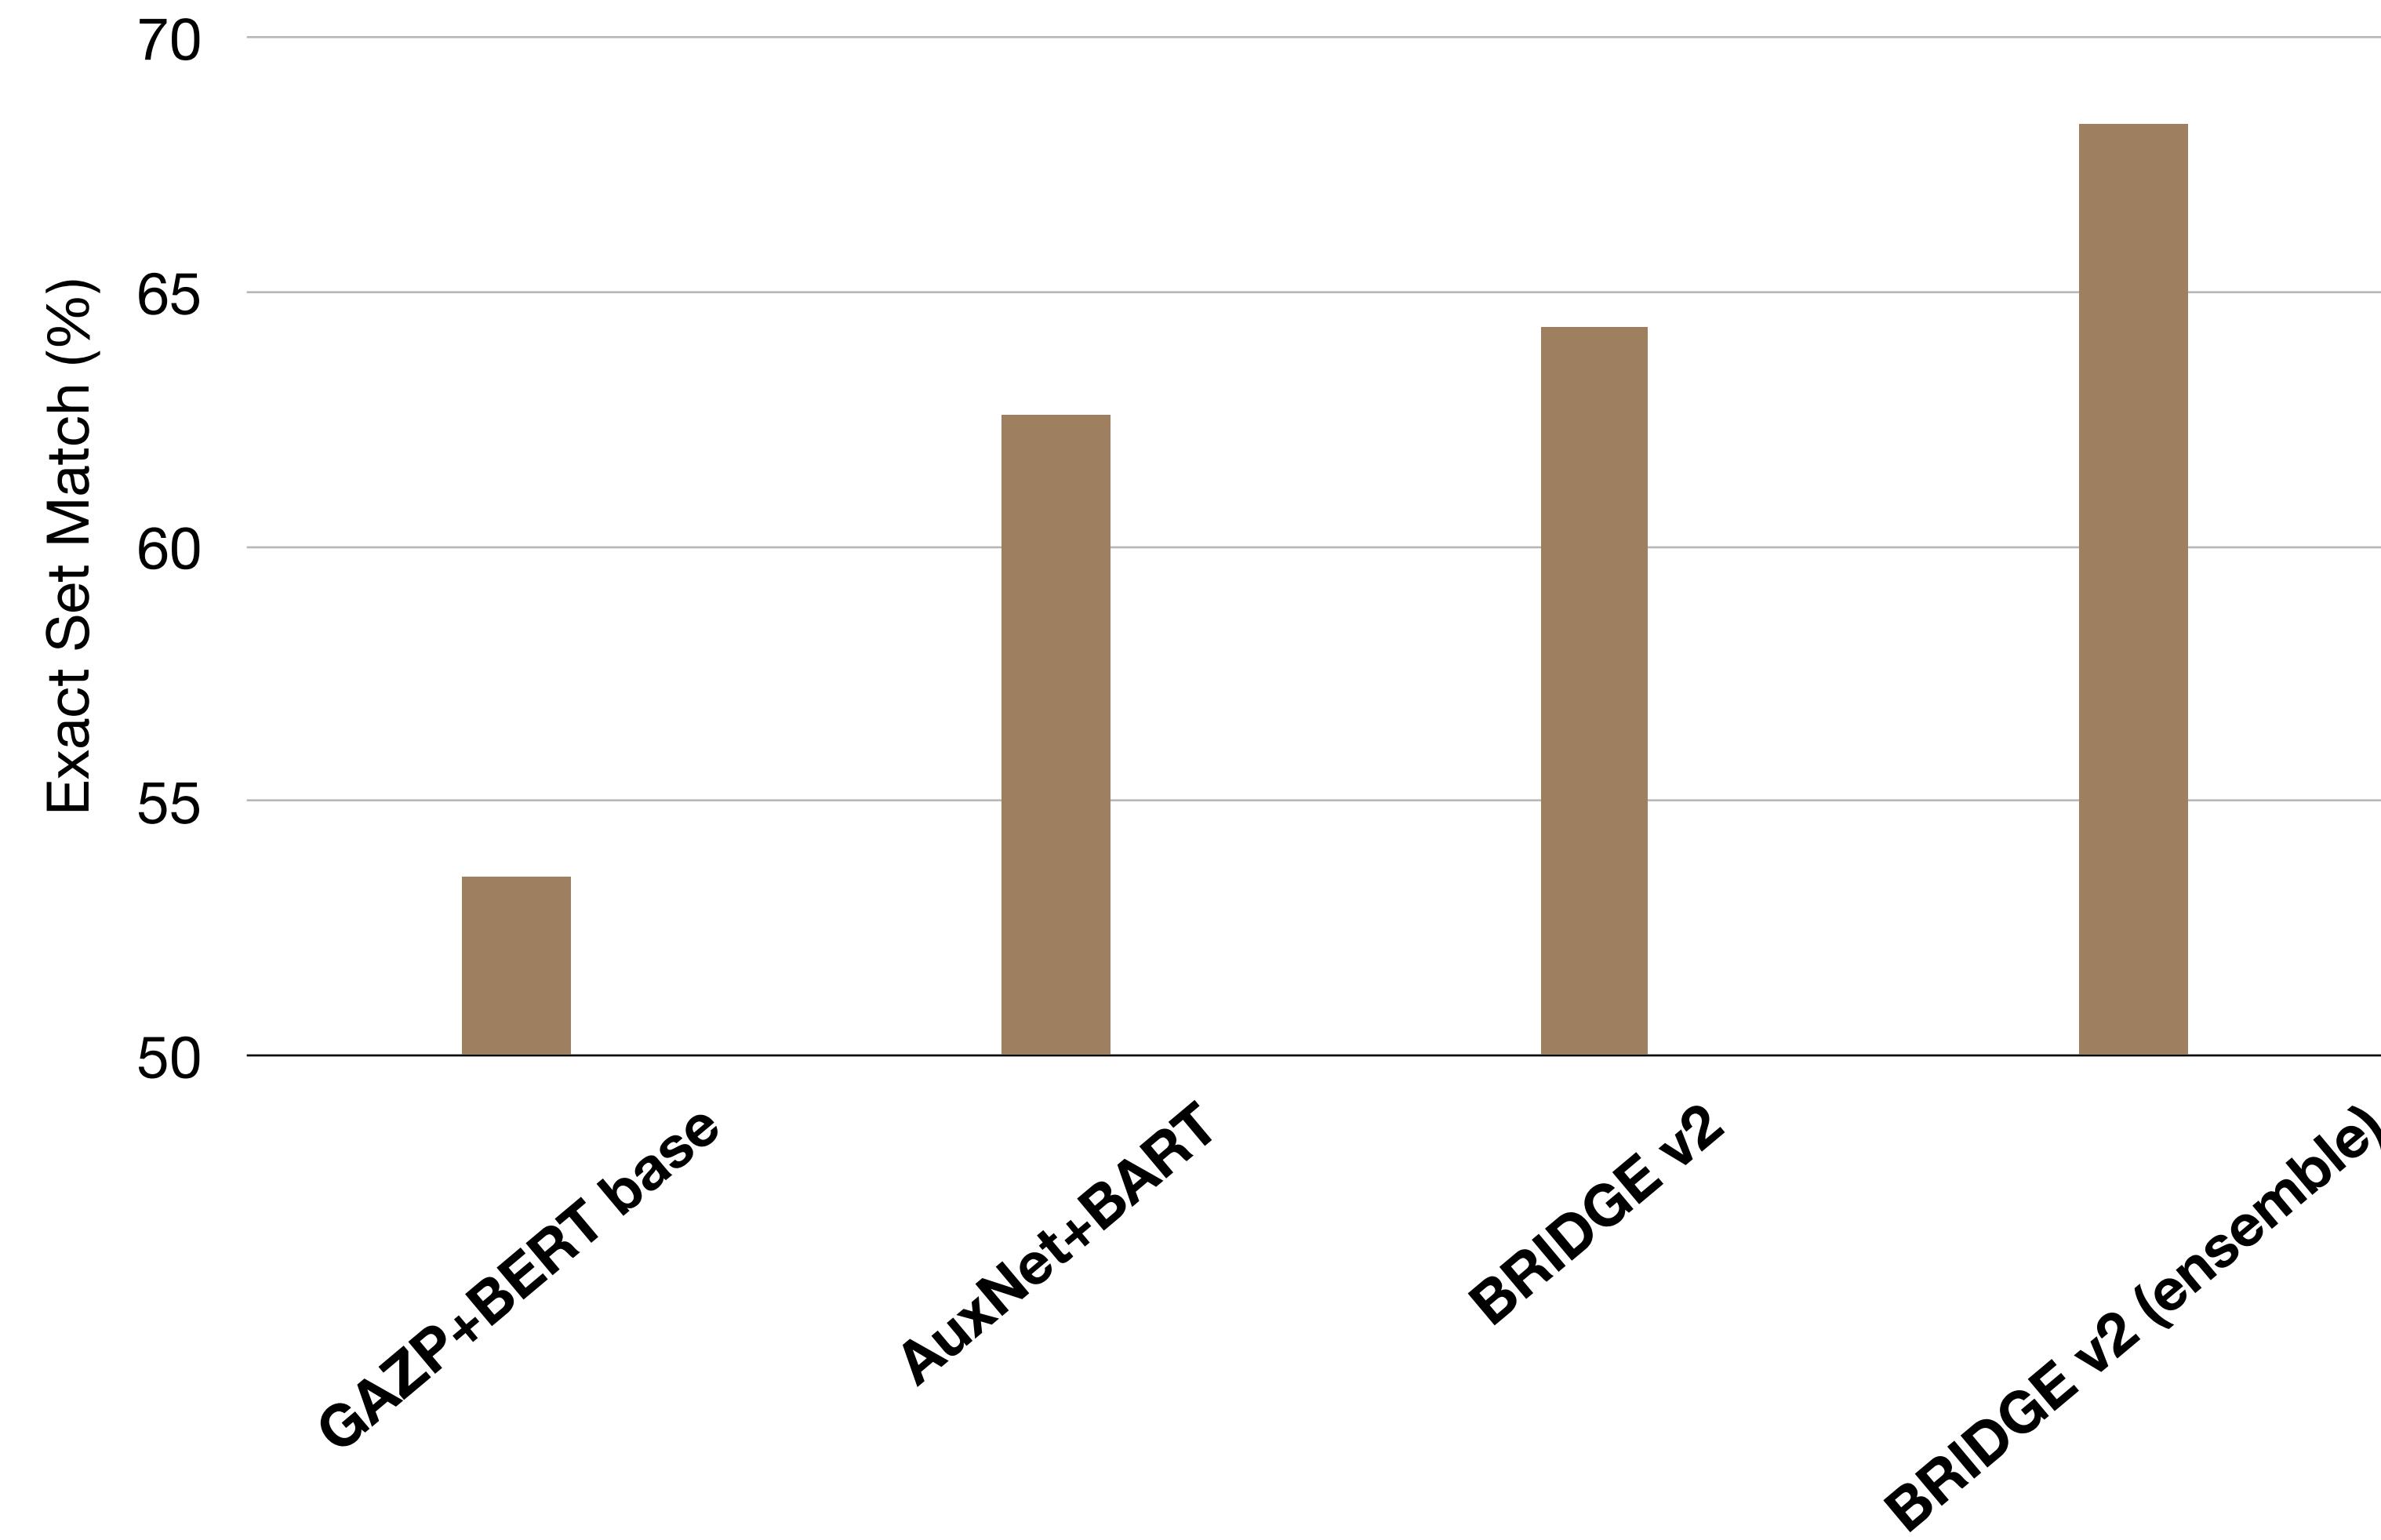


Figure 8. Performance of BRIDGE v2 compared to RATSQ v3 on the Spider Test set.

Performance on Spider Leaderboard - Execution Accuracy

Figure 7.1. Comparison to other top-performing text-to-SQL models on the Spider leaderboard based on execution accuracy (Jan 31, 2021).



Performance on WikiSQL Leaderboard



Figure 9. Comparison to other top text-to-SQL models on the WikiSQL leaderboard (Jan 31, 2020). ♠ denotes approaches that use table content during training. EG refers to “execution guided decoding”.

Model	Dev		Test	
	EM	EX	EM	EX
SQLova (Hwang et al., 2019)	81.6	87.2	80.7	86.2
X-SQL (He et al., 2019b)	83.8	89.5	83.3	88.7
IE-SQL (Ma et al., 2020)	84.6	88.7	84.6	88.8
NL2SQL ♠ (Guo and Gao, 2019)	84.3	90.3	83.7	89.2
HydraNet (Lyu et al., 2020)	83.6	89.1	83.8	89.2
BRIDGE _L ♠	86.2	91.7	85.7	91.1
SQLova+EG (Hwang et al., 2019)	84.2	90.2	83.6	89.6
NL2SQL+EG ♠ (Guo and Gao, 2019)	85.4	91.1	84.5	90.1
X-SQL+EG (He et al., 2019b)	86.2	92.3	86.0	91.8
BRIDGE _L +EG ♠	86.8	92.6	86.3	91.9
HydraNet+EG (Lyu et al., 2020)	86.6	92.4	86.5	92.2
IE-SQL+EG (Ma et al., 2020)	87.9	92.6	87.8	92.5

Performance on WikiSQL Leaderboard

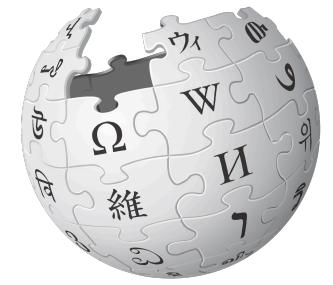


Figure 9. Comparison to other top text-to-SQL models on the WikiSQL leaderboard (Jan 31, 2020). ♠ denotes approaches that use table content during training. EG refers to “execution guided decoding”.

Model	Dev		Test	
	EM	EX	EM	EX
SQLova (Hwang et al., 2019)	81.6	87.2	80.7	86.2
X-SQL (He et al., 2019b)	83.8	89.5	83.3	88.7
IE-SQL (Ma et al., 2020)	84.6	88.7	84.6	88.8
NL2SQL ♠ (Guo and Gao, 2019)	84.3	90.3	83.7	89.2
HydraNet (Lyu et al., 2020)	83.6	89.1	83.8	89.2
BRIDGE _L ♠	86.2	91.7	85.7	91.1
SQLova+EG (Hwang et al., 2019)	84.2	90.2	83.6	89.6
NL2SQL+EG ♠ (Guo and Gao, 2019)	85.4	91.1	84.5	90.1
X-SQL+EG (He et al., 2019b)	86.2	92.3	86.0	91.8
BRIDGE _L +EG ♠	86.8	92.6	86.3	91.9
HydraNet+EG (Lyu et al., 2020)	86.6	92.4	86.5	92.2
IE-SQL+EG (Ma et al., 2020)	87.9	92.6	87.8	92.5

Best model without
execution guided
decoding

Performance on WikiSQL Leaderboard



Figure 9. Comparison to other top text-to-SQL models on the WikiSQL leaderboard (Jan 31, 2020). ♠ denotes approaches that use table content during training. EG refers to “execution guided decoding”.

Model	Dev		Test	
	EM	EX	EM	EX
SQLova (Hwang et al., 2019)	81.6	87.2	80.7	86.2
X-SQL (He et al., 2019b)	83.8	89.5	83.3	88.7
IE-SQL (Ma et al., 2020)	84.6	88.7	84.6	88.8
NL2SQL ♠ (Guo and Gao, 2019)	84.3	90.3	83.7	89.2
HydraNet (Lyu et al., 2020)	83.6	89.1	83.8	89.2
BRIDGE_L ♠	86.2	91.7	85.7	91.1
SQLova+EG (Hwang et al., 2019)	84.2	90.2	83.6	89.6
NL2SQL+EG ♠ (Guo and Gao, 2019)	85.4	91.1	84.5	90.1
X-SQL+EG (He et al., 2019b)	86.2	92.3	86.0	91.8
BRIDGE _L +EG ♠	86.8	92.6	86.3	91.9
HydraNet+EG (Lyu et al., 2020)	86.6	92.4	86.5	92.2
IE-SQL+EG (Ma et al., 2020)	87.9	92.6	87.8	92.5

Model	w/o EG		w/ EG	
	EM	EX	EM	EX
BRIDGE _L	86.2	91.7	86.8	92.6
-anchor text	84.2	90.0	85.2	91.3
-bridging	82.6	88.5	84.5	90.8

Best model without execution guided decoding

Performance on WikiSQL Leaderboard



Figure 9. Comparison to other top text-to-SQL models on the WikiSQL leaderboard (Jan 31, 2020). ♠ denotes approaches that use table content during training. EG refers to “execution guided decoding”.

Model	Dev		Test	
	EM	EX	EM	EX
SQLova (Hwang et al., 2019)	81.6	87.2	80.7	86.2
X-SQL (He et al., 2019b)	83.8	89.5	83.3	88.7
IE-SQL (Ma et al., 2020)	84.6	88.7	84.6	88.8
NL2SQL ♠ (Guo and Gao, 2019)	84.3	90.3	83.7	89.2
HydraNet (Lyu et al., 2020)	83.6	89.1	83.8	89.2
BRIDGE_L ♠	86.2	91.7	85.7	91.1
SQLova+EG (Hwang et al., 2019)	84.2	90.2	83.6	89.6
NL2SQL+EG ♠ (Guo and Gao, 2019)	85.4	91.1	84.5	90.1
X-SQL+EG (He et al., 2019b)	86.2	92.3	86.0	91.8
BRIDGE _L +EG ♠	86.8	92.6	86.3	91.9
HydraNet+EG (Lyu et al., 2020)	86.6	92.4	86.5	92.2
IE-SQL+EG (Ma et al., 2020)	87.9	92.6	87.8	92.5

Top-3 model using
execution guided
decoding

Cross-Database Performance

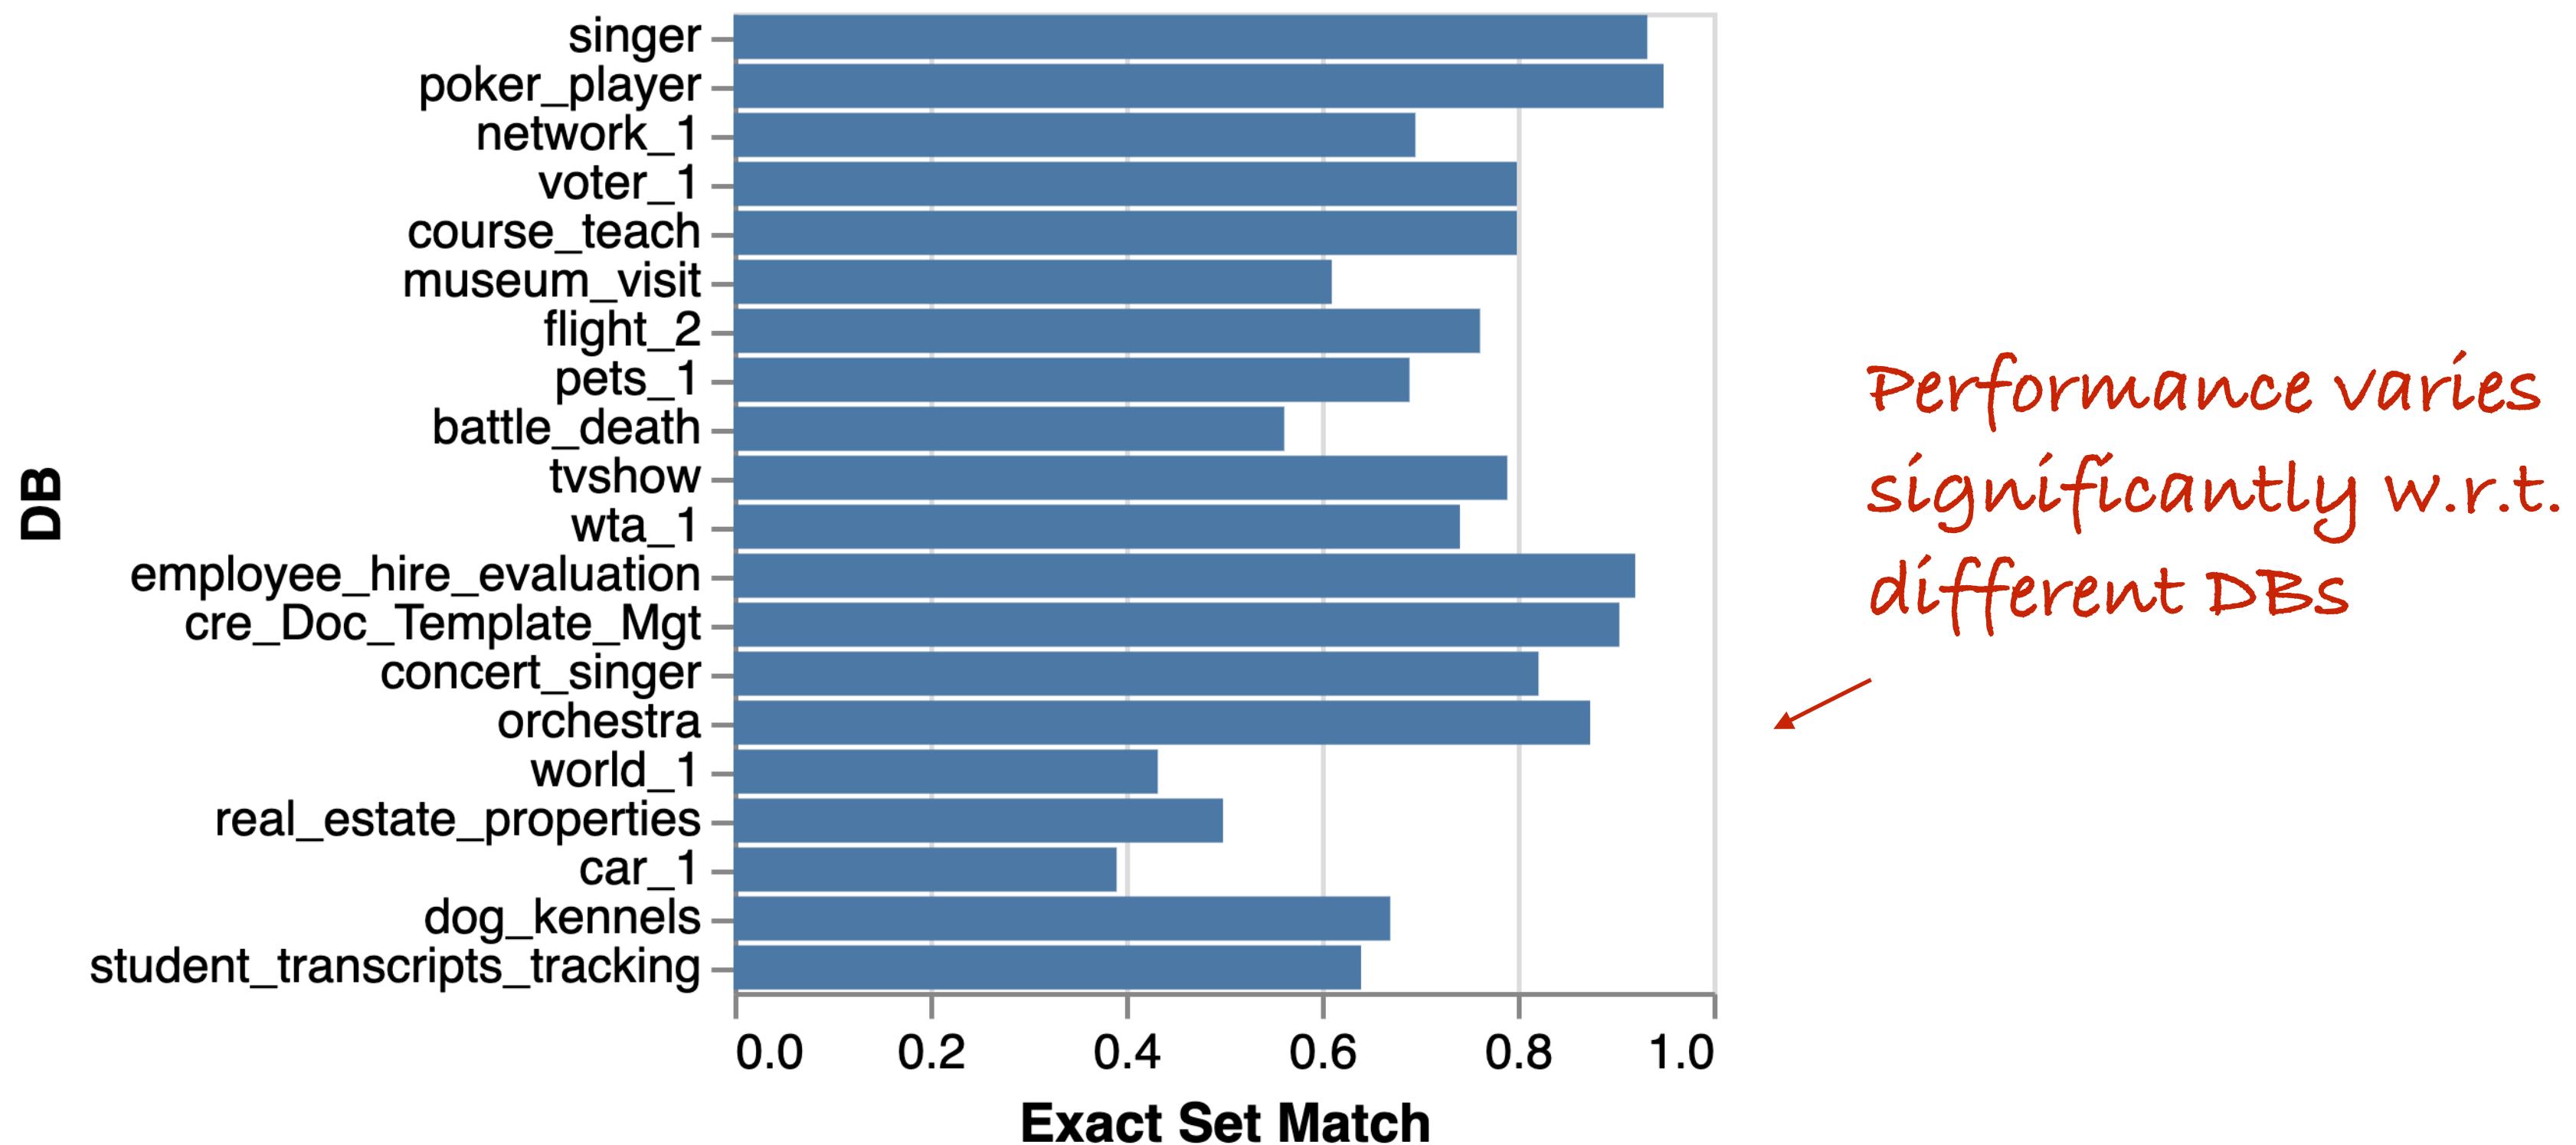
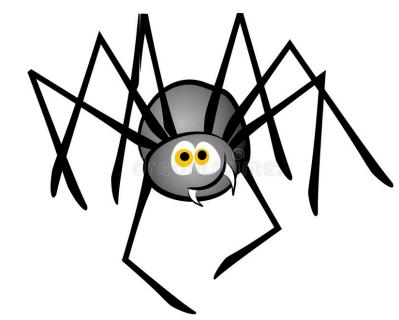


Figure 10. Performance of BRIDGE on each database on the Spider dev set.

Error Analysis

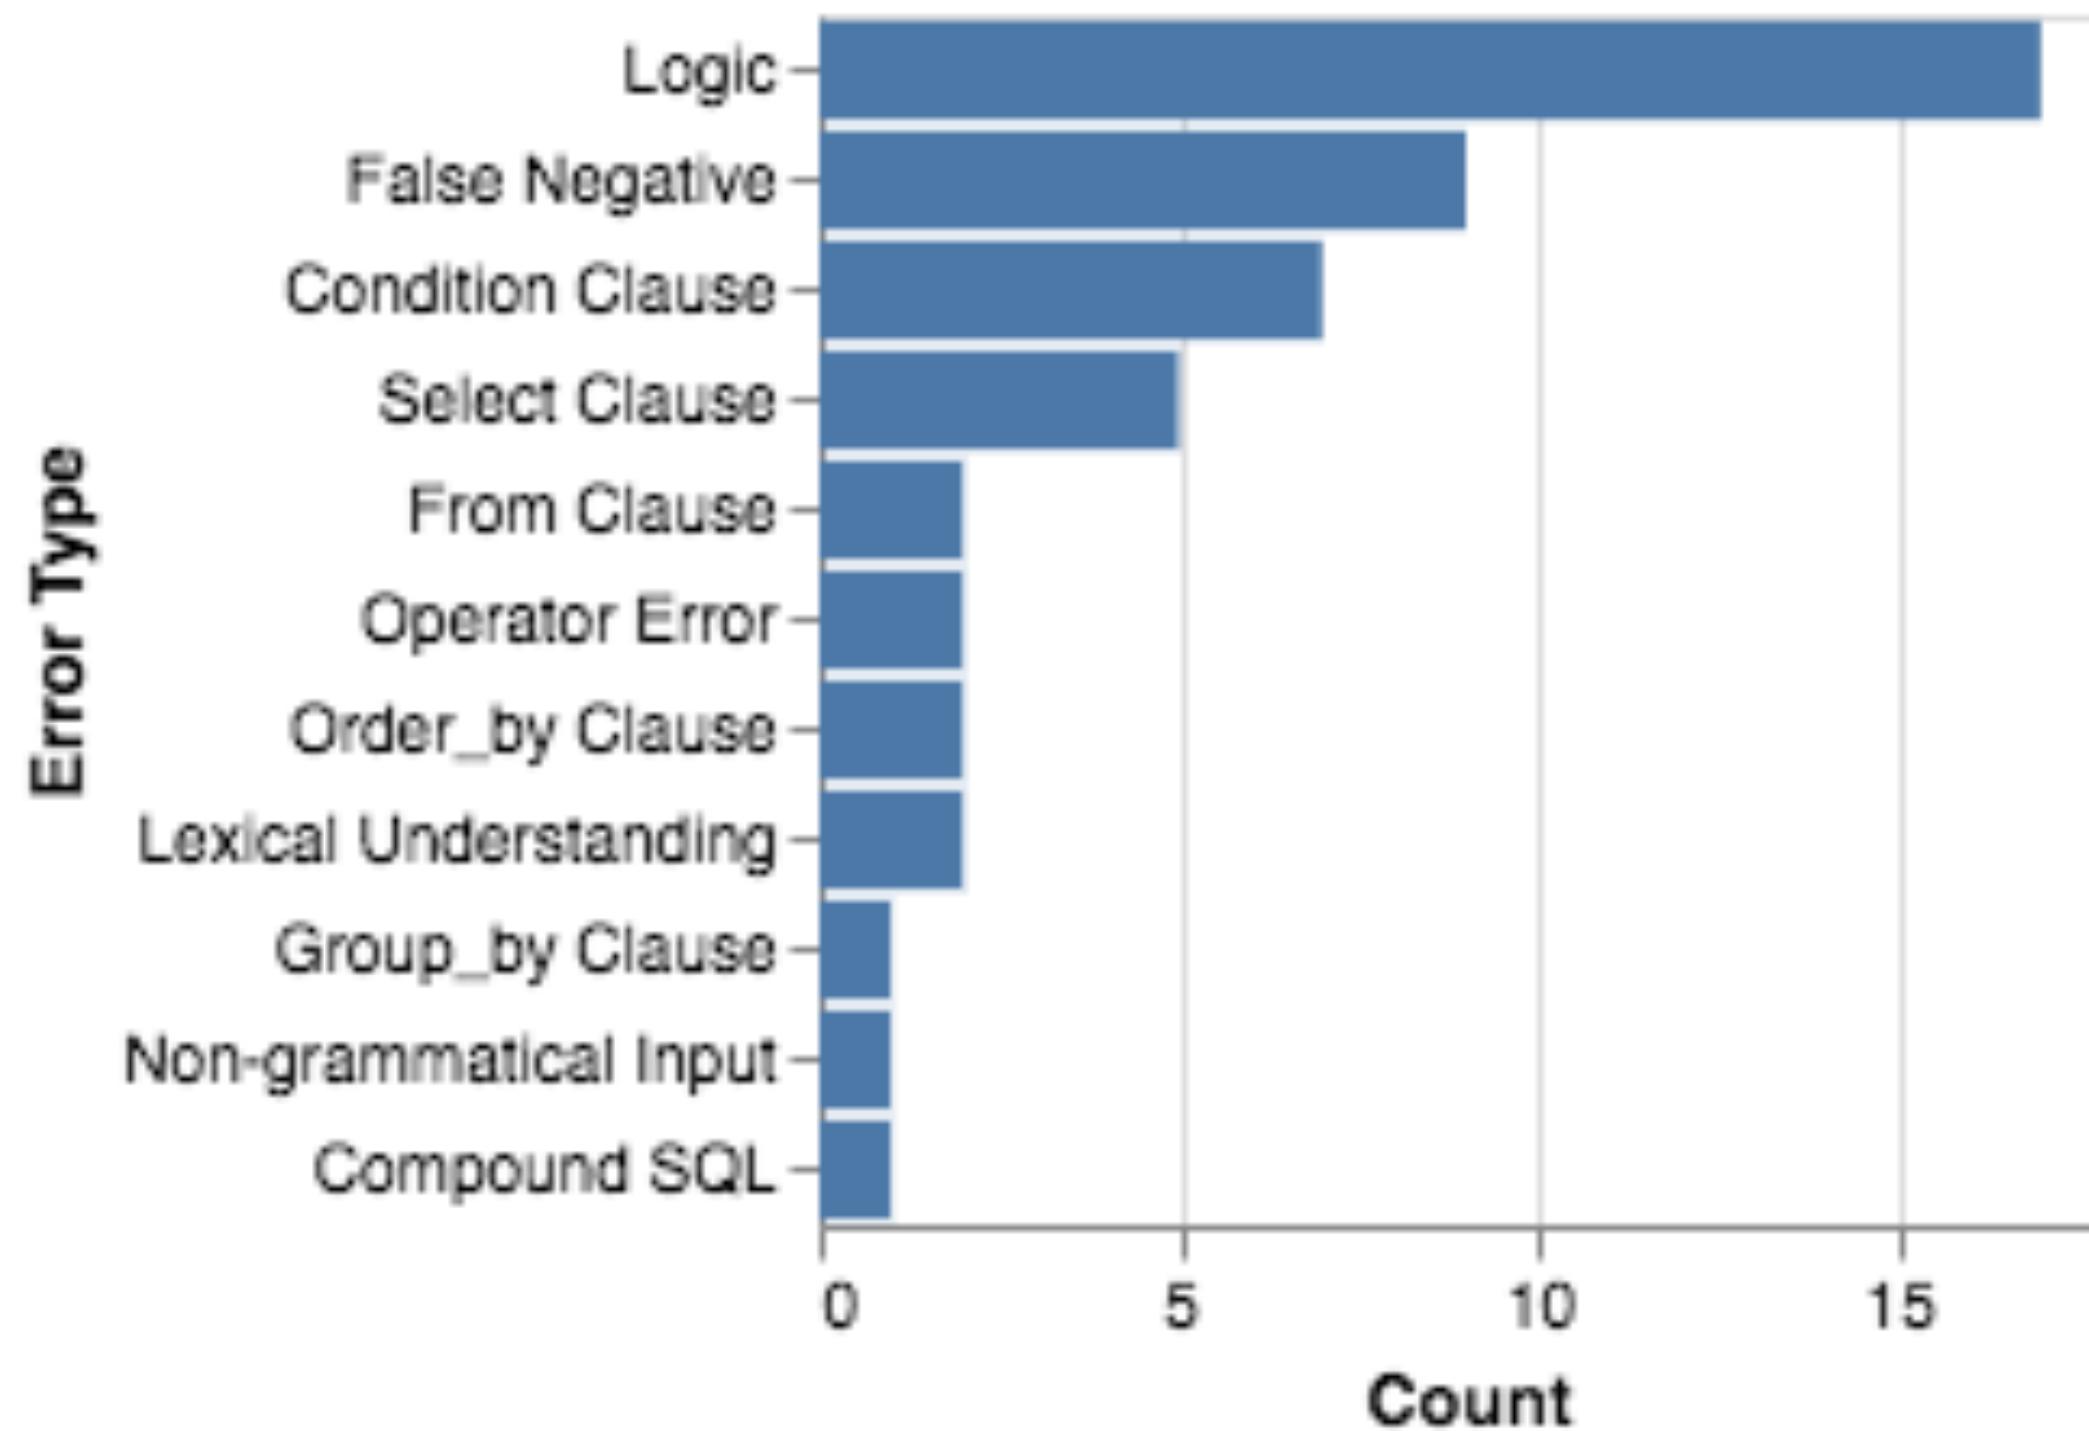
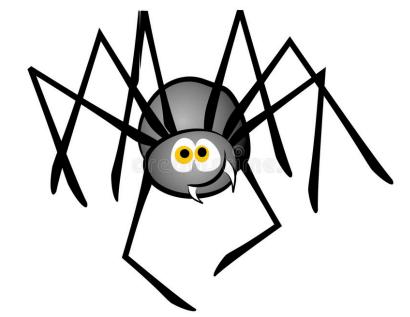


Figure 11. Manual error categorization for 50 wrong predictions on the Spider dev set.

Error Analysis

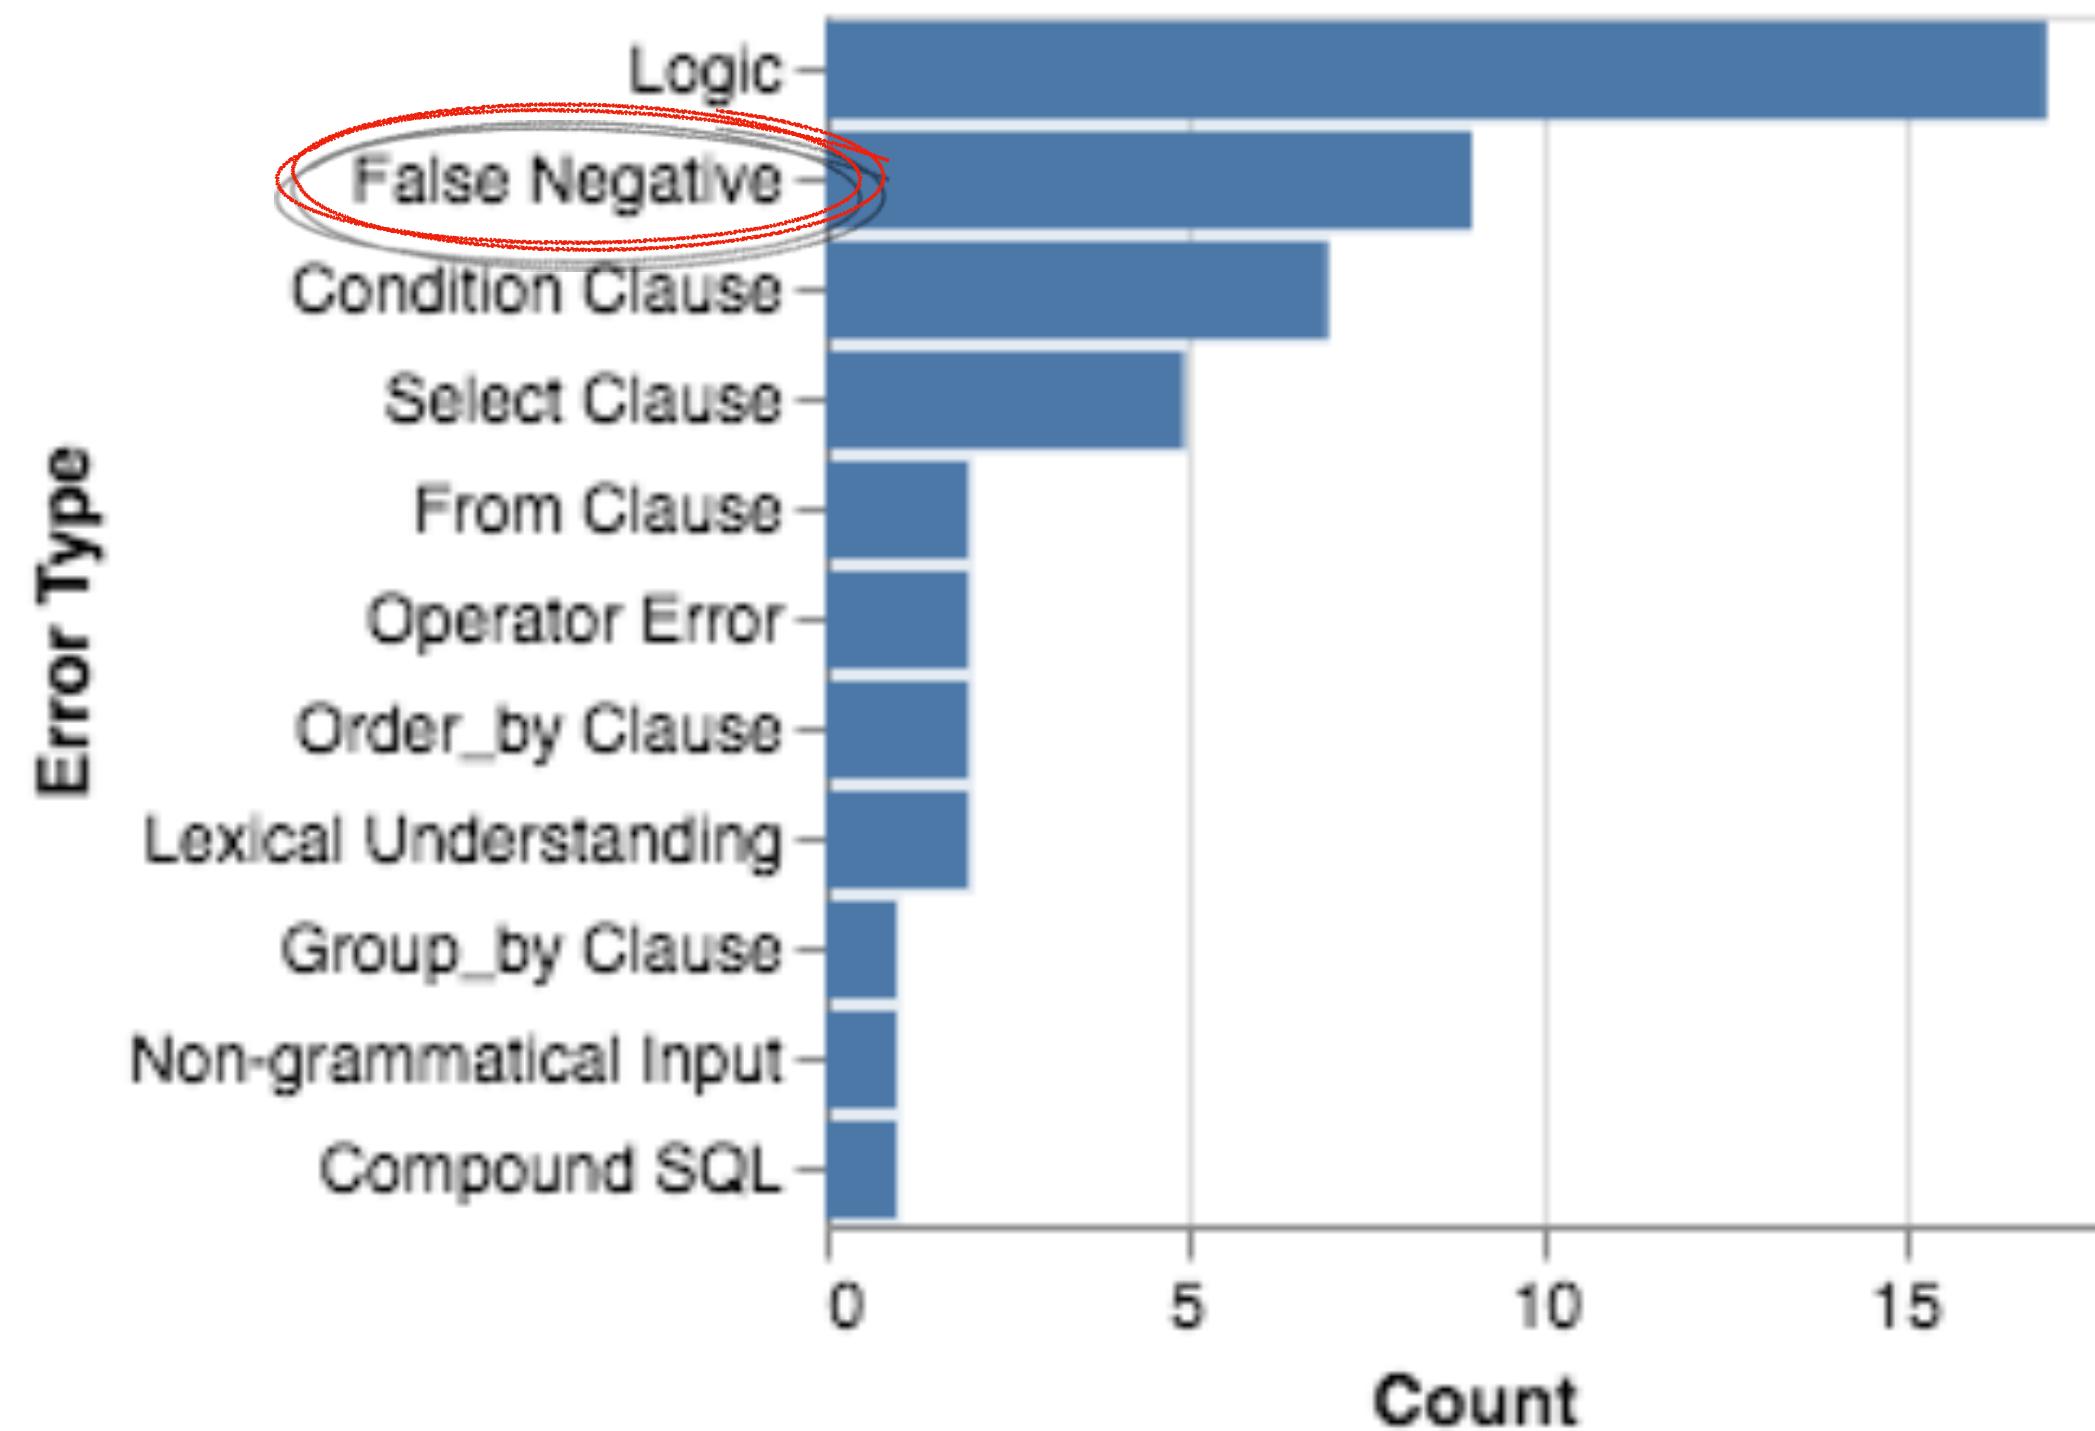
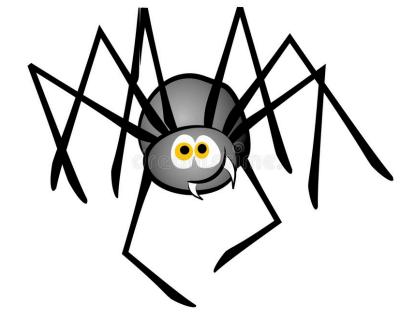


Figure 11. Manual error categorization for 50 wrong predictions on the Spider dev set.

Error Analysis

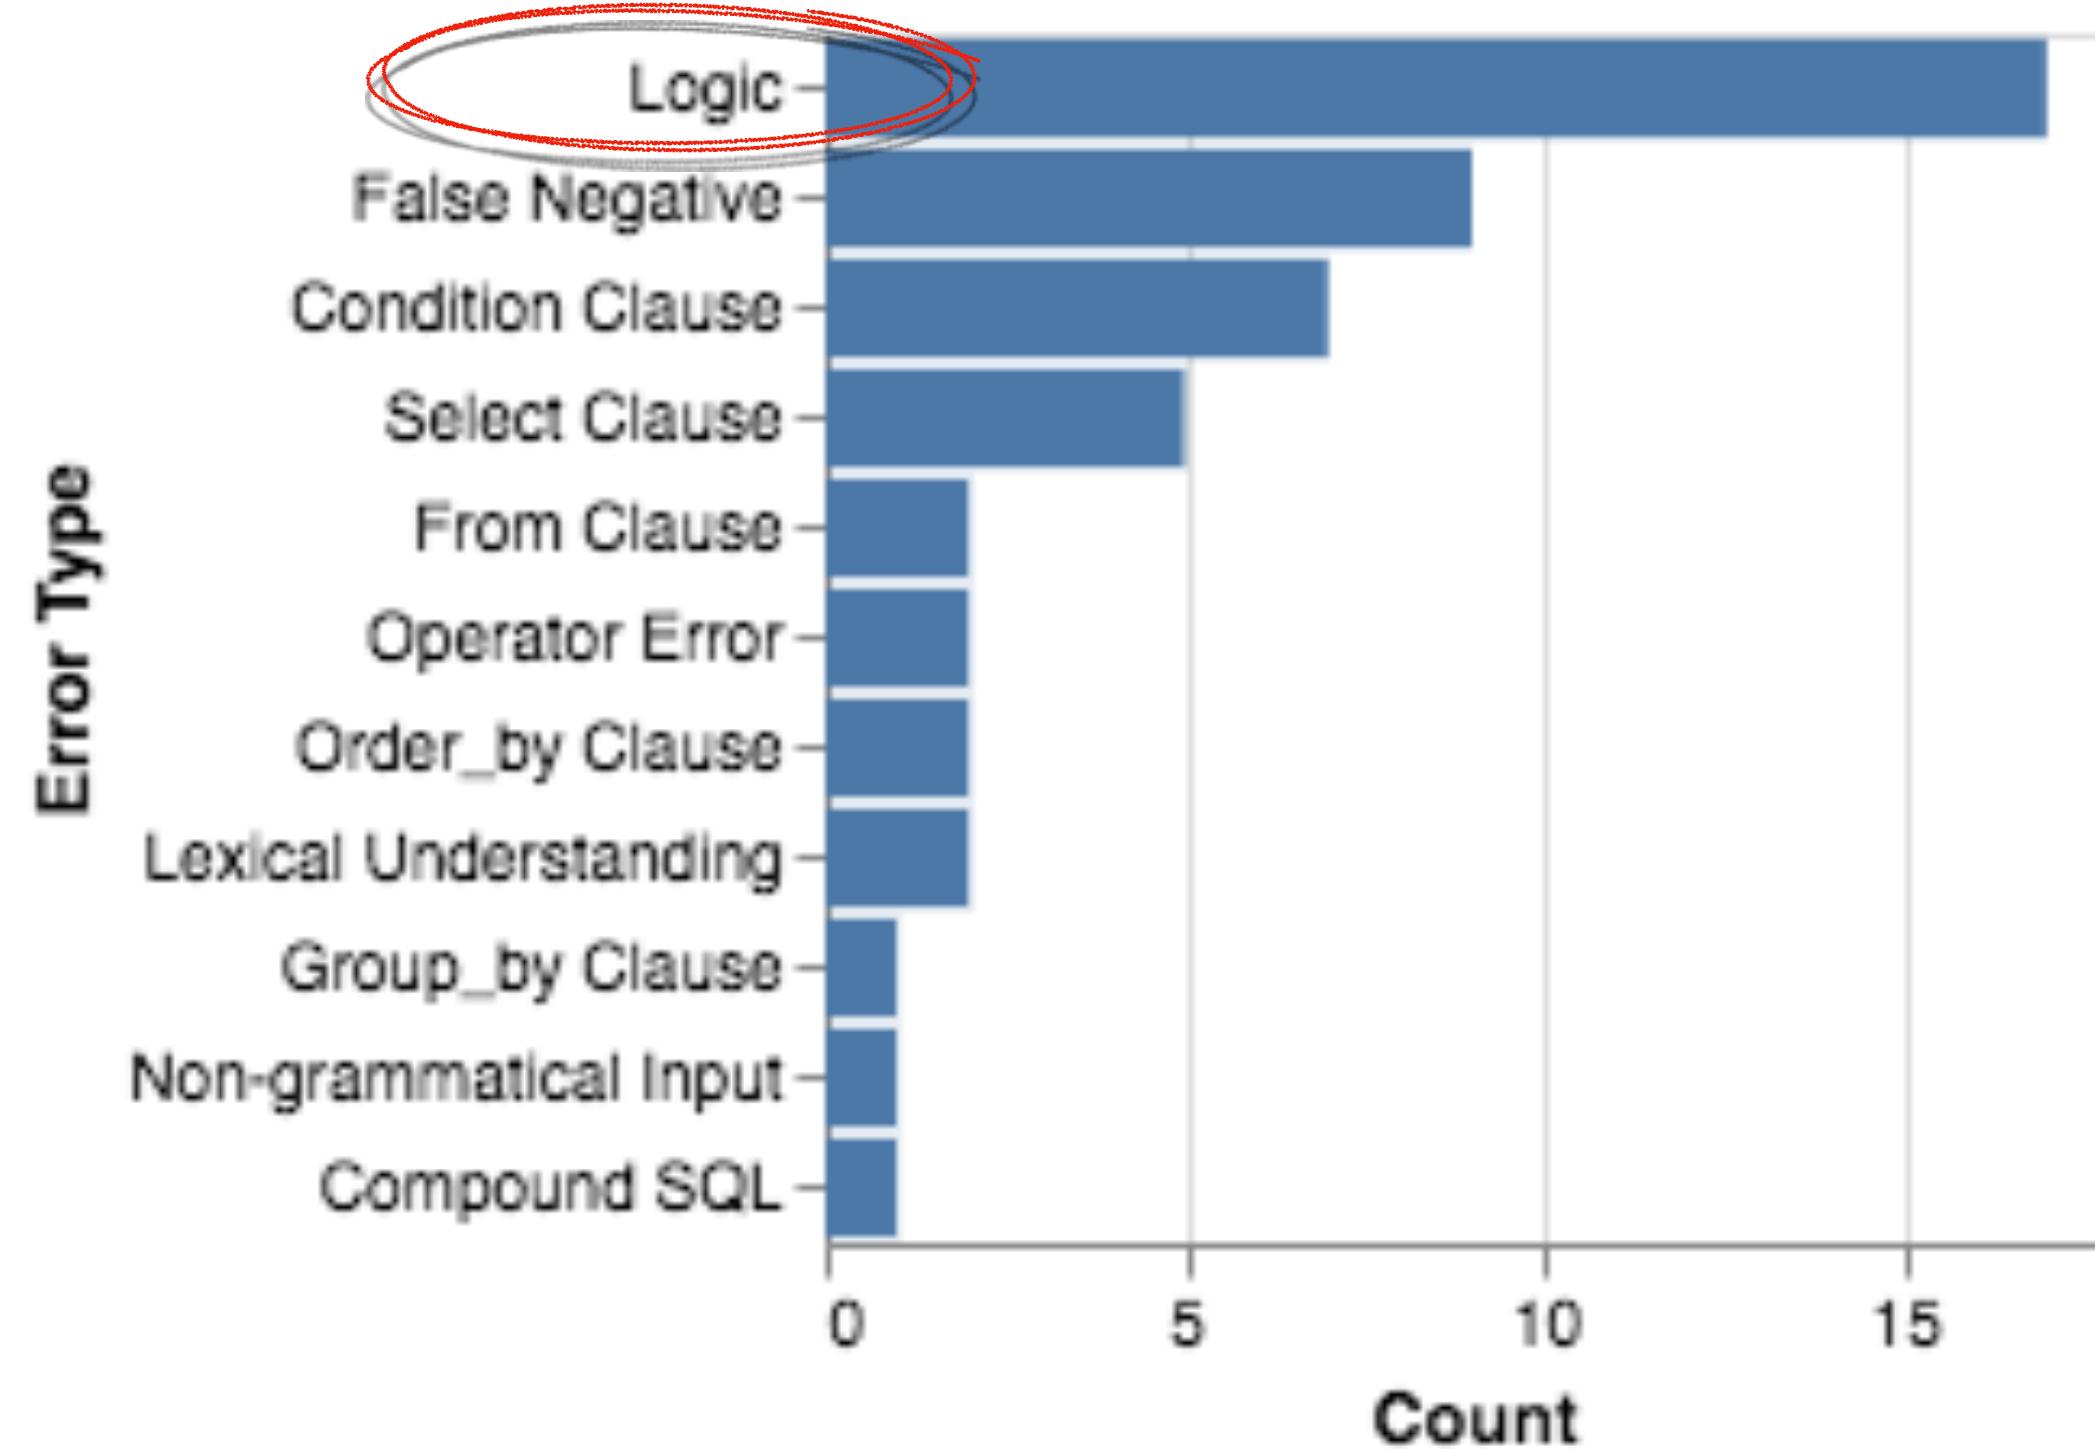
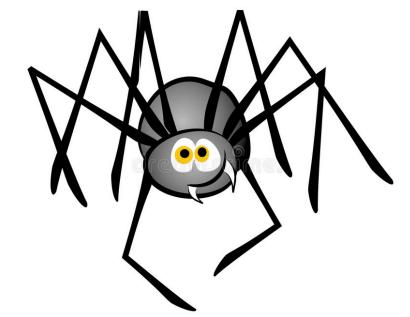
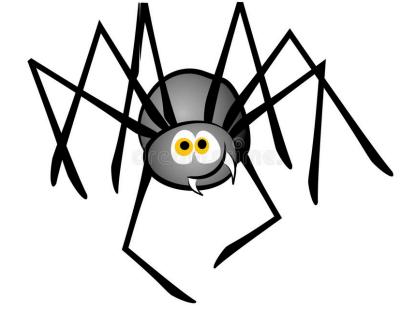


Figure 11. Manual error categorization for 50 wrong predictions on the Spider dev set.

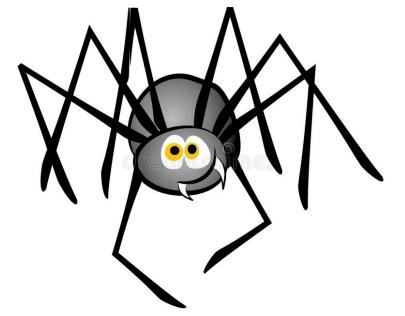
Qualitative Examples



Error Category I - Logic

-
- Logic**
- 👤 *Find the number of concerts happened in the stadium with the highest capacity.*
concert_singer
 - ✗ SELECT COUNT(*) FROM stadium JOIN concert ON stadium.Stadium_ID = concert.Stadium_ID ORDER BY stadium.Capacity DESC LIMIT 1
 - ✓ SELECT COUNT(*) FROM concert WHERE stadium_id = (SELECT stadium_id FROM stadium ORDER BY capacity DESC LIMIT 1)
 - 👤 *Show the names of all of the high schooler Kyle's friends.* network_1
 - ✗ SELECT Highschooler.name FROM Friend JOIN Highschooler ON Friend.friend_id = Highschooler.ID WHERE Highschooler.name = "Kyle"
 - ✓ SELECT T3.name FROM Friend AS T1 JOIN Highschooler AS T2 ON T1.student_id = T2.id JOIN Highschooler AS T3 ON T1.friend_id = T3.id WHERE T2.name = "Kyle"
-

Qualitative Examples



Error Category II - Lexical Understanding

Lexical Understanding



Count the number of countries for which Spanish is the predominantly spoken language.

world_1



`SELECT COUNT(*) FROM countrylanguage WHERE countrylanguage.Language = "Spanish"`



`SELECT COUNT(*), MAX(Percentage) FROM countrylanguage WHERE LANGUAGE = "Spanish"`
`GROUP BY CountryCode`



What are the full names of all left handed players, in order of birth date? **WTA_1**



`SELECT first_name, last_name FROM players ORDER BY birth_date`



`SELECT first_name, last_name FROM players WHERE hand = 'L' ORDER BY birth_date`

Qualitative Examples

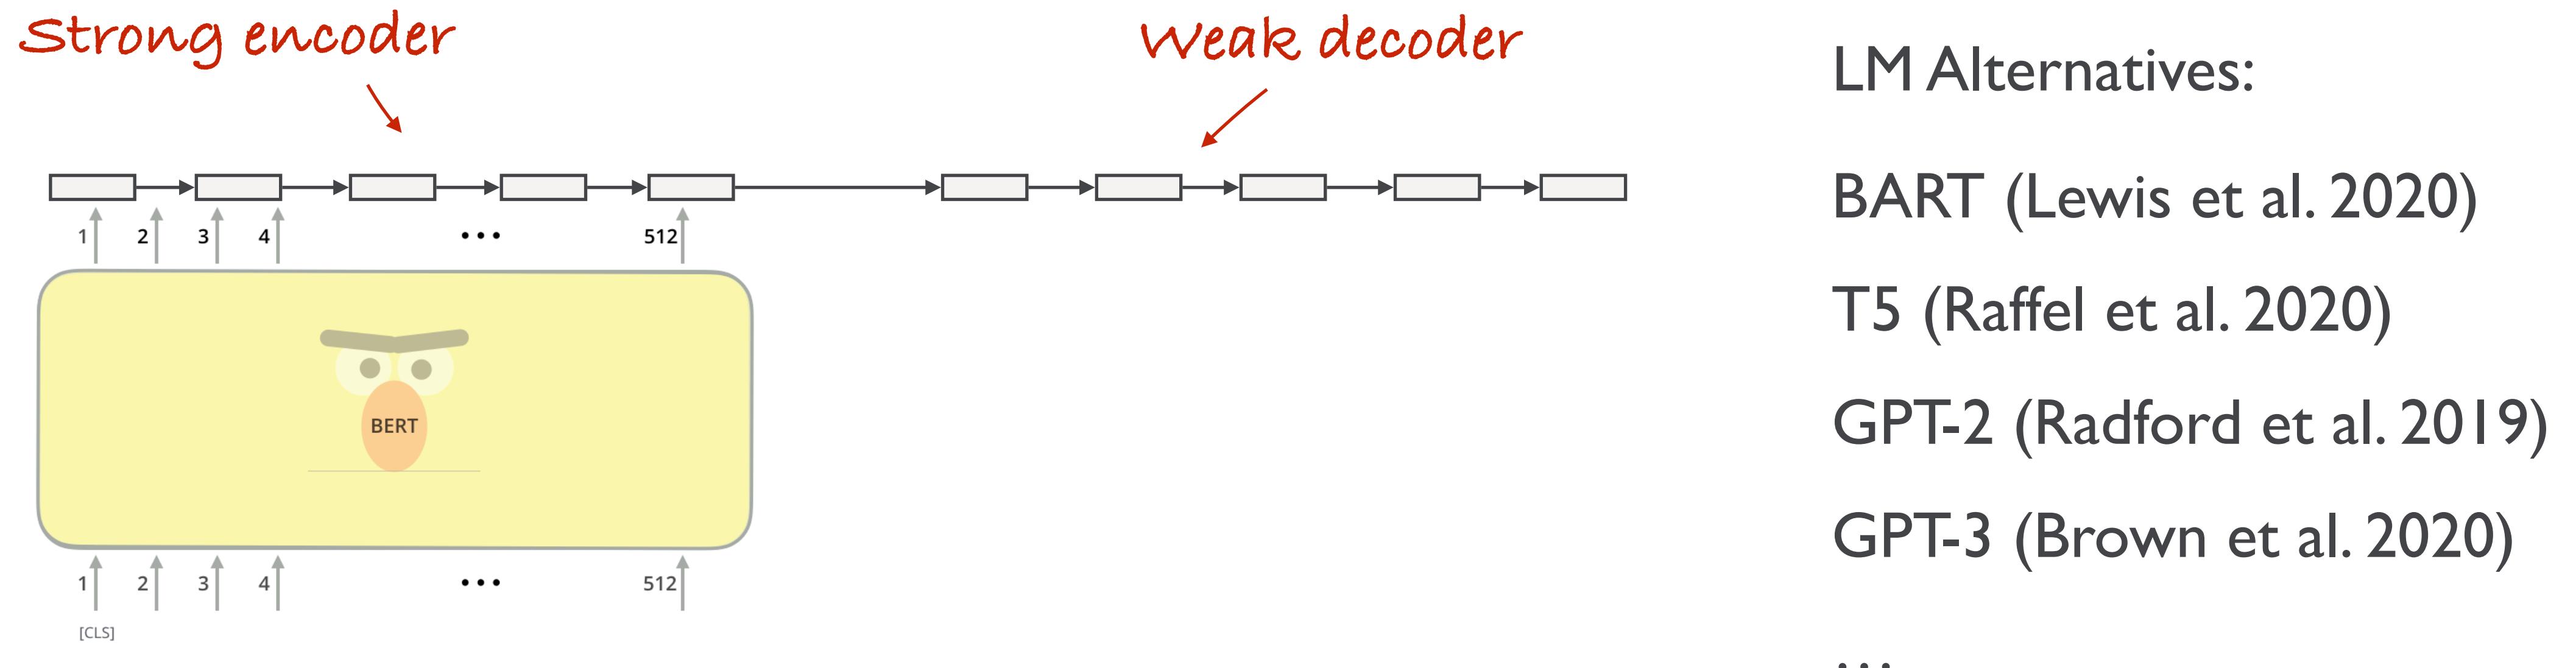


Error Category III - Robustness

Robustness	
	<p>👤 <i>What is the model of the car with the smallest amount of horsepower?</i> car_1</p>
✗	<pre>SELECT cars_data.Horsepower FROM cars_data ORDER BY cars_data.Horsepower LIMIT 1</pre>
✓	<pre>SELECT T1.Model FROM CAR_NAMES AS T1 JOIN CARS_DATA AS T2 ON T1.MakeId = T2.Id ORDER BY T2.horsepower ASC LIMIT 1</pre>
	<p>👤 <i>What is the total population and average area of countries in the continent of North America whose area is bigger than 3000?</i> concert_singer</p>
✗	<pre>SELECT SUM(country.Population), AVG(country.Population) FROM country WHERE country.Continent = "North America" AND country.SurfaceArea > 3000</pre>
✓	<pre>SELECT SUM(country.population), AVG(country.surfacearea) FROM country WHERE country.Continent = "north america" and country.SurfaceArea > 3000</pre>

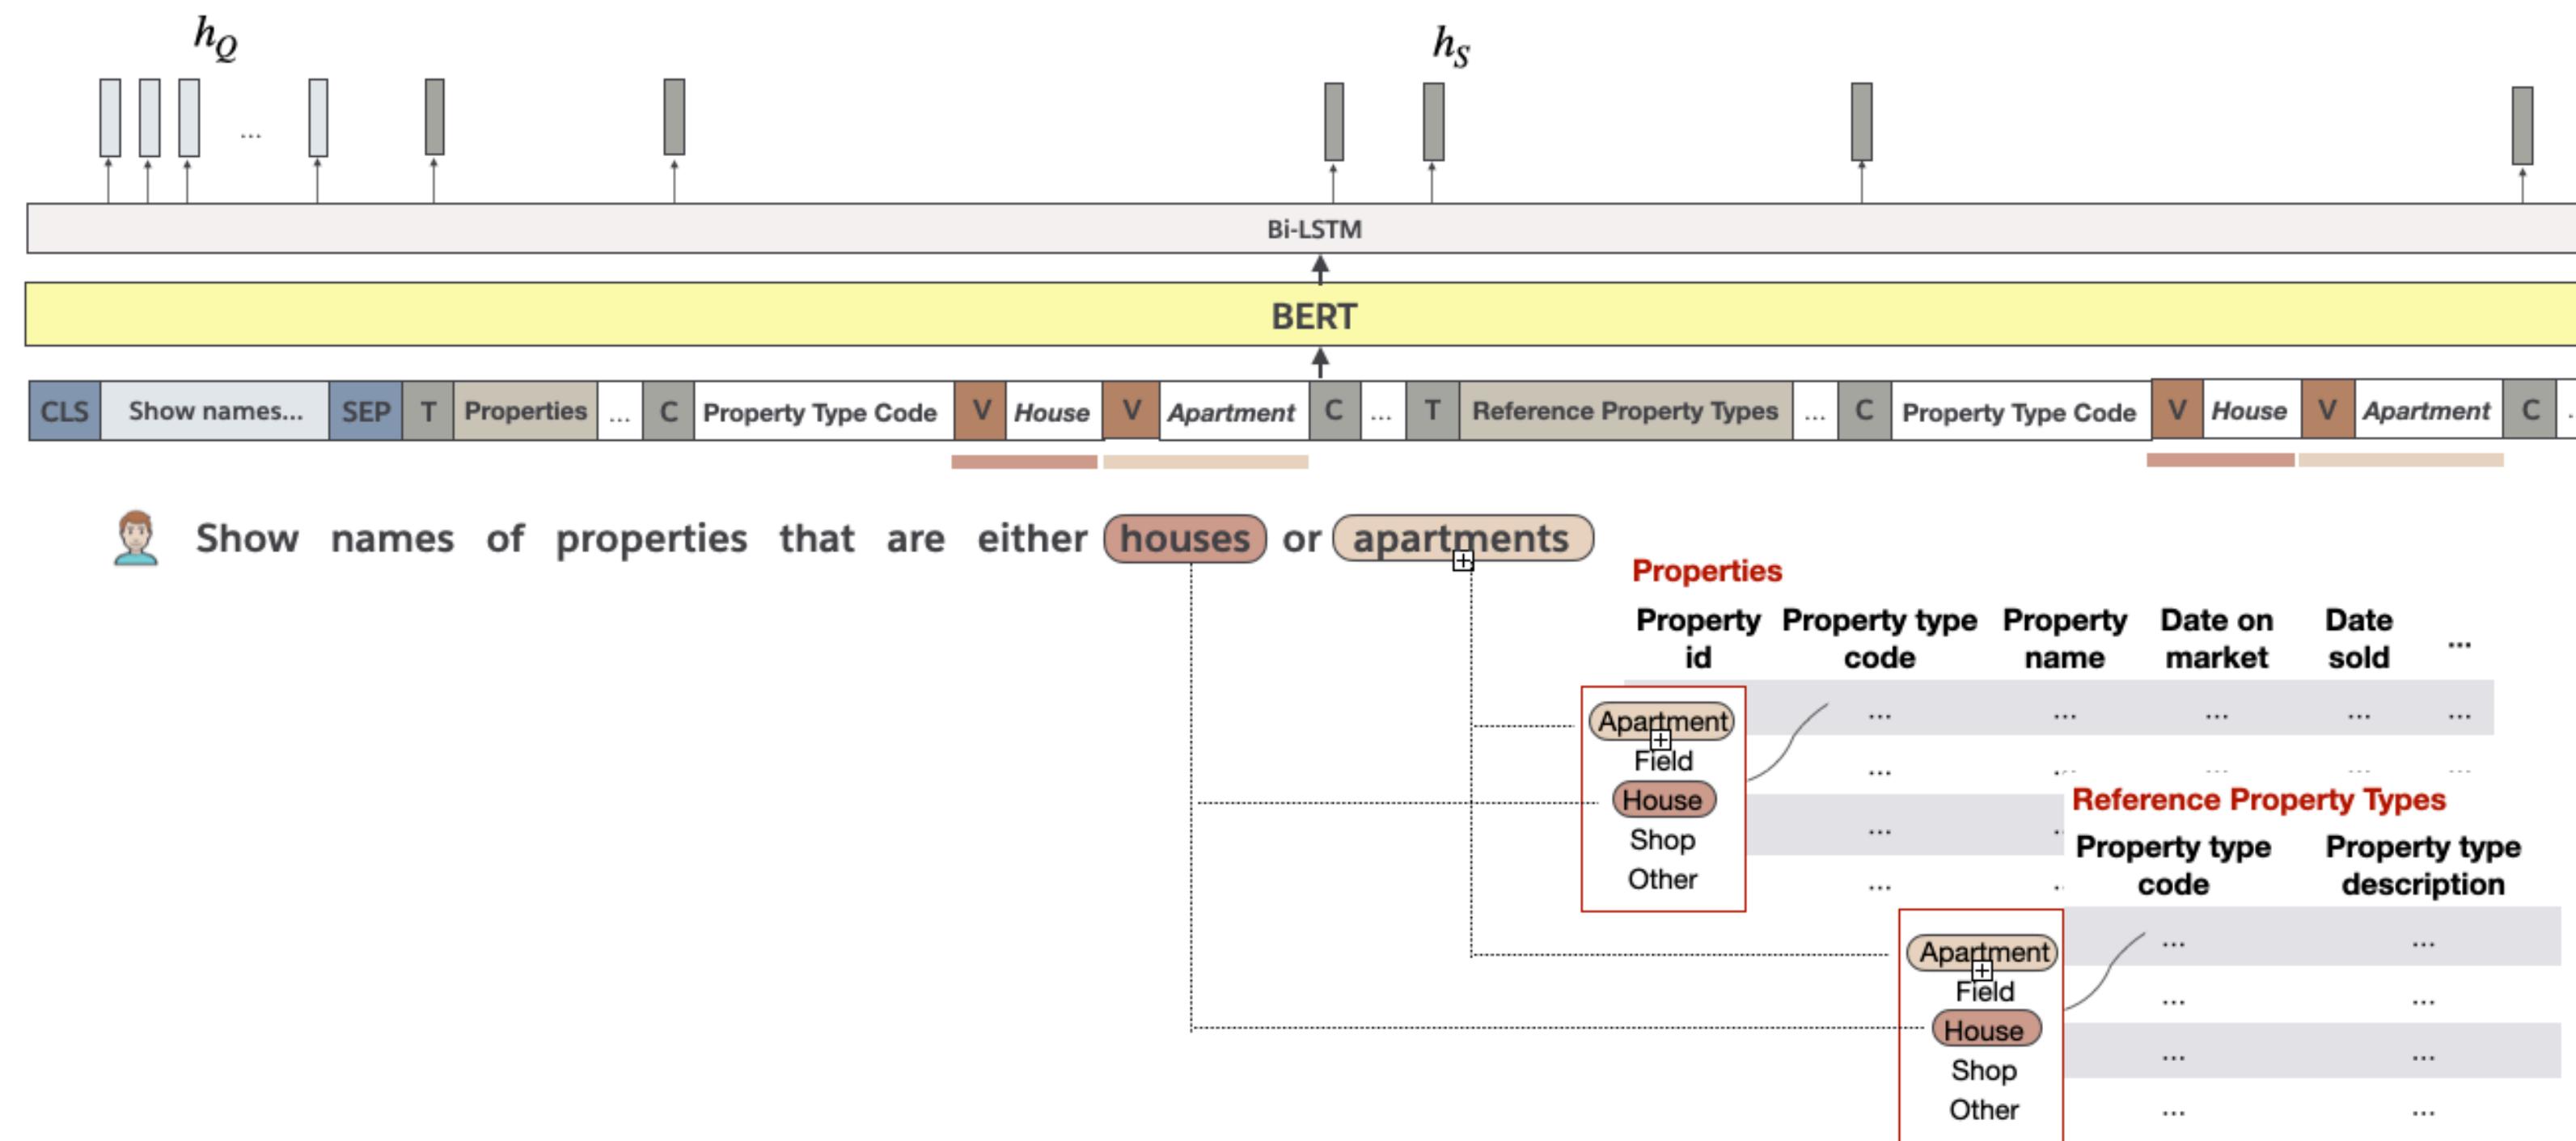
Discussion

- I. BRIDGE uses a sequential encoder for jointly encoding **text**, **DB schema** and relevant **DB cells**, and a sequential decoder for generating **SQL queries**. The decoder has significantly fewer parameters than the encoder.



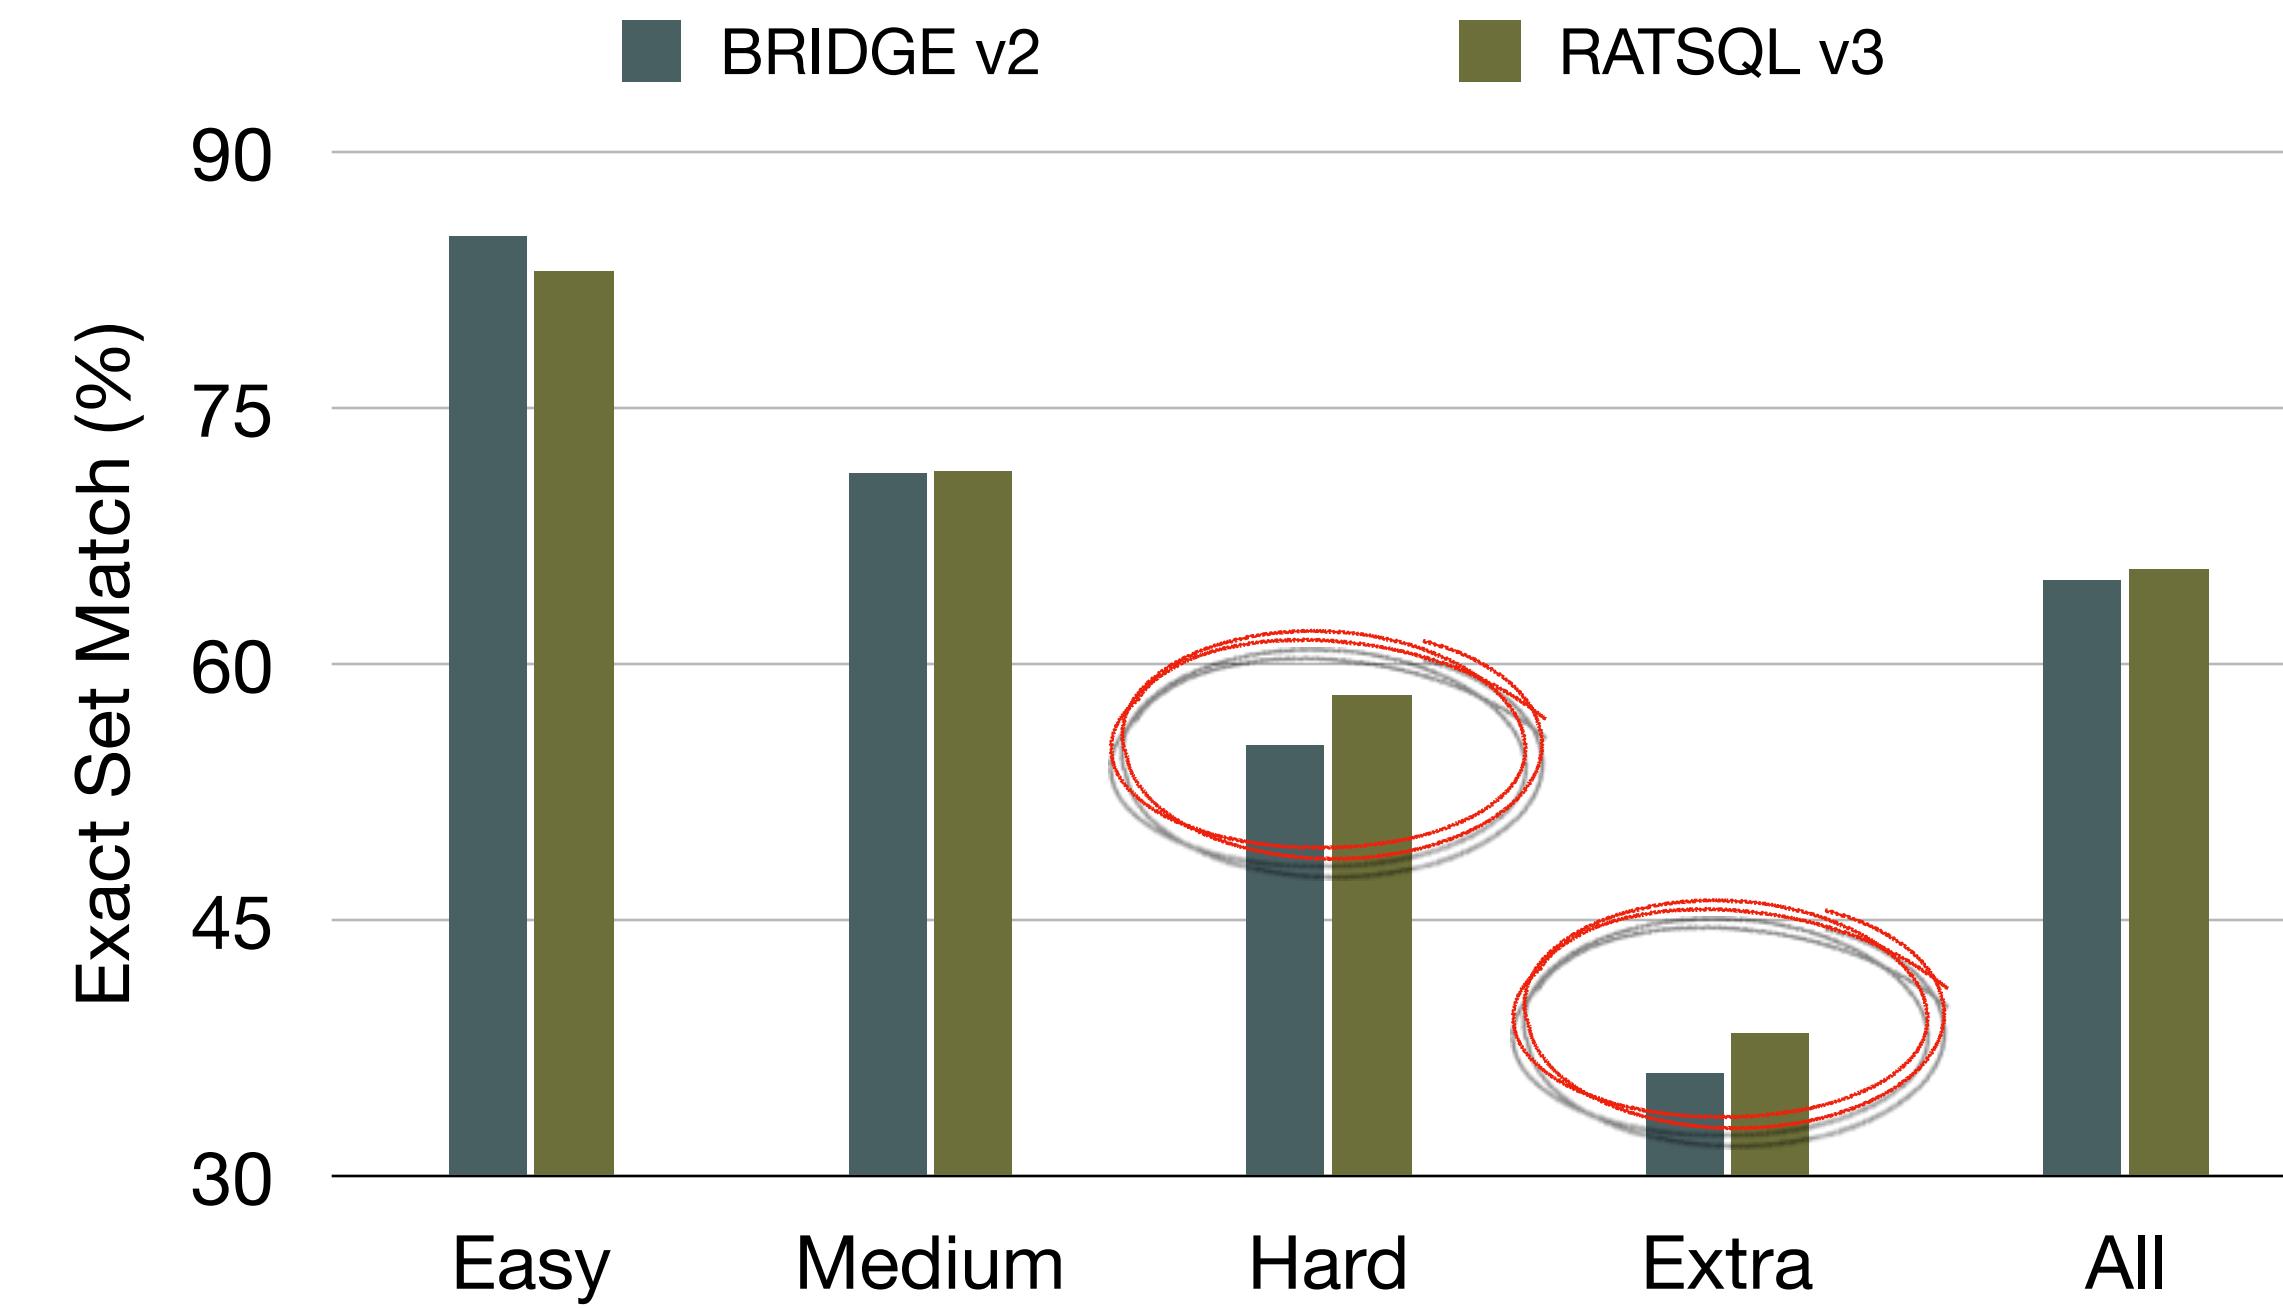
Discussion

II. Learning to recognize relevant cells (addressing acronyms and other lexical variations)



Discussion

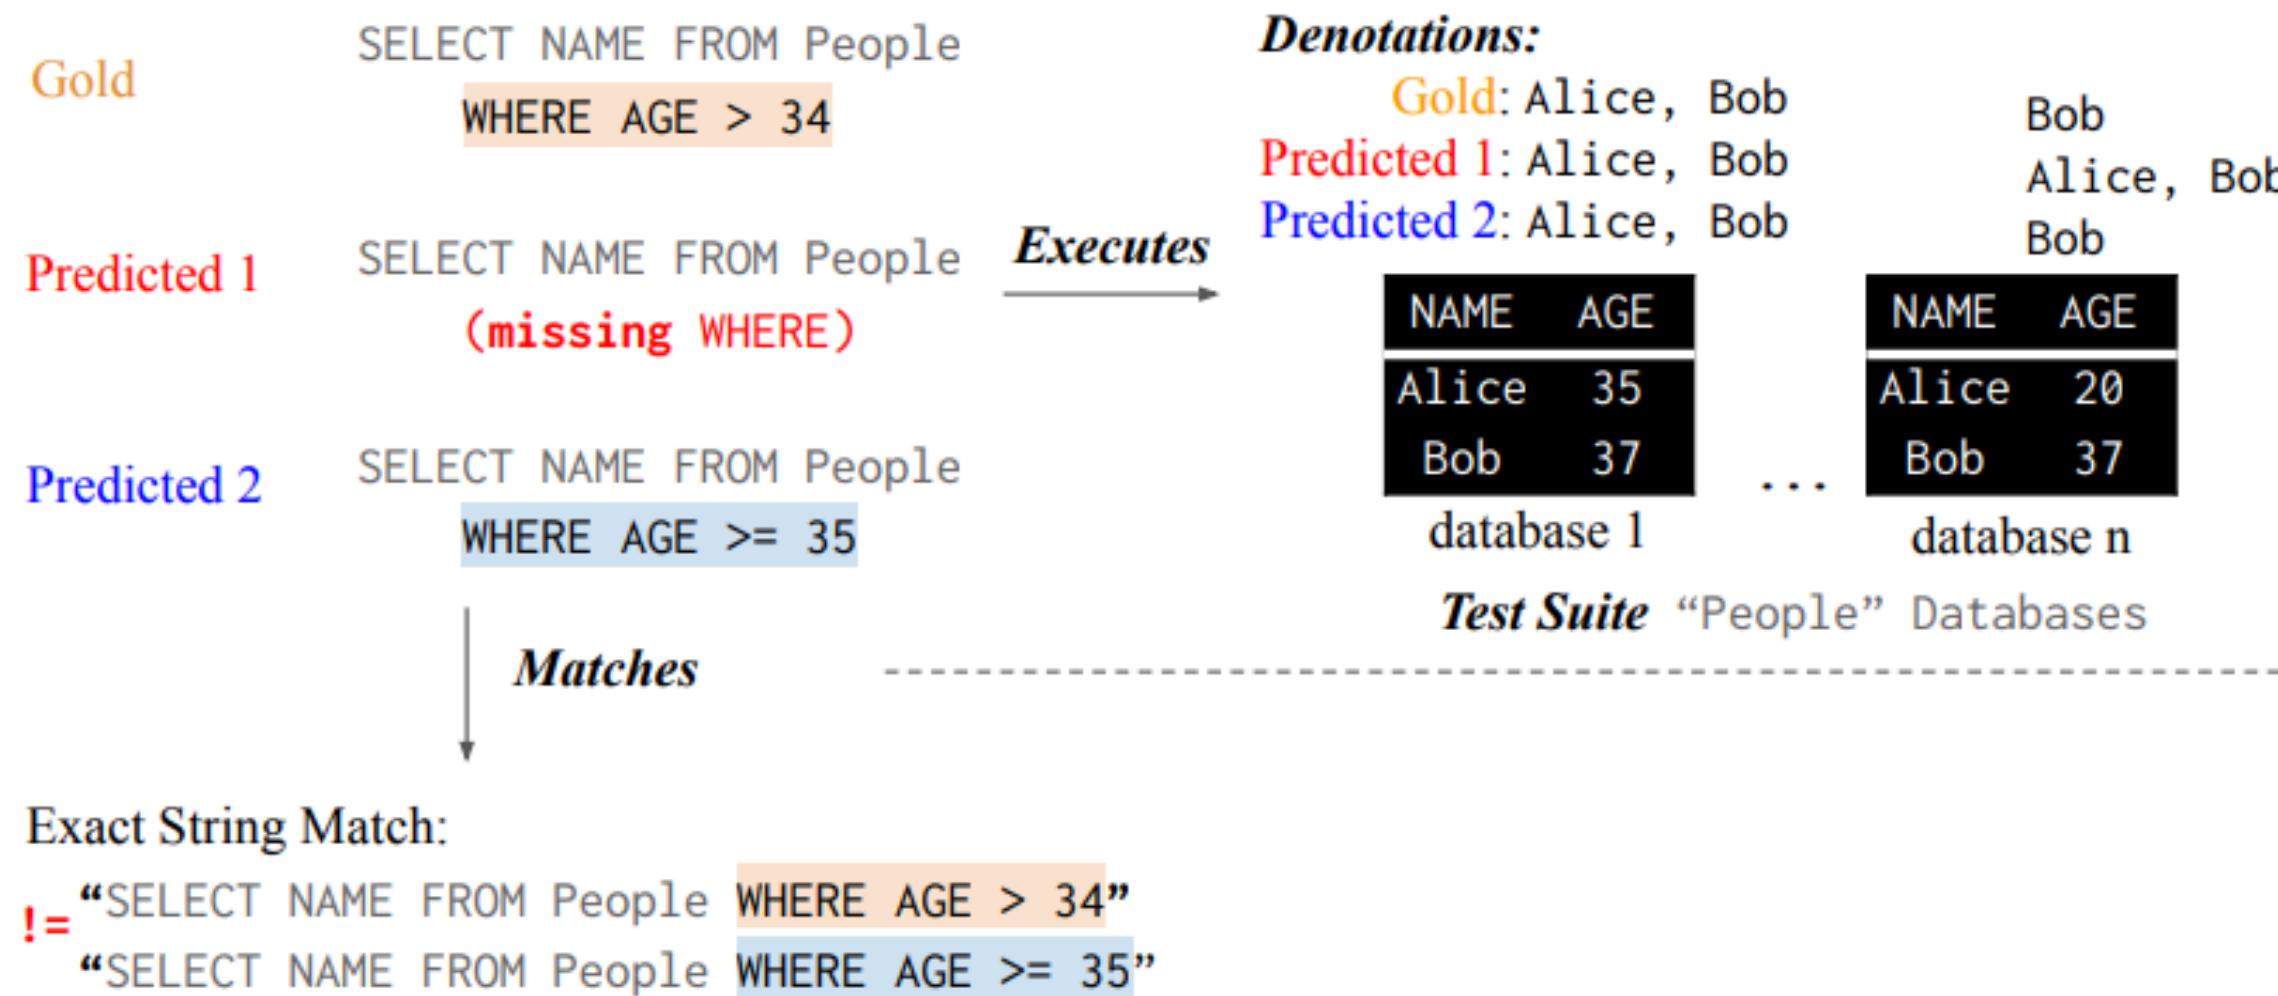
III. Compositional inductive biases (CIBs) show significant benefit for synthesizing hard SQL queries. Previous work have shown that CIBs are effective for improving model's compositional generalization and overcoming data sparsity.



Discussion

IV. More comprehensive model evaluation and benchmarking.

Query: "Who is above 34 years old?"



Distilled Test Suites by (Zhong et al. 2020)

Discussion

V. Existing benchmark datasets are not perfect.

- Sparse schema component coverage
- Sparse logic relation coverage
- Data synthesis?
- Interpolation?

The screenshot shows a database interface with the following components:

- Left Sidebar (Tables):** Lists the tables in the schema: machine_repair, machine, repair, and a query result table.
- Machine Table:** Shows data for machines across five years (1991-1995). The columns are Machine_ID, Making_Year, Class, Team, Machine_series, and value_points.
- Repair Table:** Shows data for repairs. The columns are repair_ID, name, Launch_Date, and Notes.
- Query Result Table:** Shows a subset of data from the machine table, filtered by Machine_series = "RS125". The columns are Class, Machine_ID, Machine_series, Making_Year, Team, quality_rank, and value_points.
- Chat Sidebar:**
 - Message 1:** Chat started by Photon • 8:53:51 PM
 - Message 2:** DB switched to src-2
 - Message 3:** Hello! Please input your question in NL or SQL to query the DB
 - Message 4:** DB switched to machine_repair
 - Message 5:** What is the launch date for machines in series RS125
 - Message 6:** SELECT * FROM machine WHERE machine.Machine_series = "RS125"
 - Message 7:** Did I get it right?
 - Buttons:** Yes and No
 - Text Input:** Type Here

Discussion

VI. More future directions

- Train with execution feedback
- Overcome data sparsity
- Interpretability and Explainability
- Process context and pragmatics
- Question answering over DBs, documents, and other modality of information
- ...

<https://github.com/salesforce/TabularSemanticParsing>

 Star

51

 Fork

11

Pre-trained transformer LMs can effectively capture language-database grounding when the cross-modal data are serialized and tagged with special tokens

Two strategies significantly contribute to the overall text-to-SQL performance

- The bridging mechanism that appends field values mentions (anchor texts) to the corresponding field names in the serialized representation
- Search-space pruning based on SQL syntax and schema consistency

Co-authors



Richard Socher

You.com



Caiming Xiong

Salesforce AI
Research

